

Supervised learning

MACHINE LEARNING FOR EVERYONE



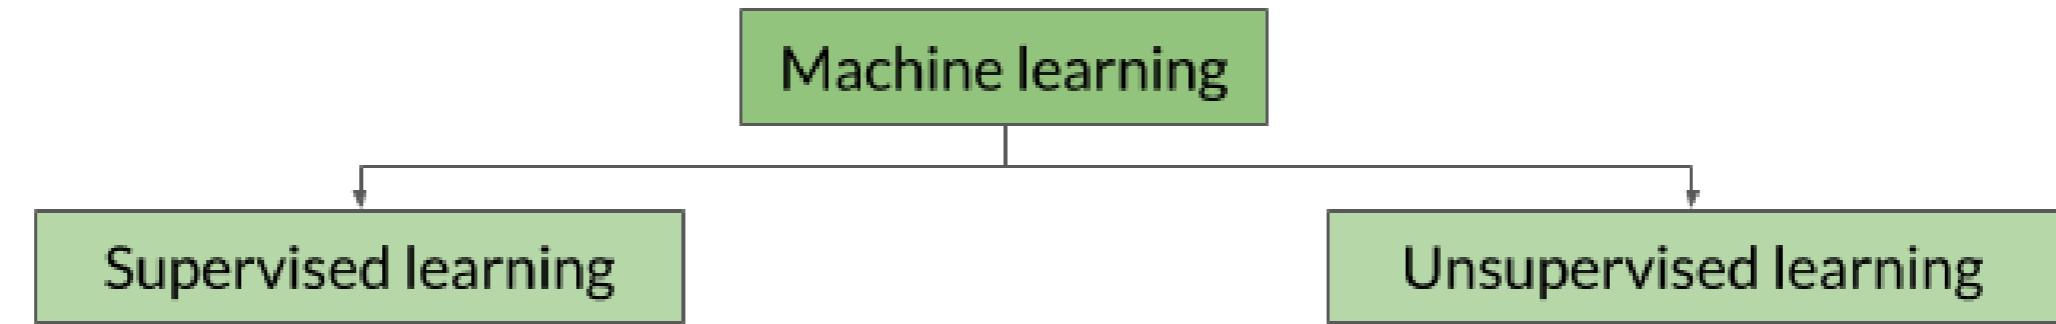
Hadrien Lacroix

Content Developer at DataCamp

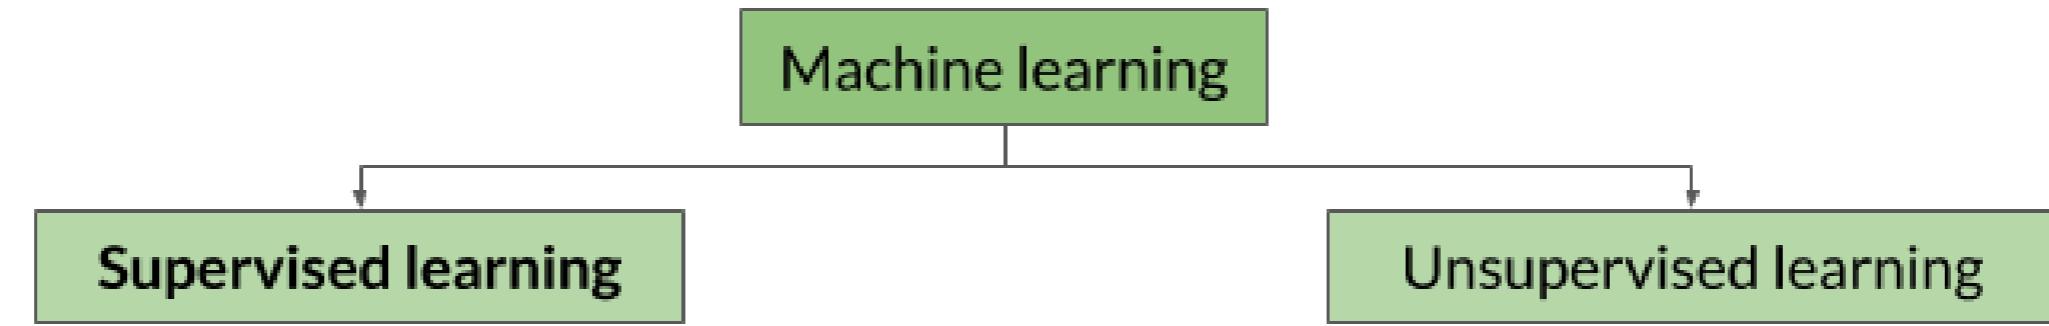
Modeling



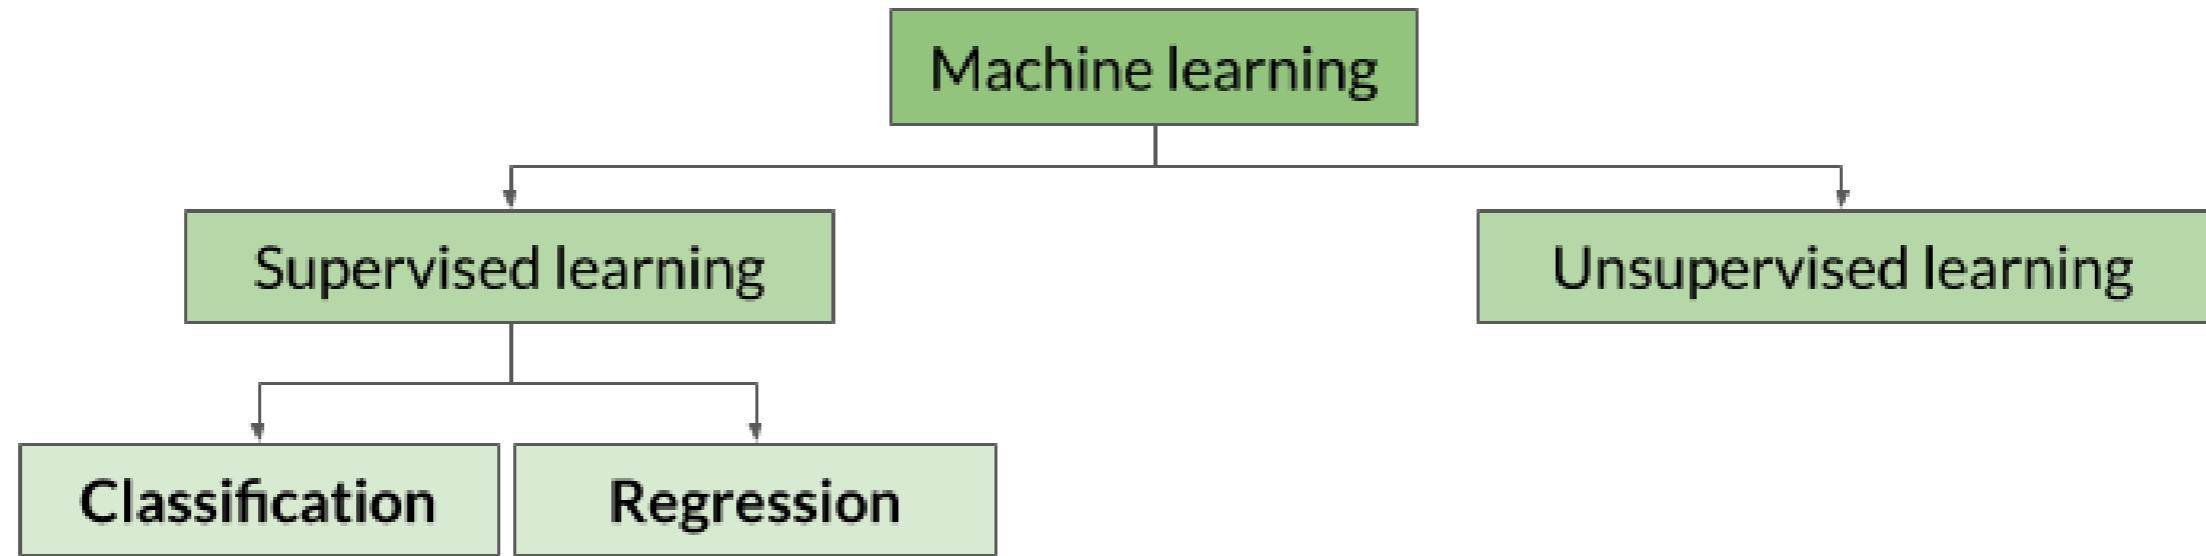
Types



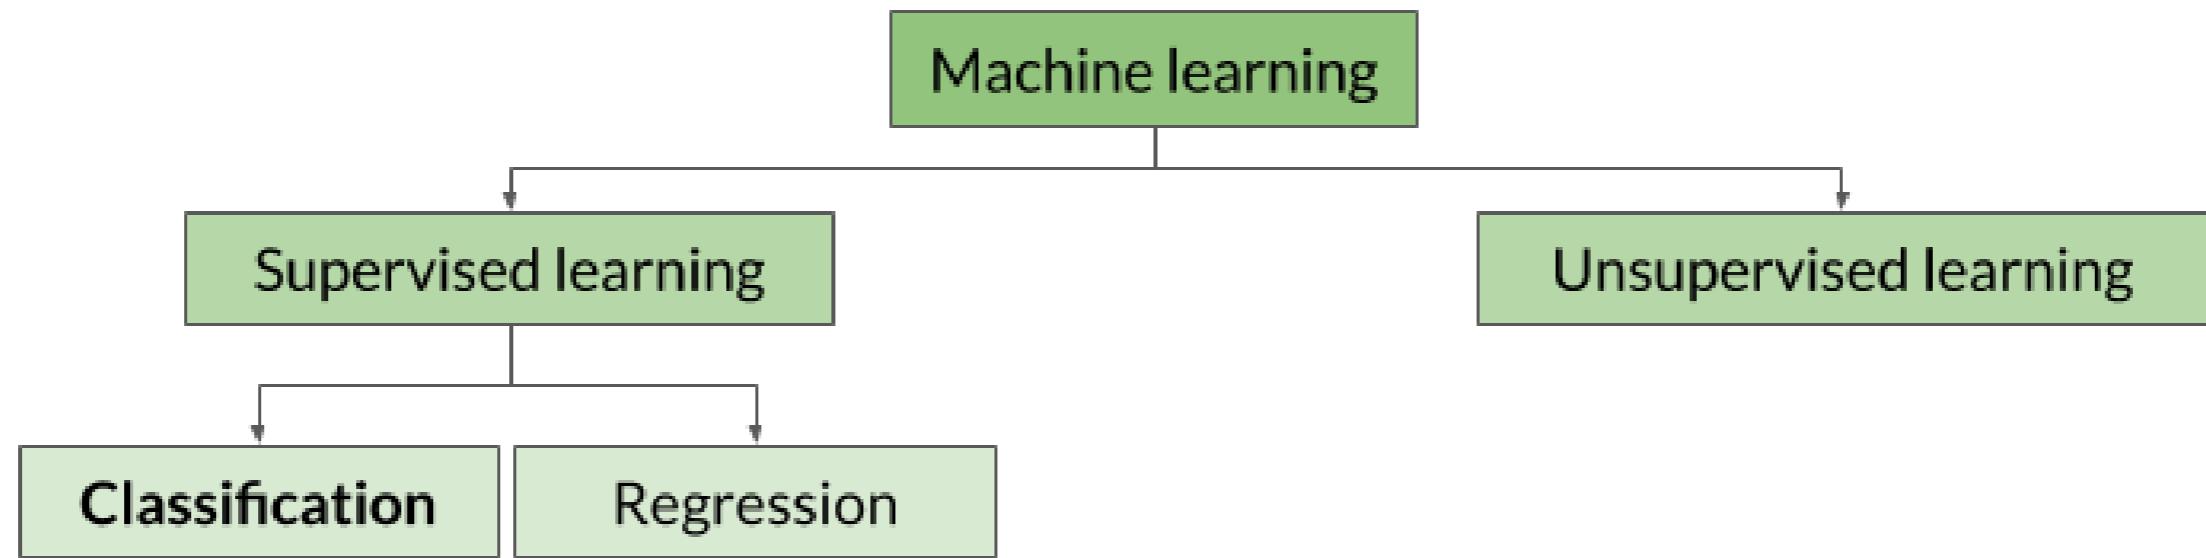
What is supervised learning?



Classification and regression



Classification



Classification

- **Classification** = assigning a **category**
 - Will this customer **stop** its subscription?
 - Yes, No
 - Is this mole **cancerous**?
 - Yes, No
 - What **kind** of wine is that?
 - Red, White, Rosé
 - What **flower** is that?
 - Rose, Tulip, Carnation, Lily

Observations

	Applicant ID	High school GPA	Test results	Accepted
First observation	0	3.5	2.4	False
Second observation	1	4	2.2	False
Third observation	2	4.2	4.3	True
Fourth observation	3	4.8	2.9	False
n observations

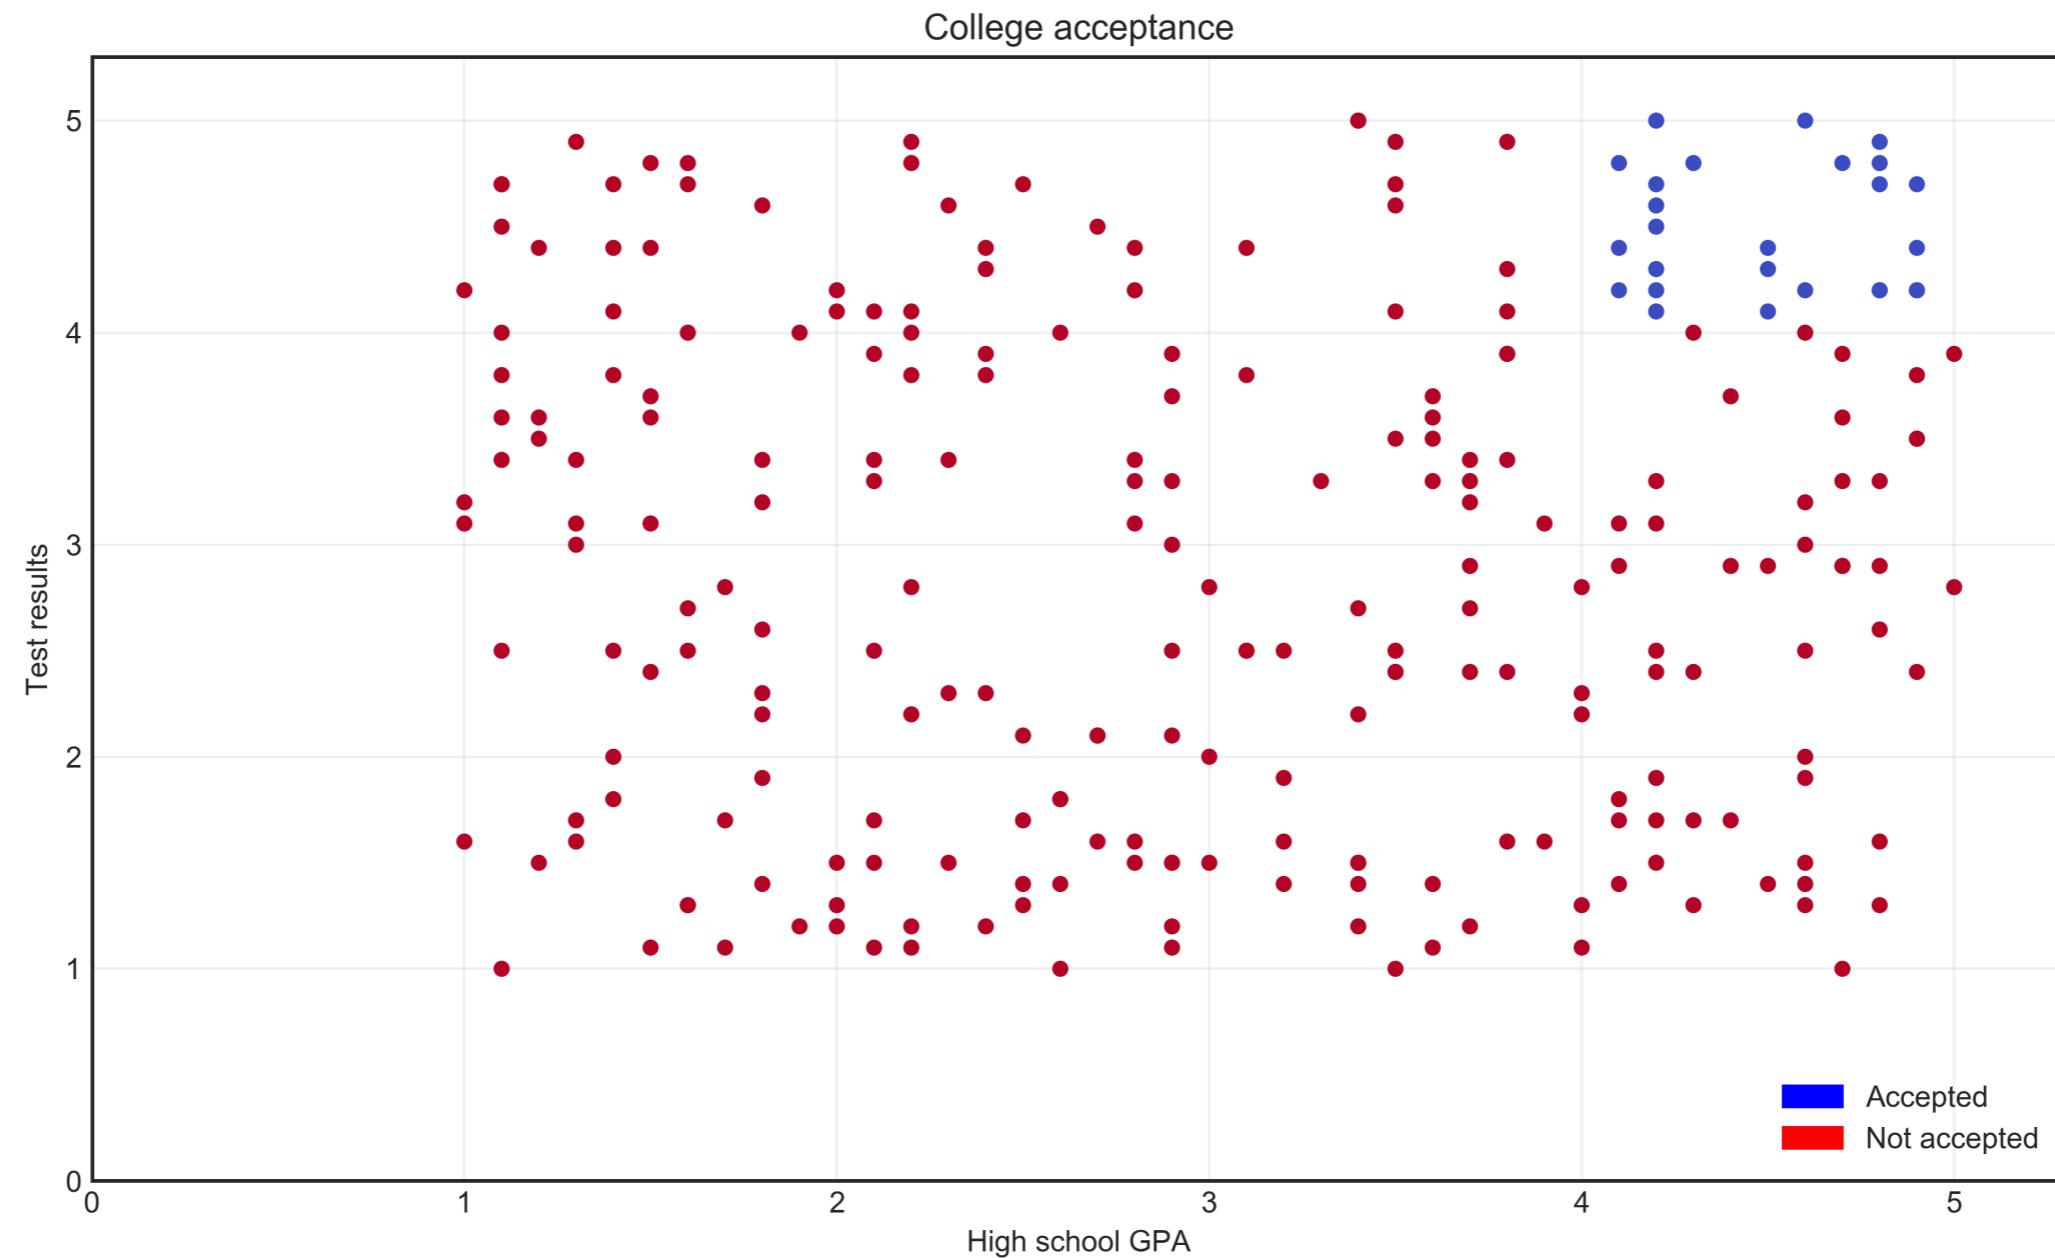
Features

Applicant ID	High school GPA	Test results	Accepted
0	3.5	2.4	False
1	4	2.2	False
2	4.2	4.3	True
3	4.8	2.9	False
...

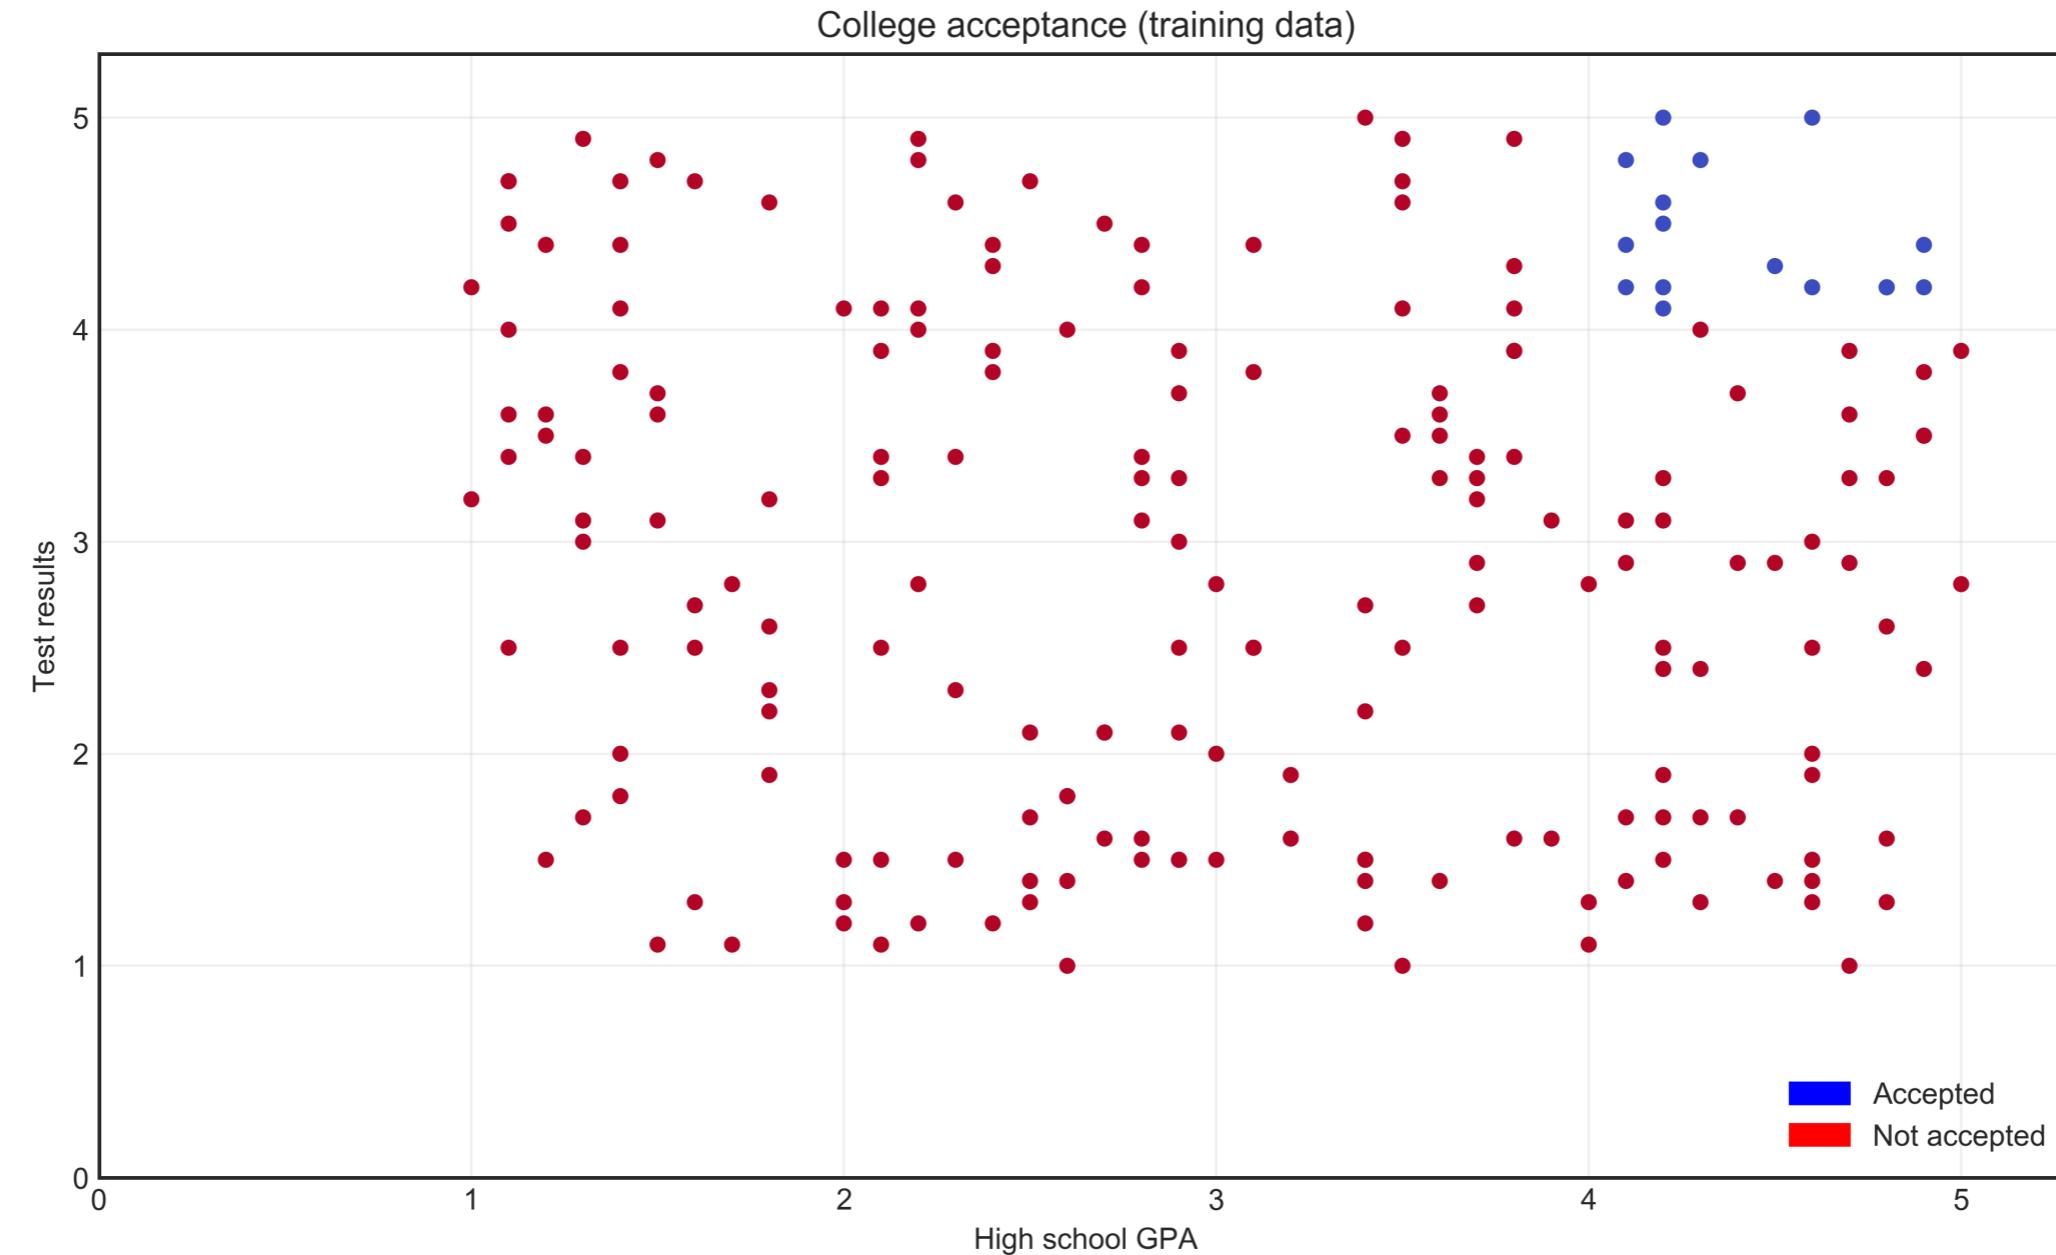
Target

Applicant ID	High school GPA	Test results	Accepted
0	3.5	2.4	False
1	4	2.2	False
2	4.2	4.3	True
3	4.8	2.9	False
...

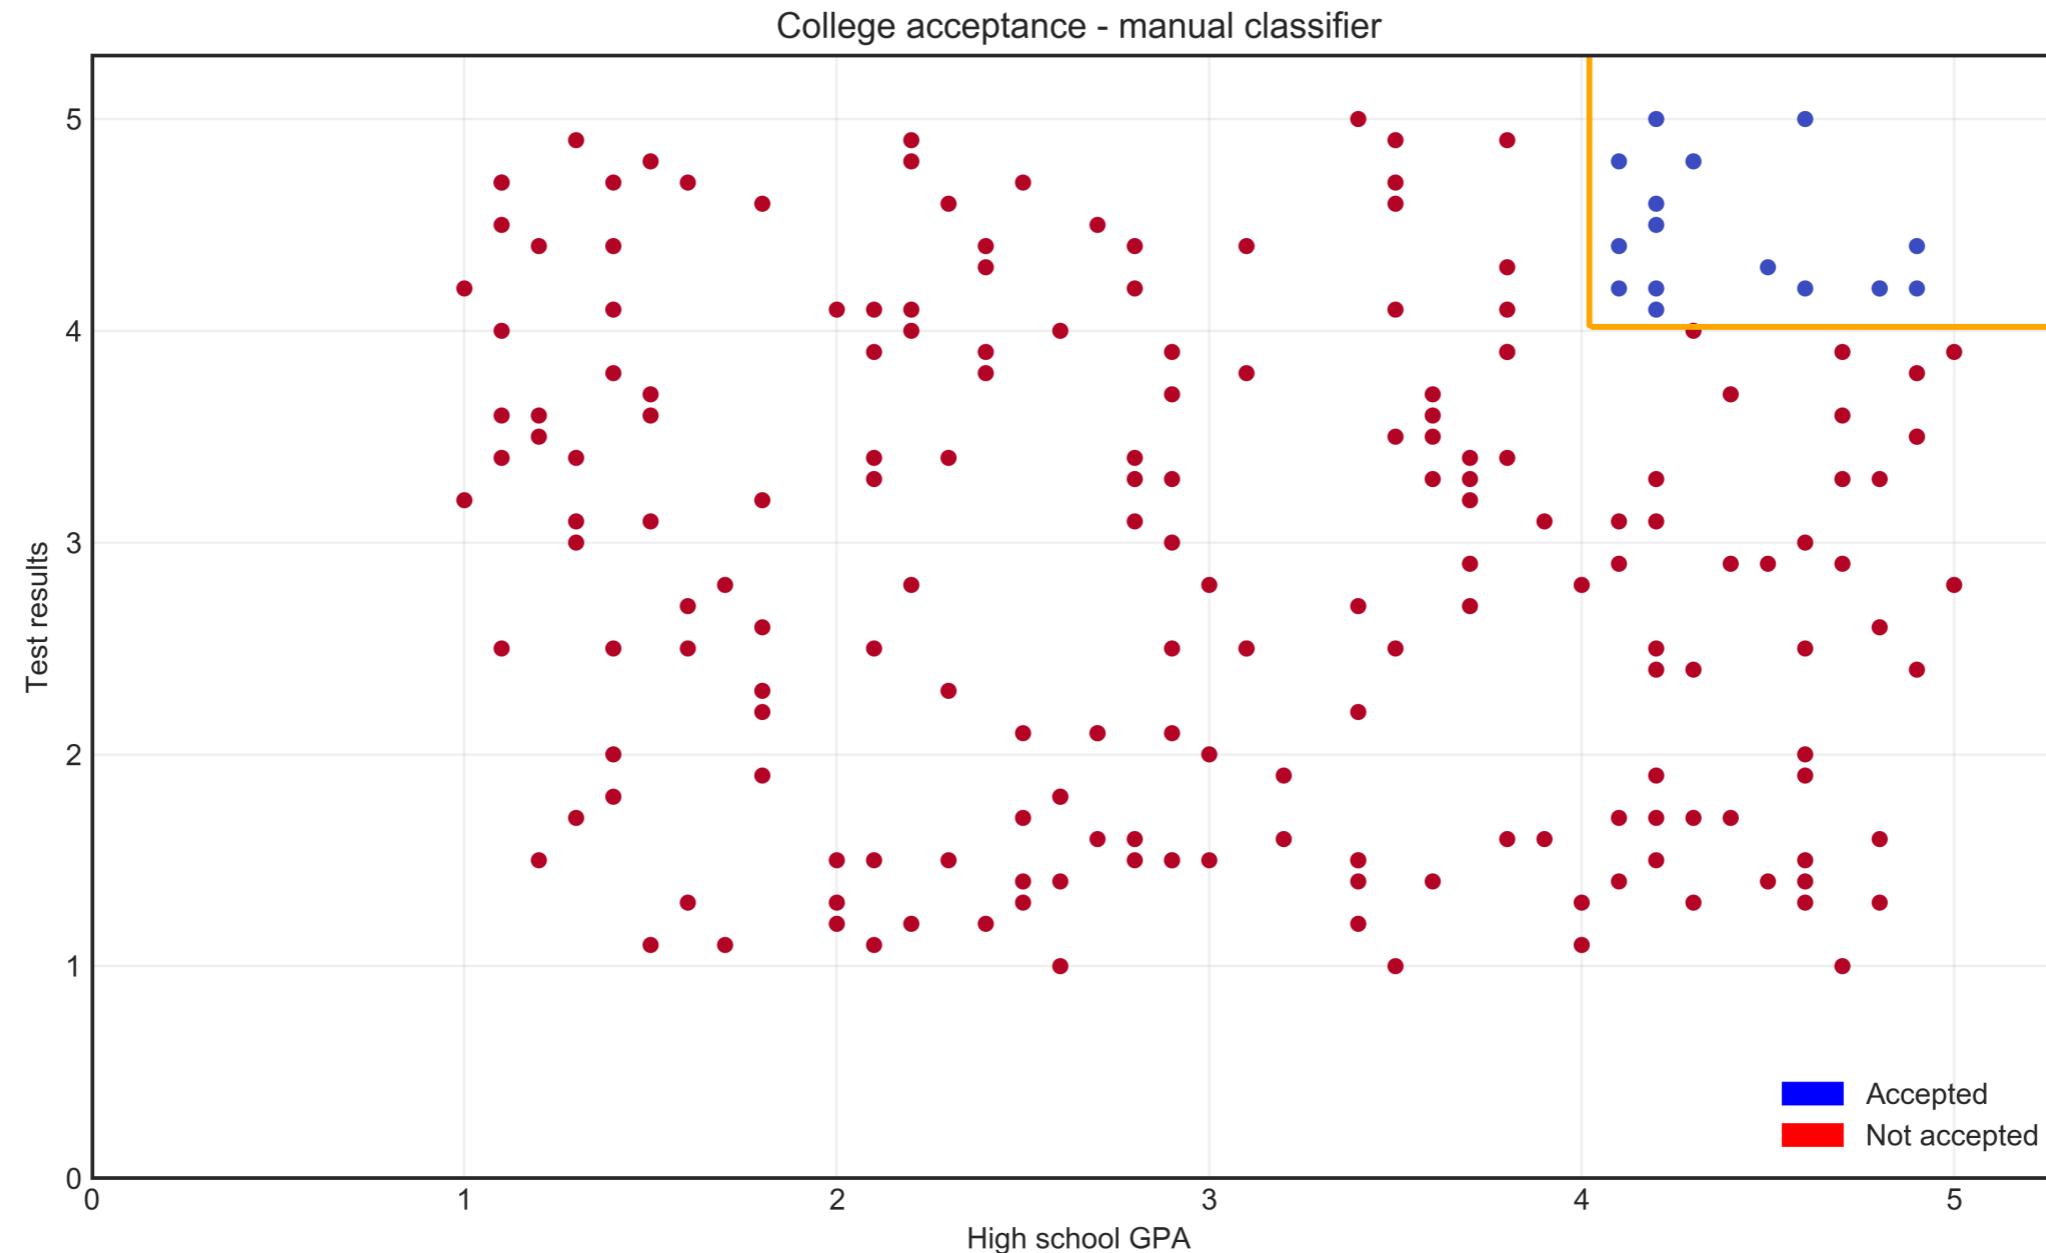
Graphing our data



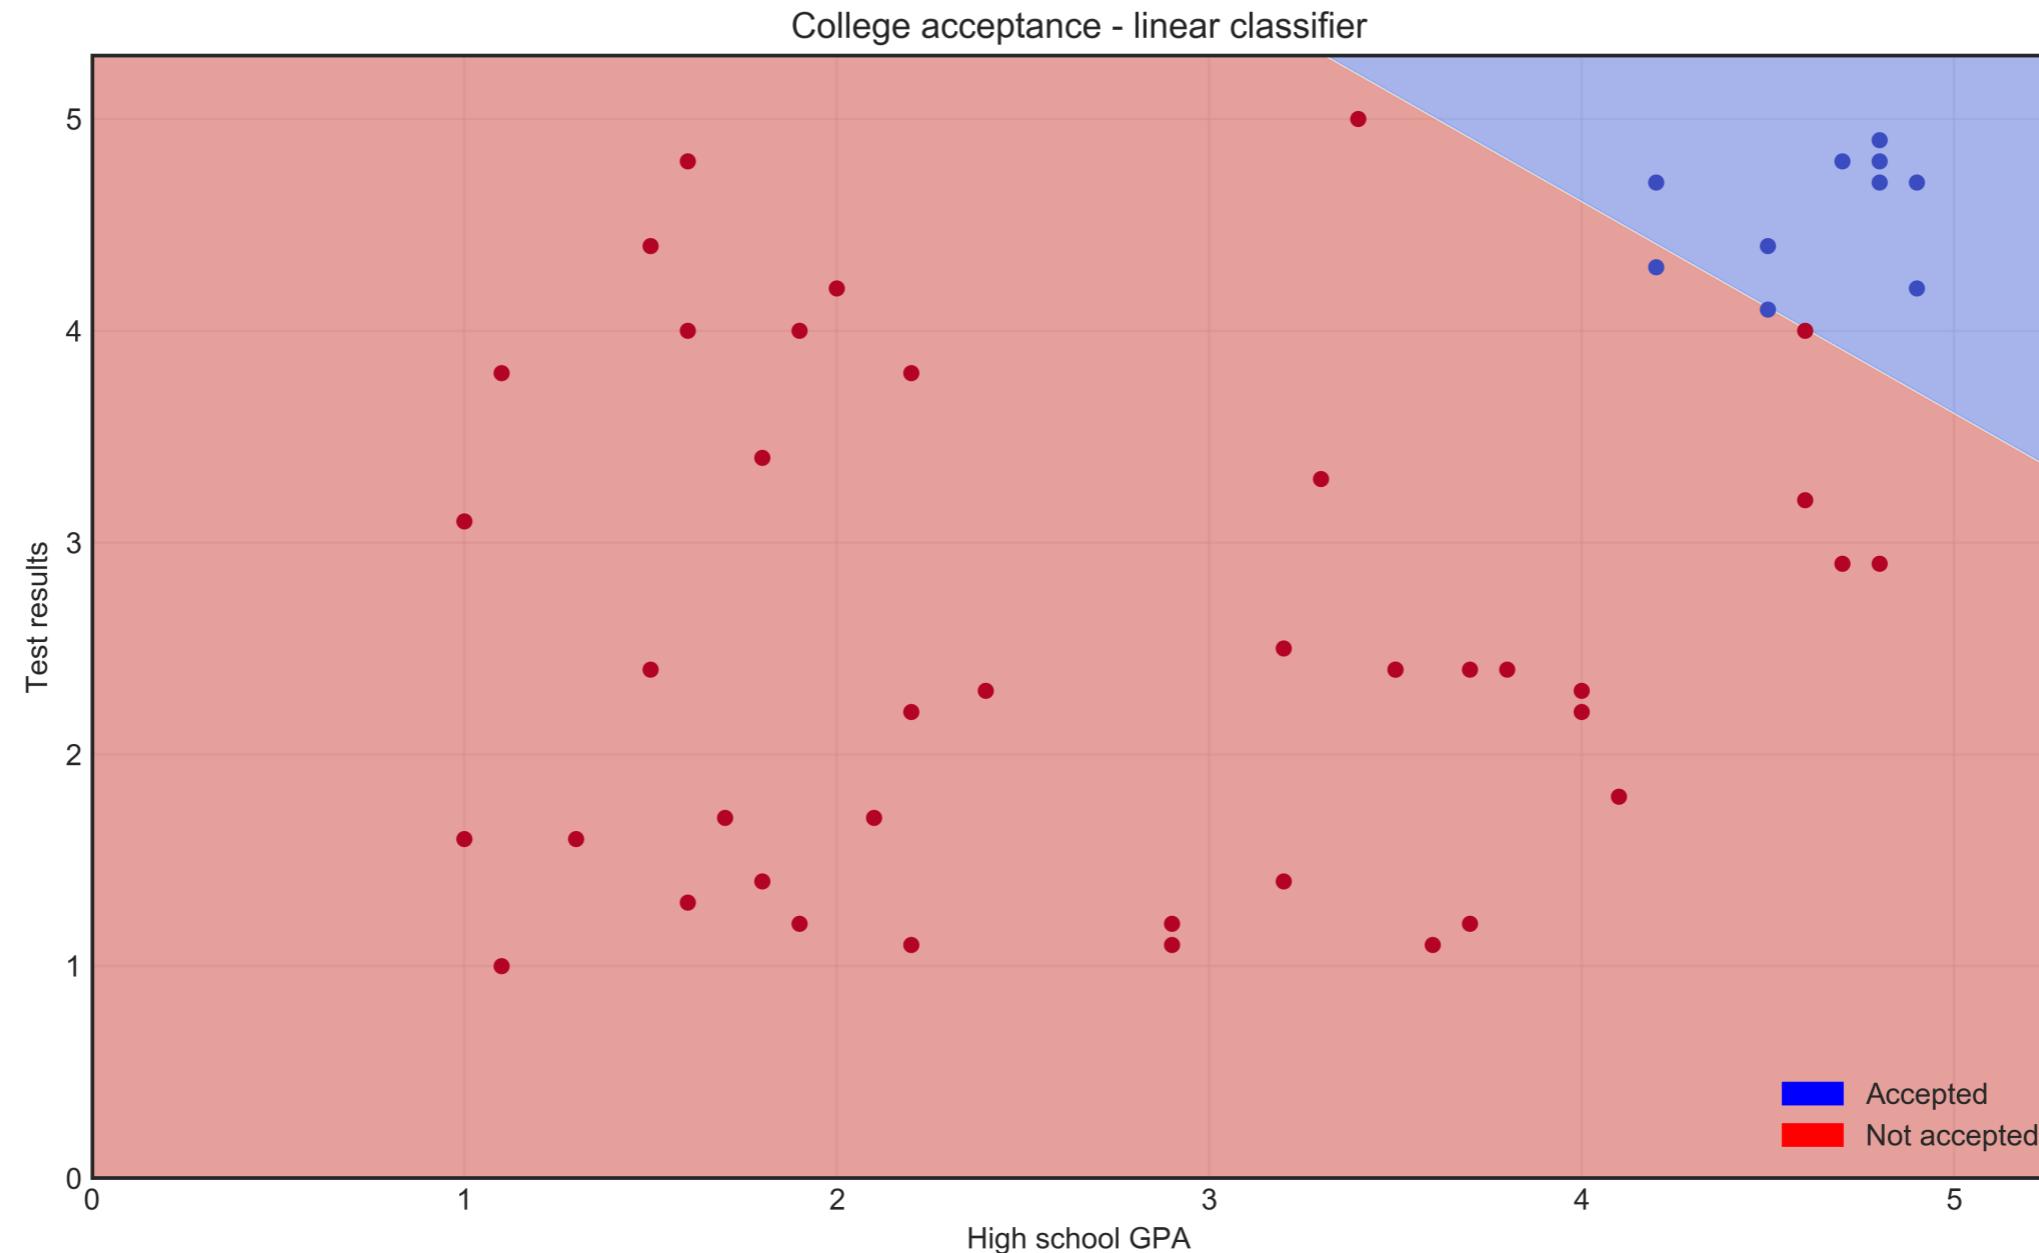
Splitting data



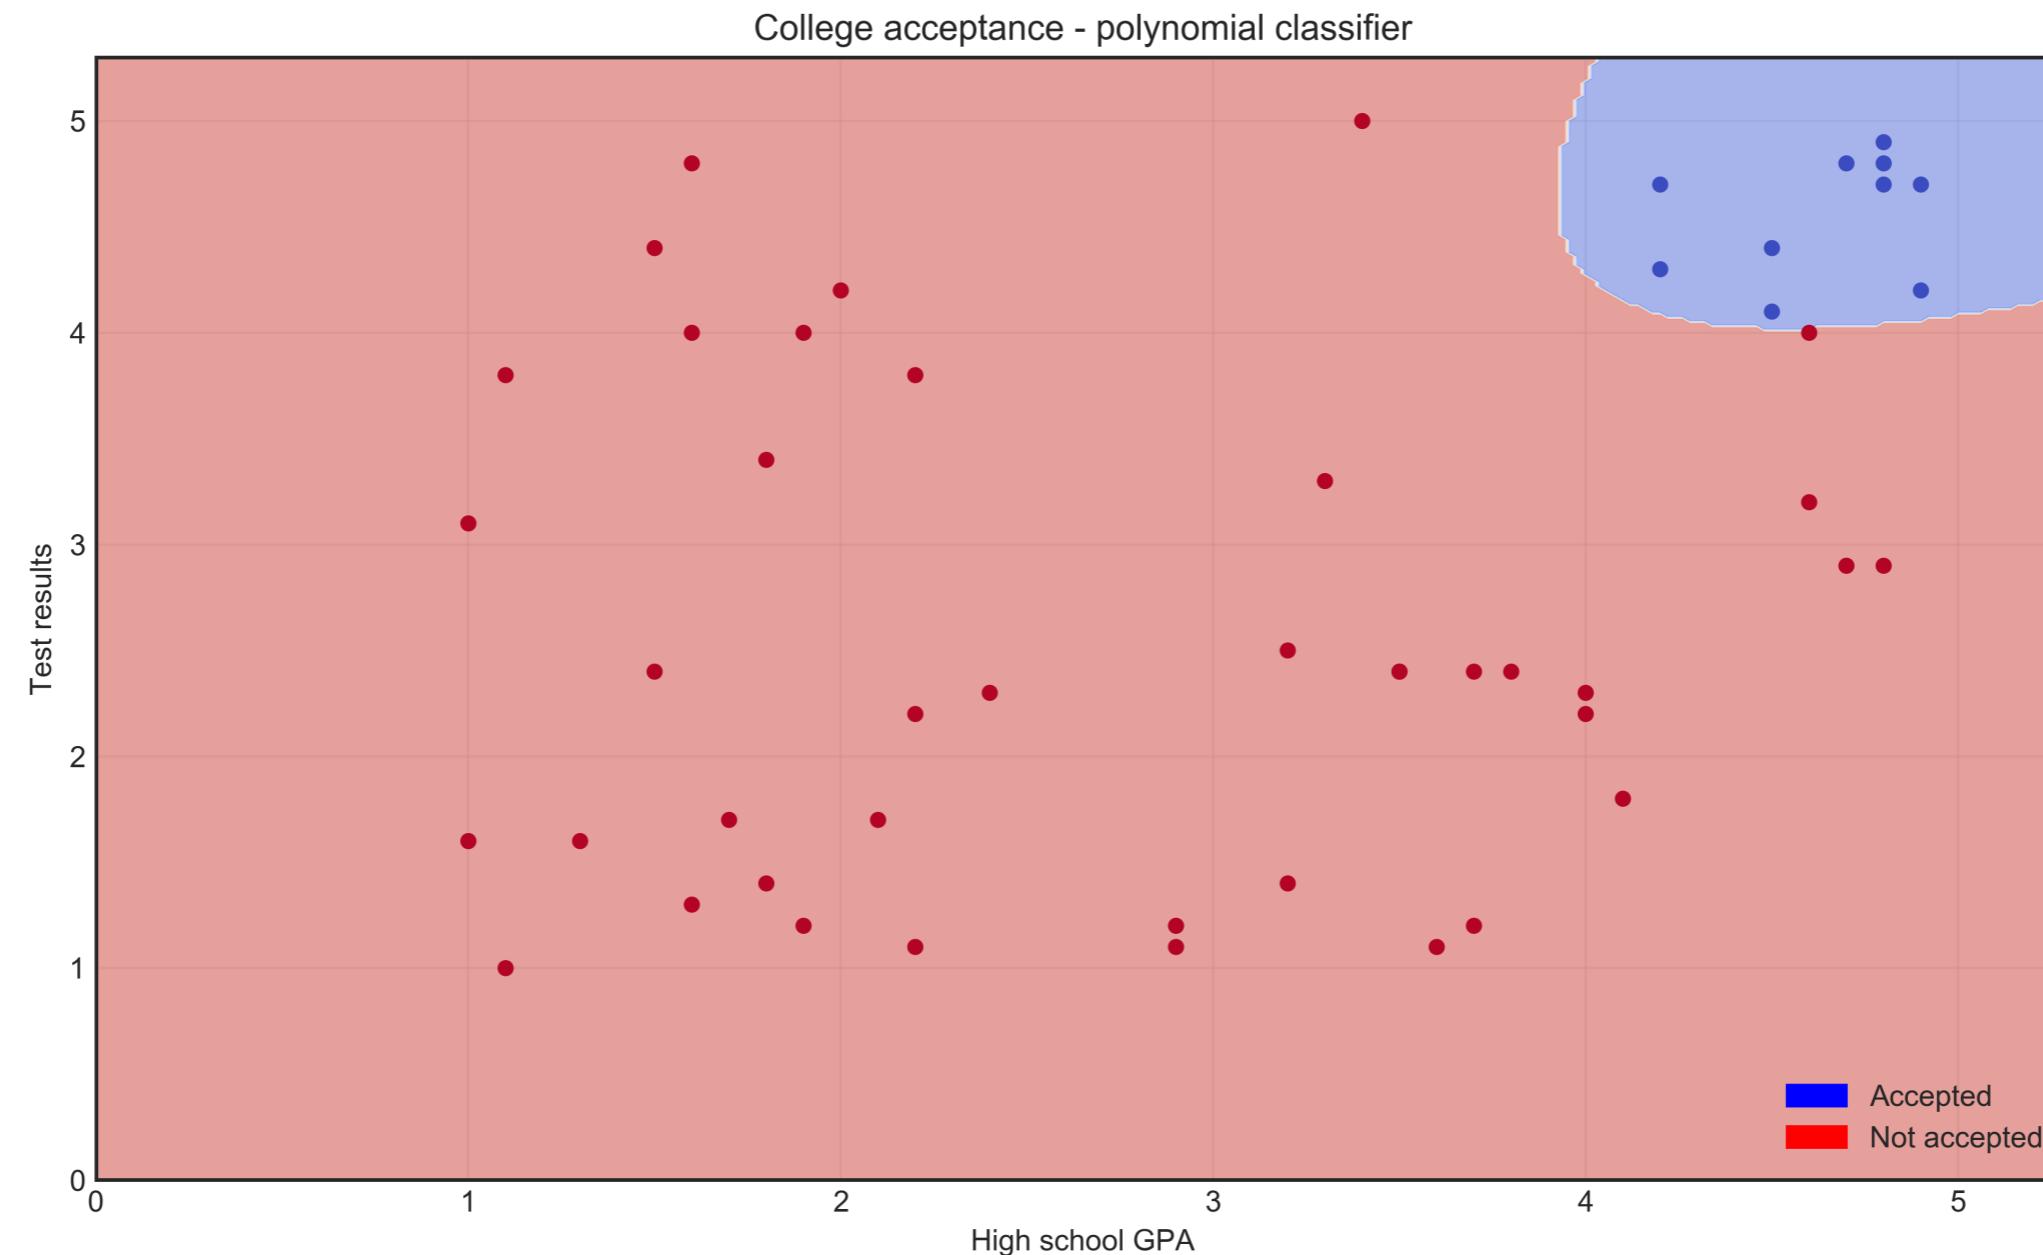
Manual classifier



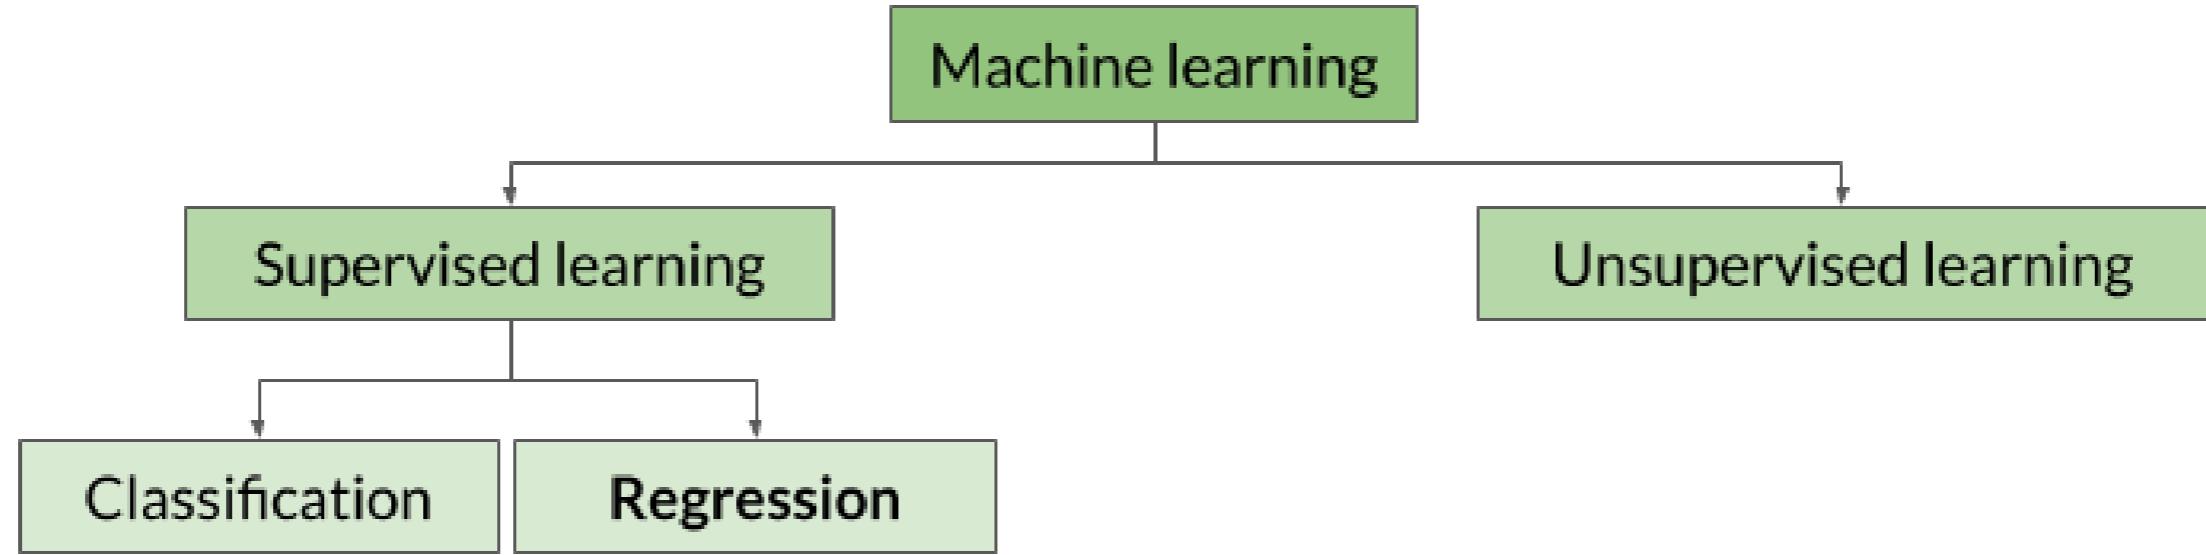
Support vector machine - linear classifier



Support vector machine - polynomial classifier



Regression



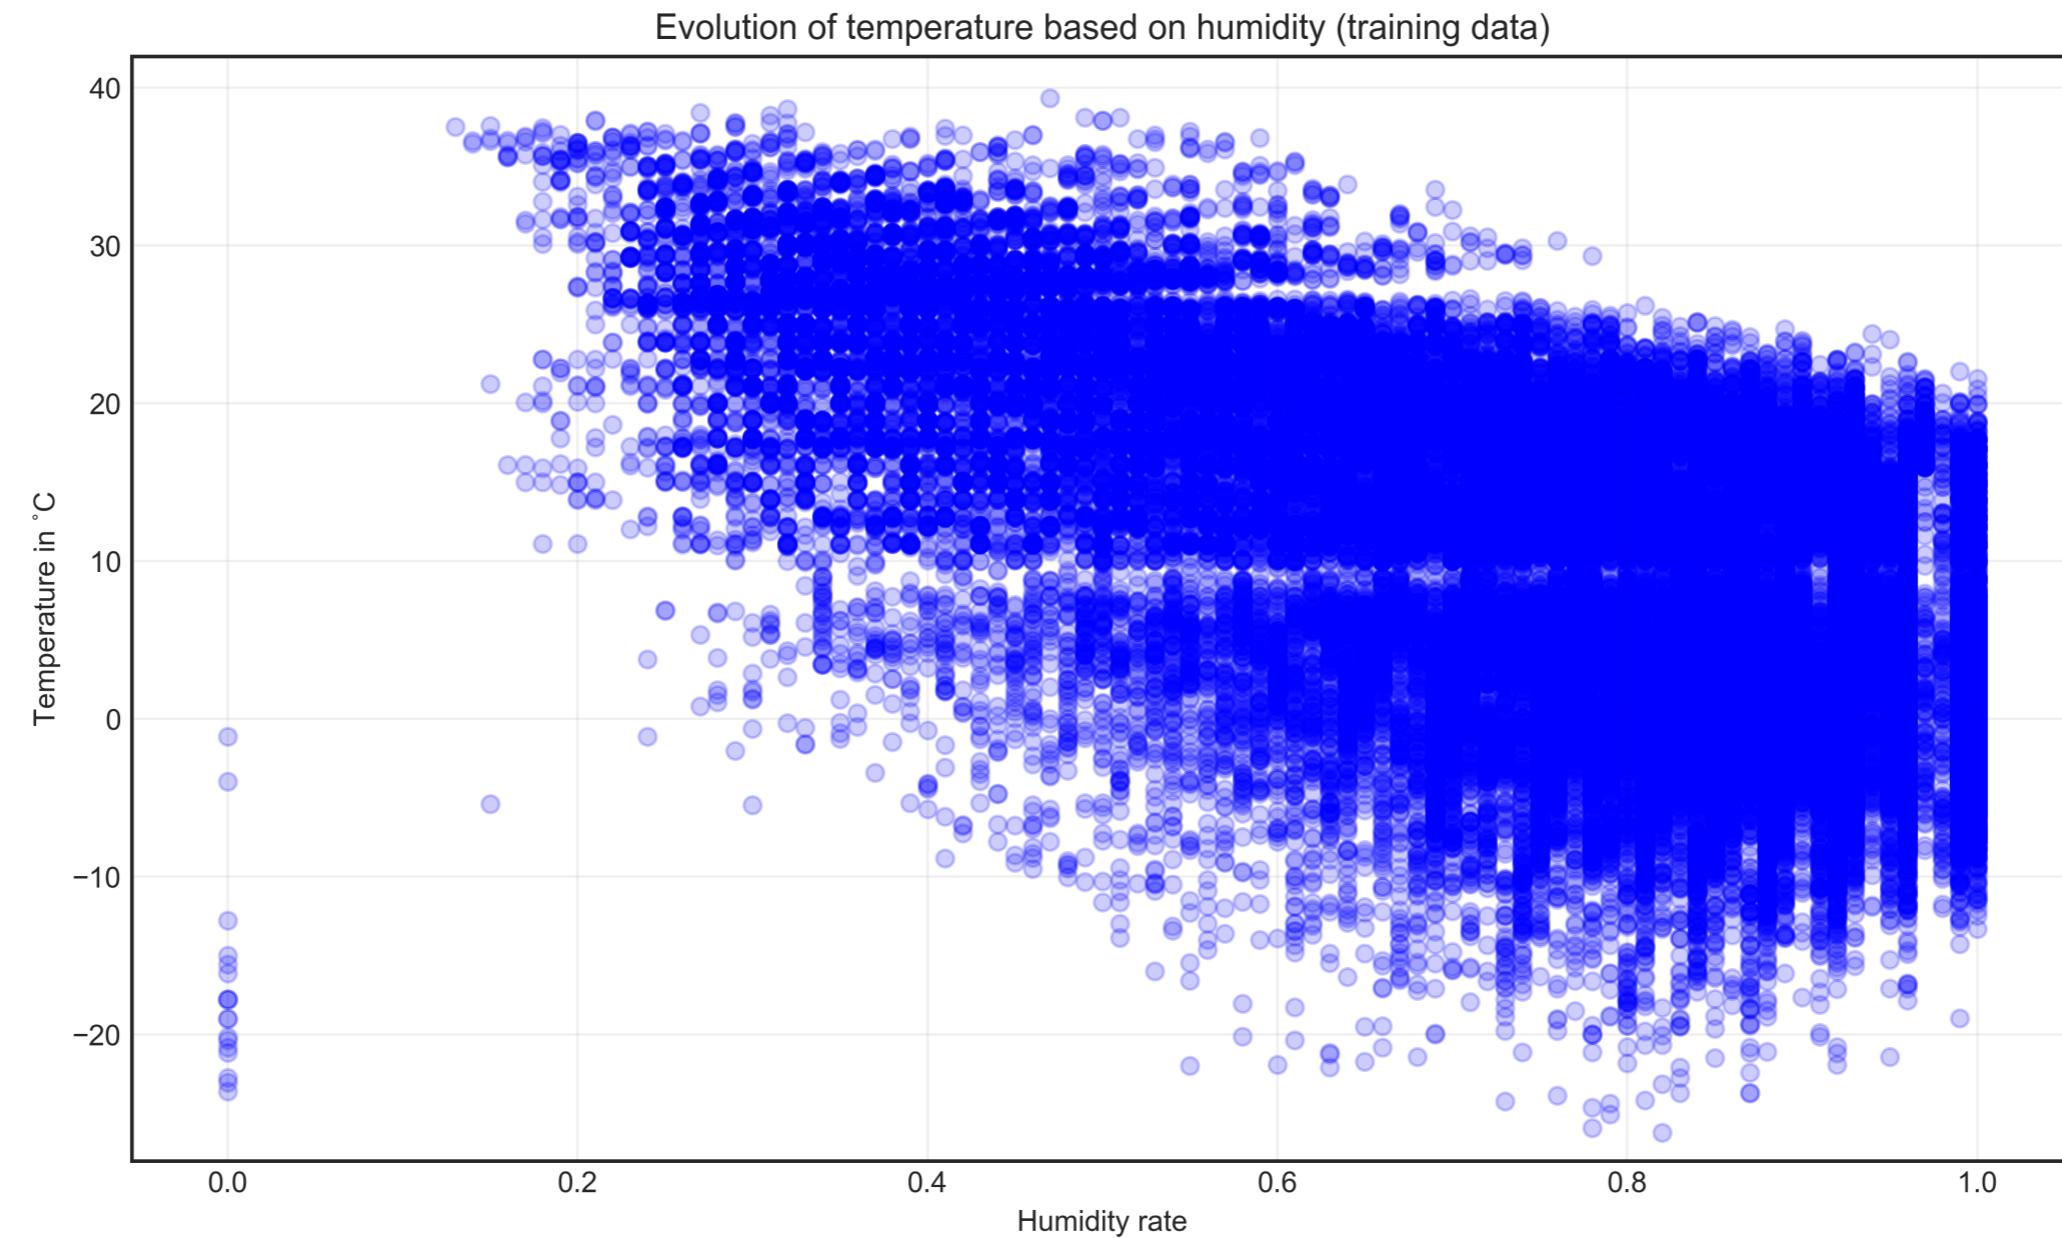
Regression

- **Regression** = assigning a **continuous** variable
 - How much will this stock be **worth**?
 - What is this exoplanet's **mass**?
 - How **tall** will this child be as an adult?

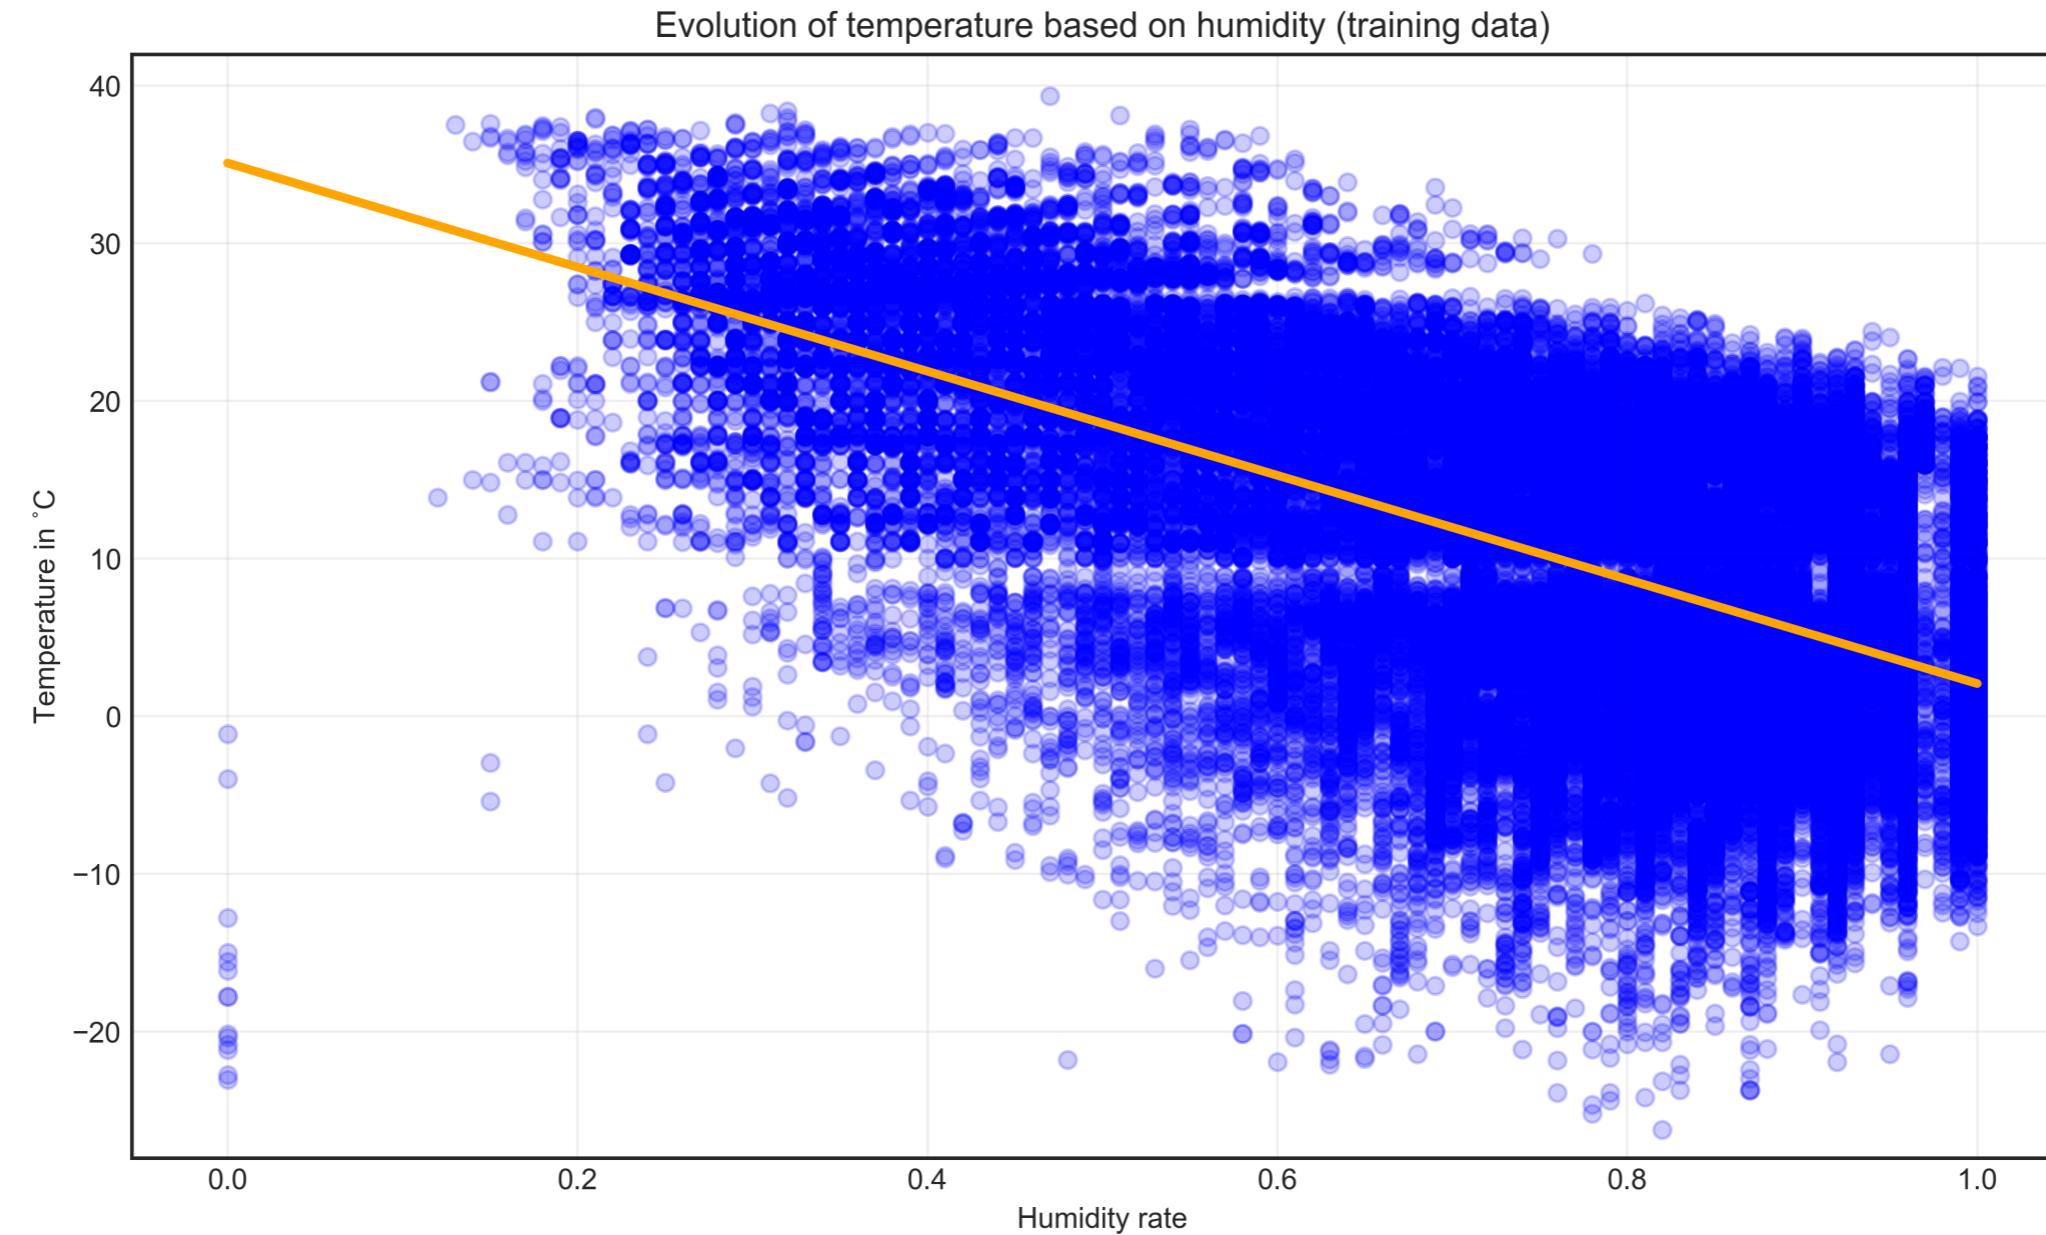
Predicting temperature

Reading ID	Humidity rate	Temperature in °C
0	0.89	7.388889
1	0.86	7.227778
2	0.89	9.377778
3	0.83	5.944444
...

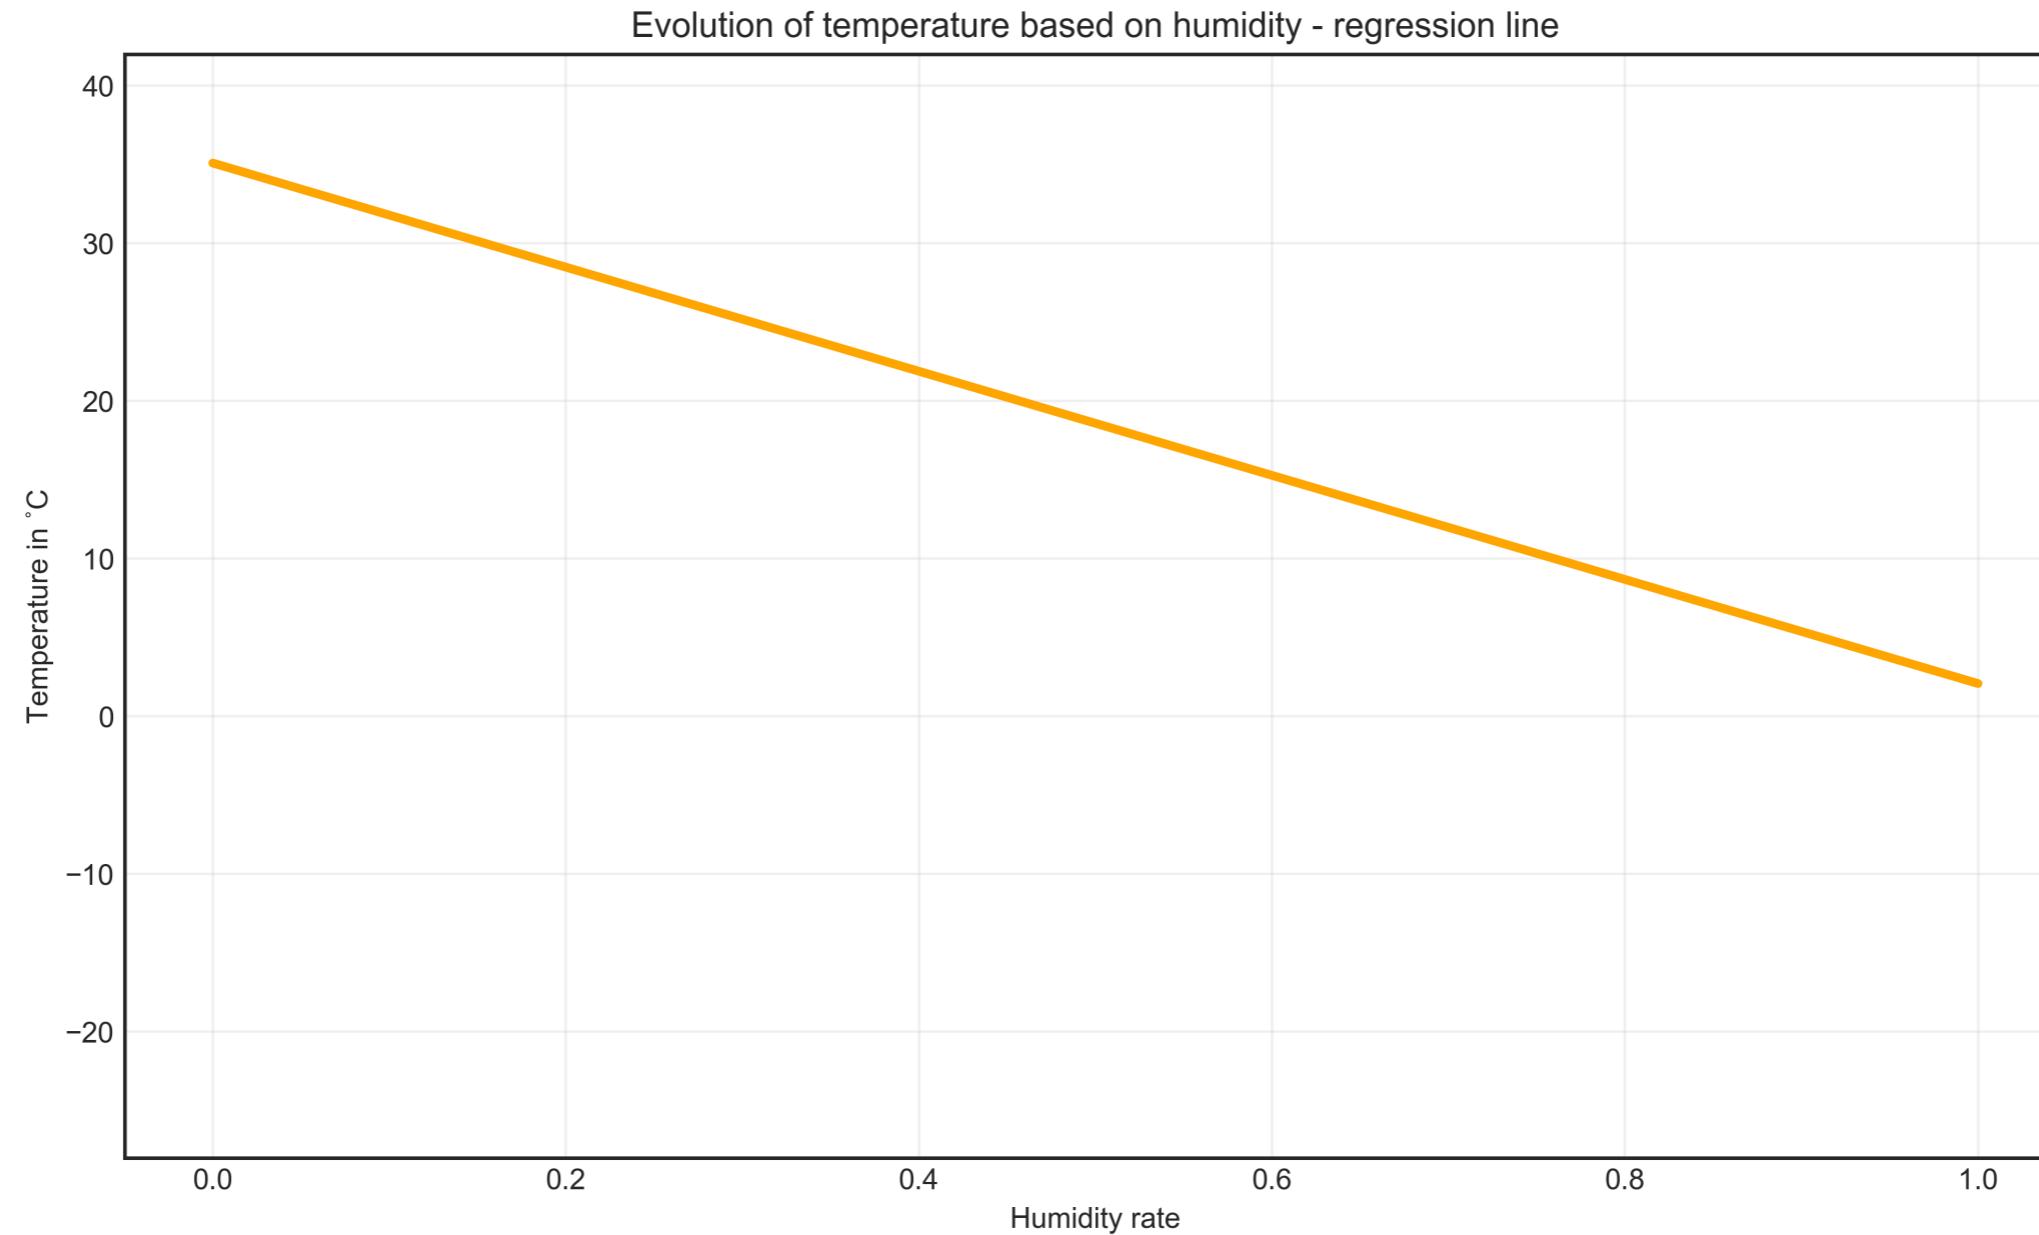
Training data



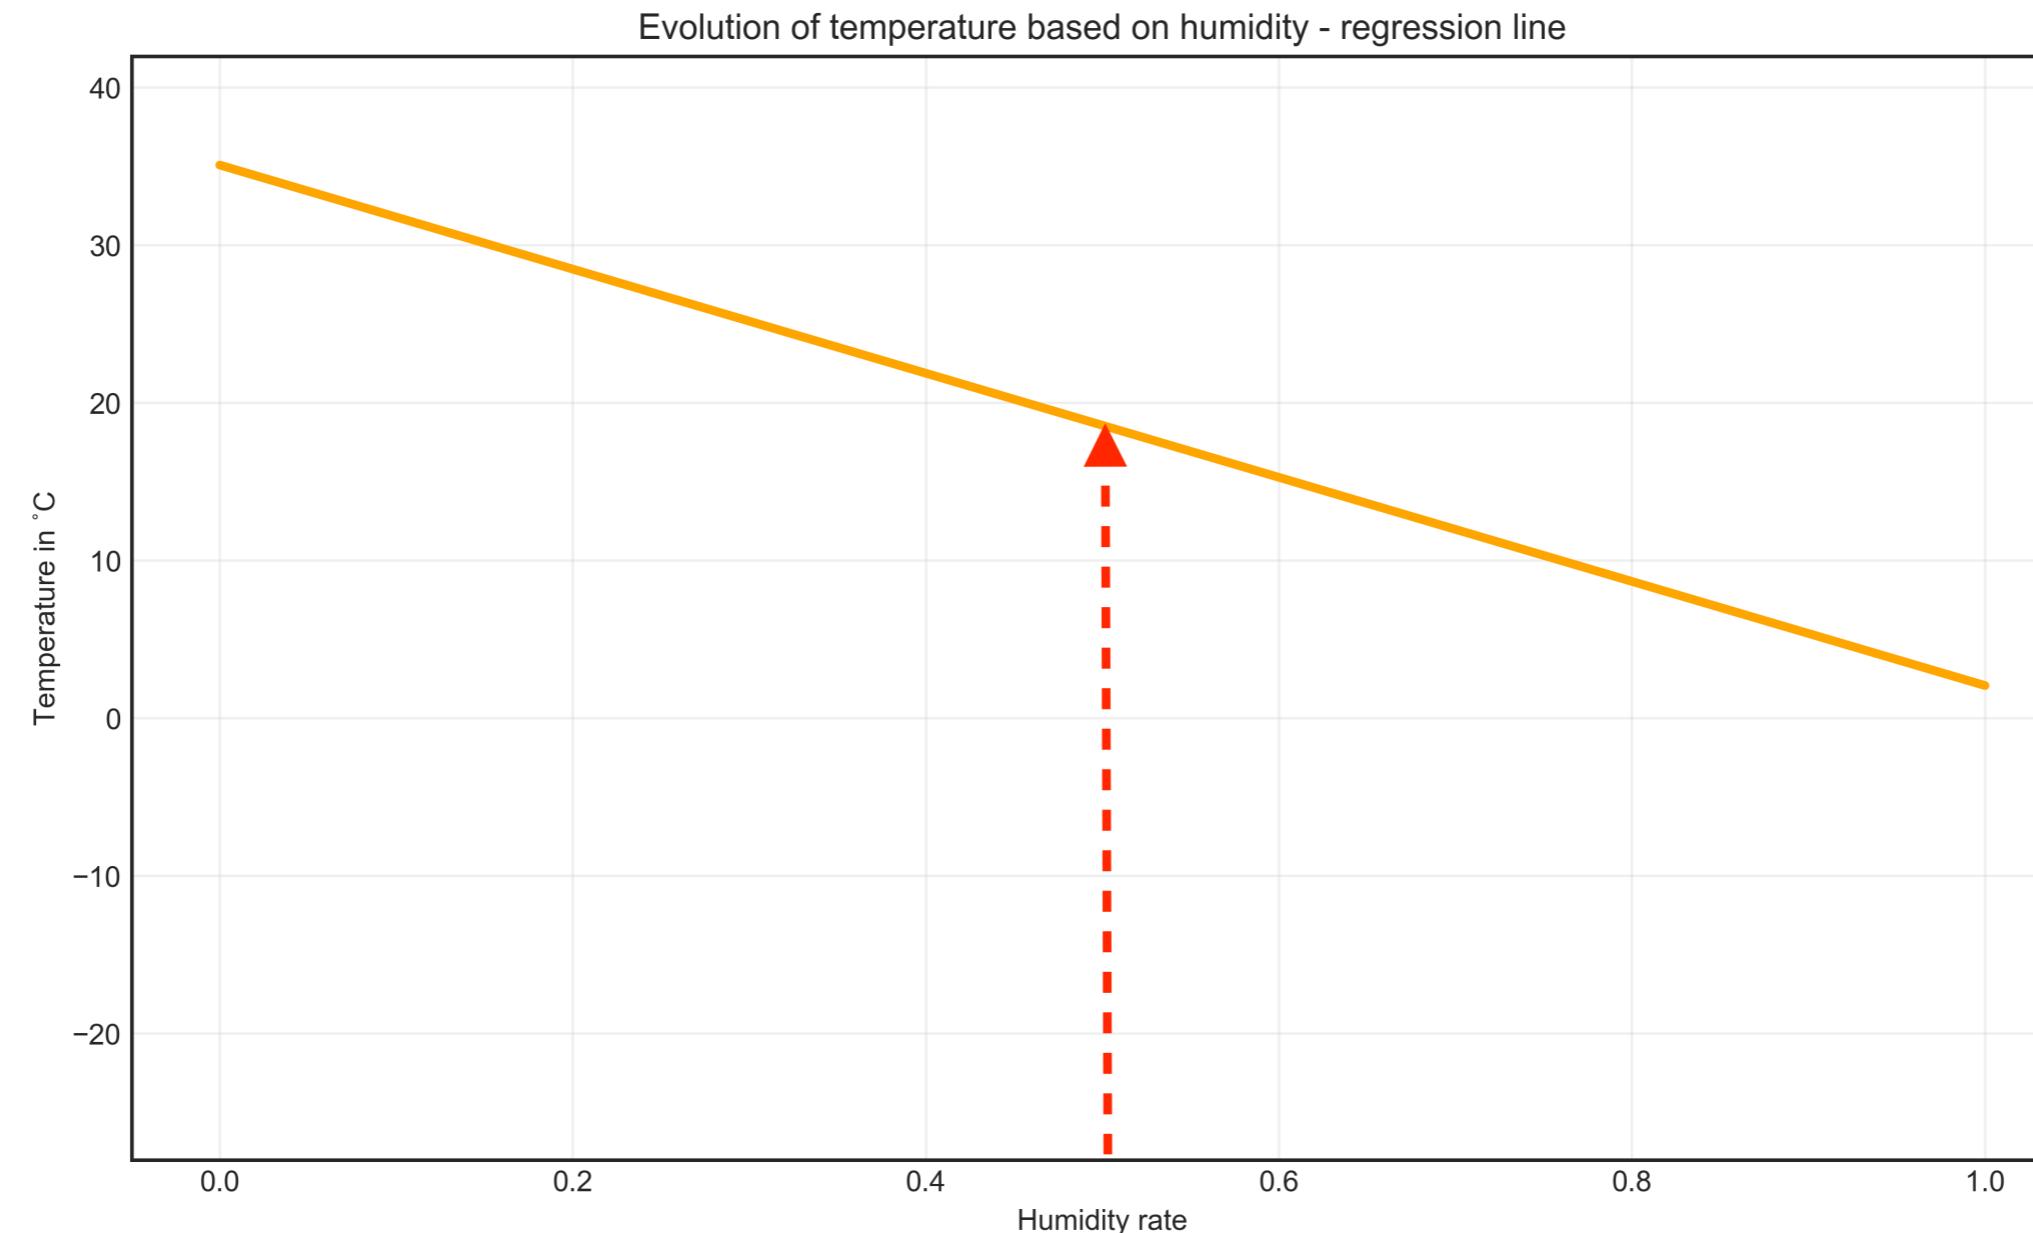
Linear regression



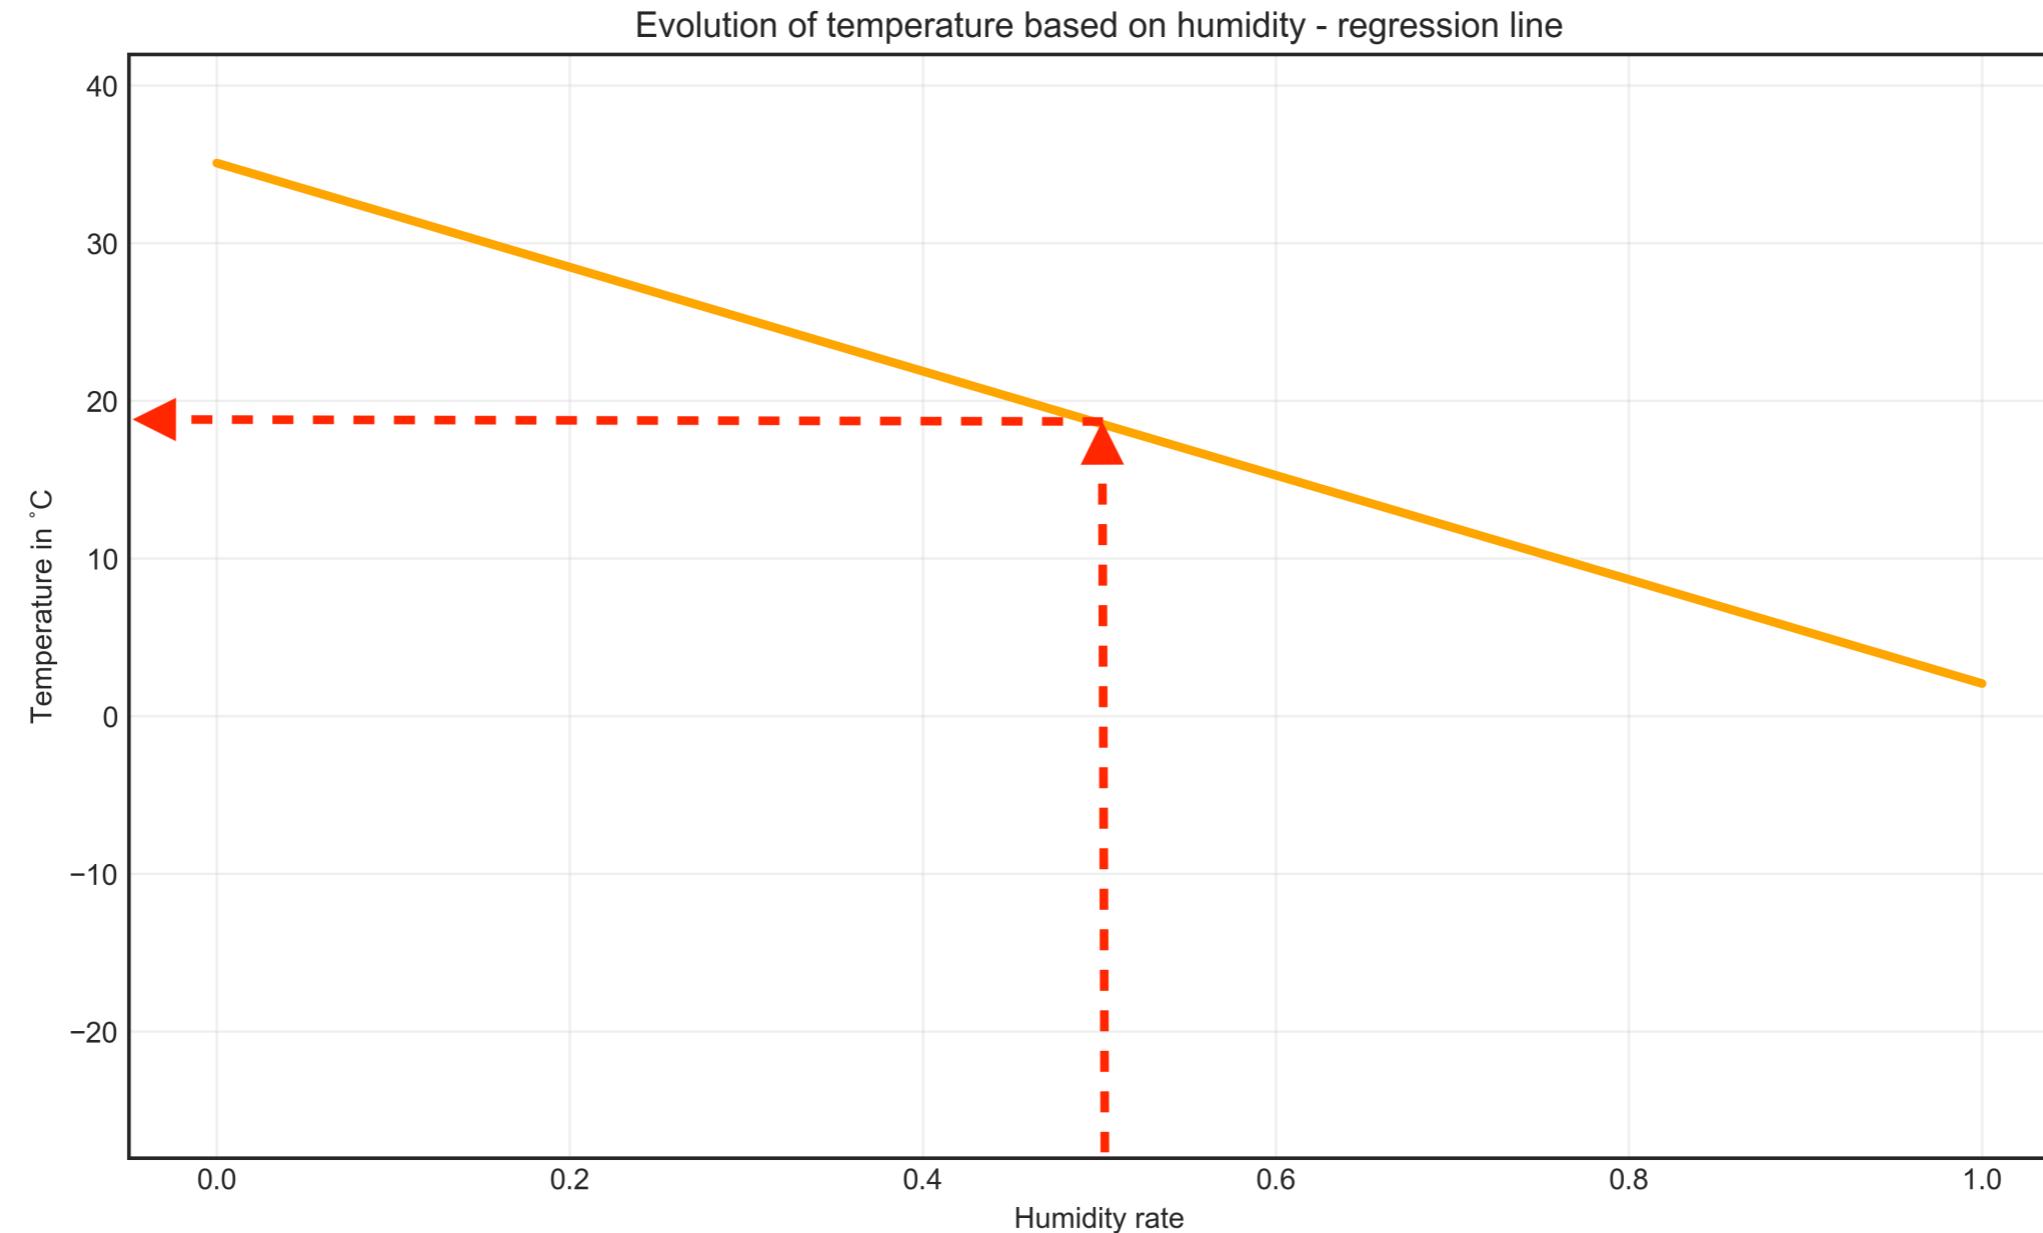
Model



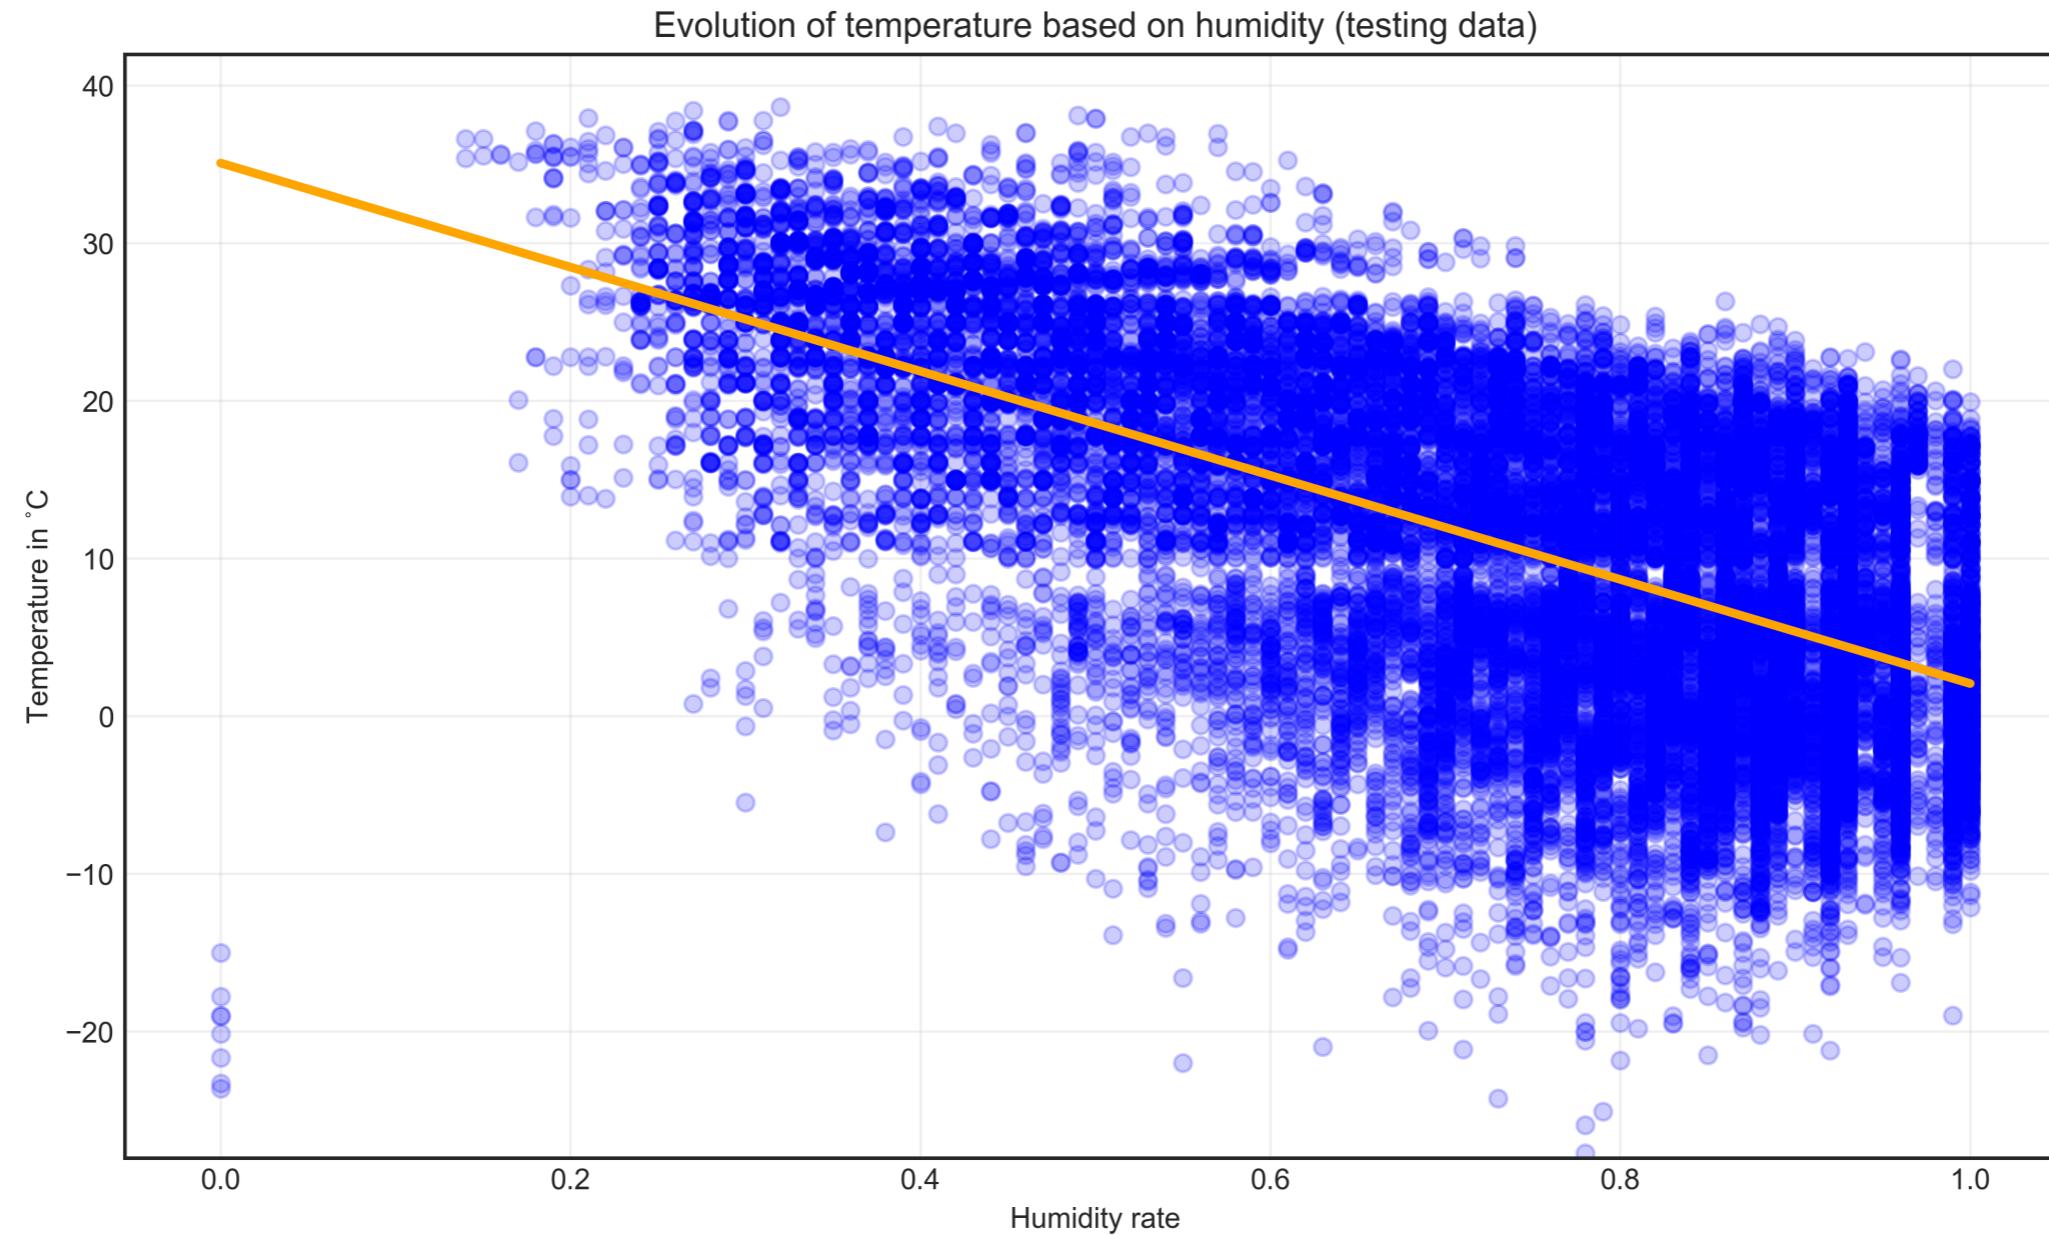
Given humidity...



...find temperature



Testing data



Classification vs regression

- Regression = **continuous**
 - **Any value** within a finite (*height*) or infinite (*time*) interval
 - $20^{\circ}\text{F}, 20.1^{\circ}\text{F}, 20.01^{\circ}\text{F}...$
- Classification = **category**
 - One of few **specific values**
 - *Cold, Mild, Hot*

Let's practice!

MACHINE LEARNING FOR EVERYONE

Unsupervised learning

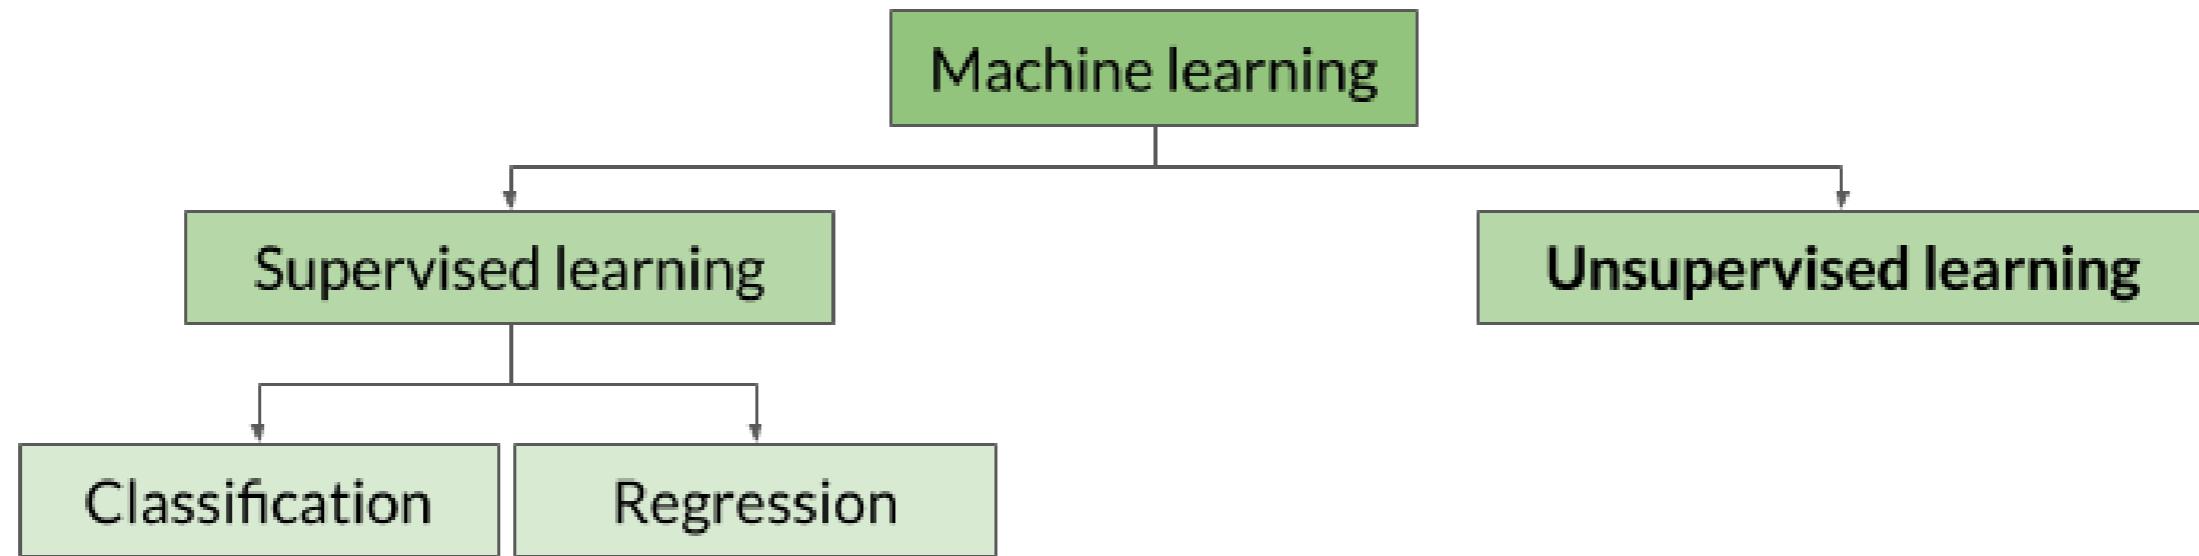
MACHINE LEARNING FOR EVERYONE



Hadrien Lacroix

Content Developer at DataCamp

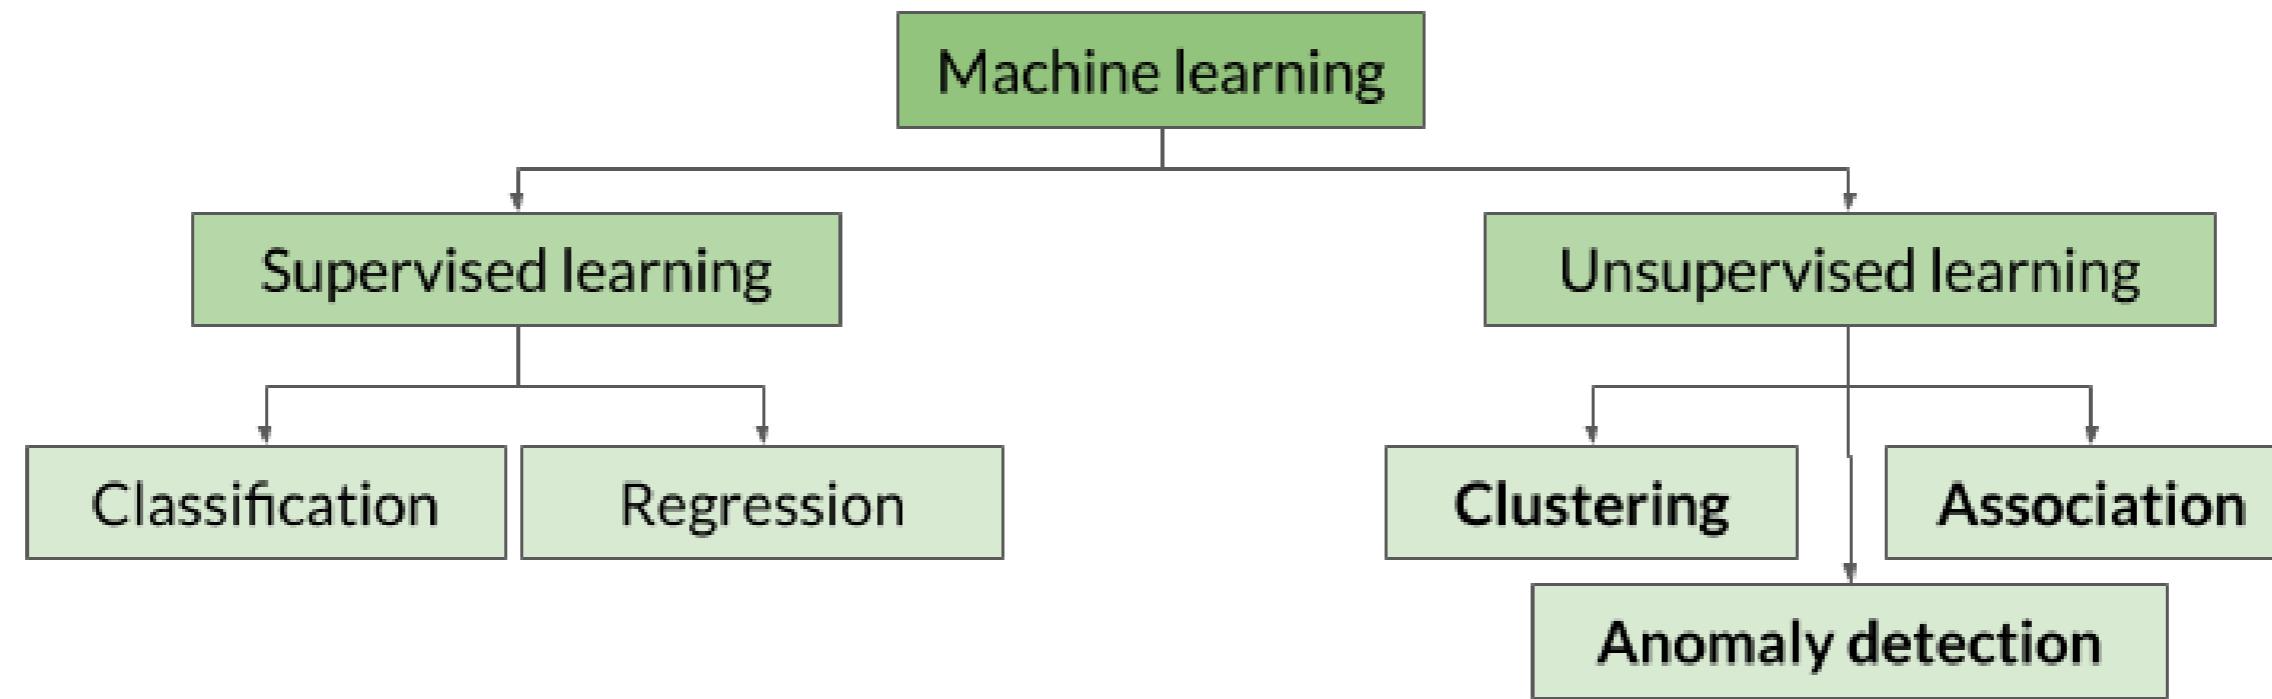
Unsupervised learning



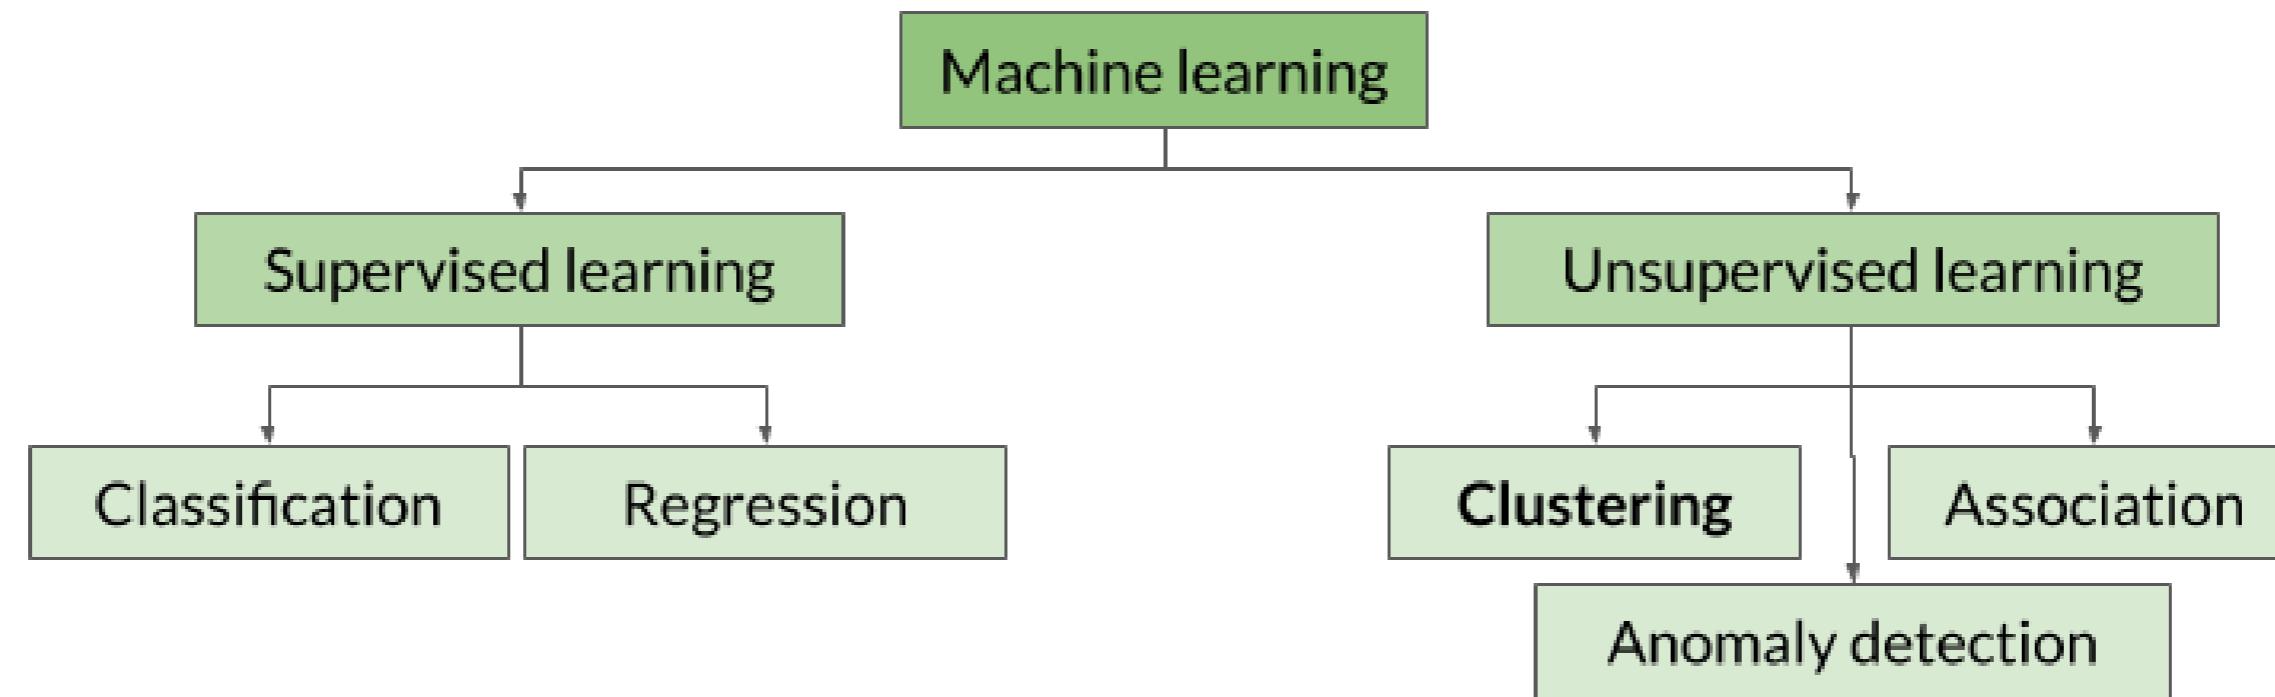
Unsupervised learning

- Unsupervised learning = **no target column**
 - No guidance
 - Looks at the whole dataset
 - Tries to detect patterns

Applications



Clustering



Clustering example

White Swiss Shepherd



Brown Japanese Bobtail



Brown Akita



Black Norwegian Forest



Black German Shepherd

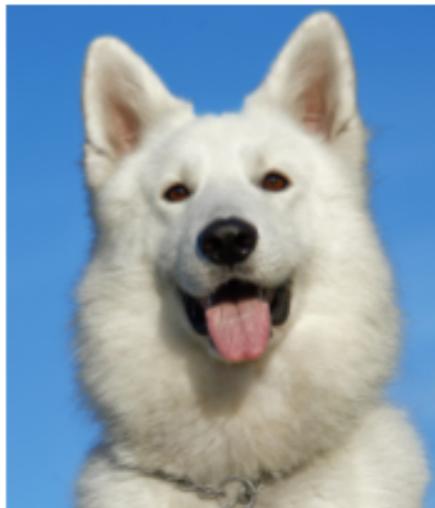


Grey Kurilian Bobtail

Species cluster

DOGS

White Swiss Shepherd



Black German Shepherd



Brown Akita



CATS



Black Norwegian Forest



Brown Japanese Bobtail



Grey Kurilian Bobtail

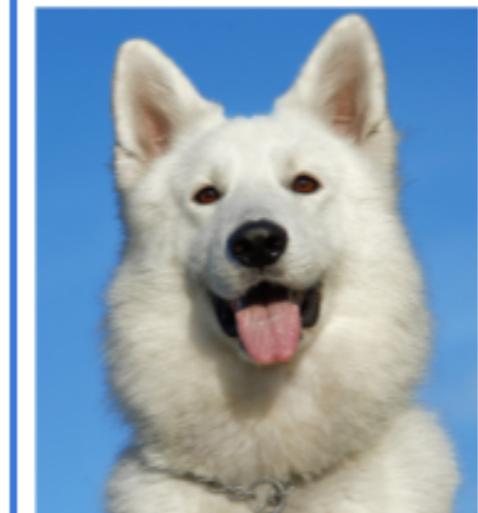
Color cluster

BLACK

Black Norwegian Forest Black German Shepherd



WHITE



White Swiss Shepherd

Grey Kurilian Bobtail



GREY

Brown Japanese Bobtail



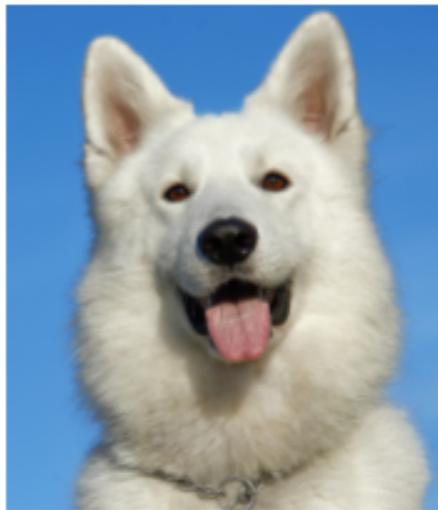
Brown Akita

BROWN

Origin cluster



White Swiss Shepherd



Black German Shepherd



Black Norwegian Forest



Brown Akita



Brown Japanese Bobtail



Grey Kurilian Bobtail

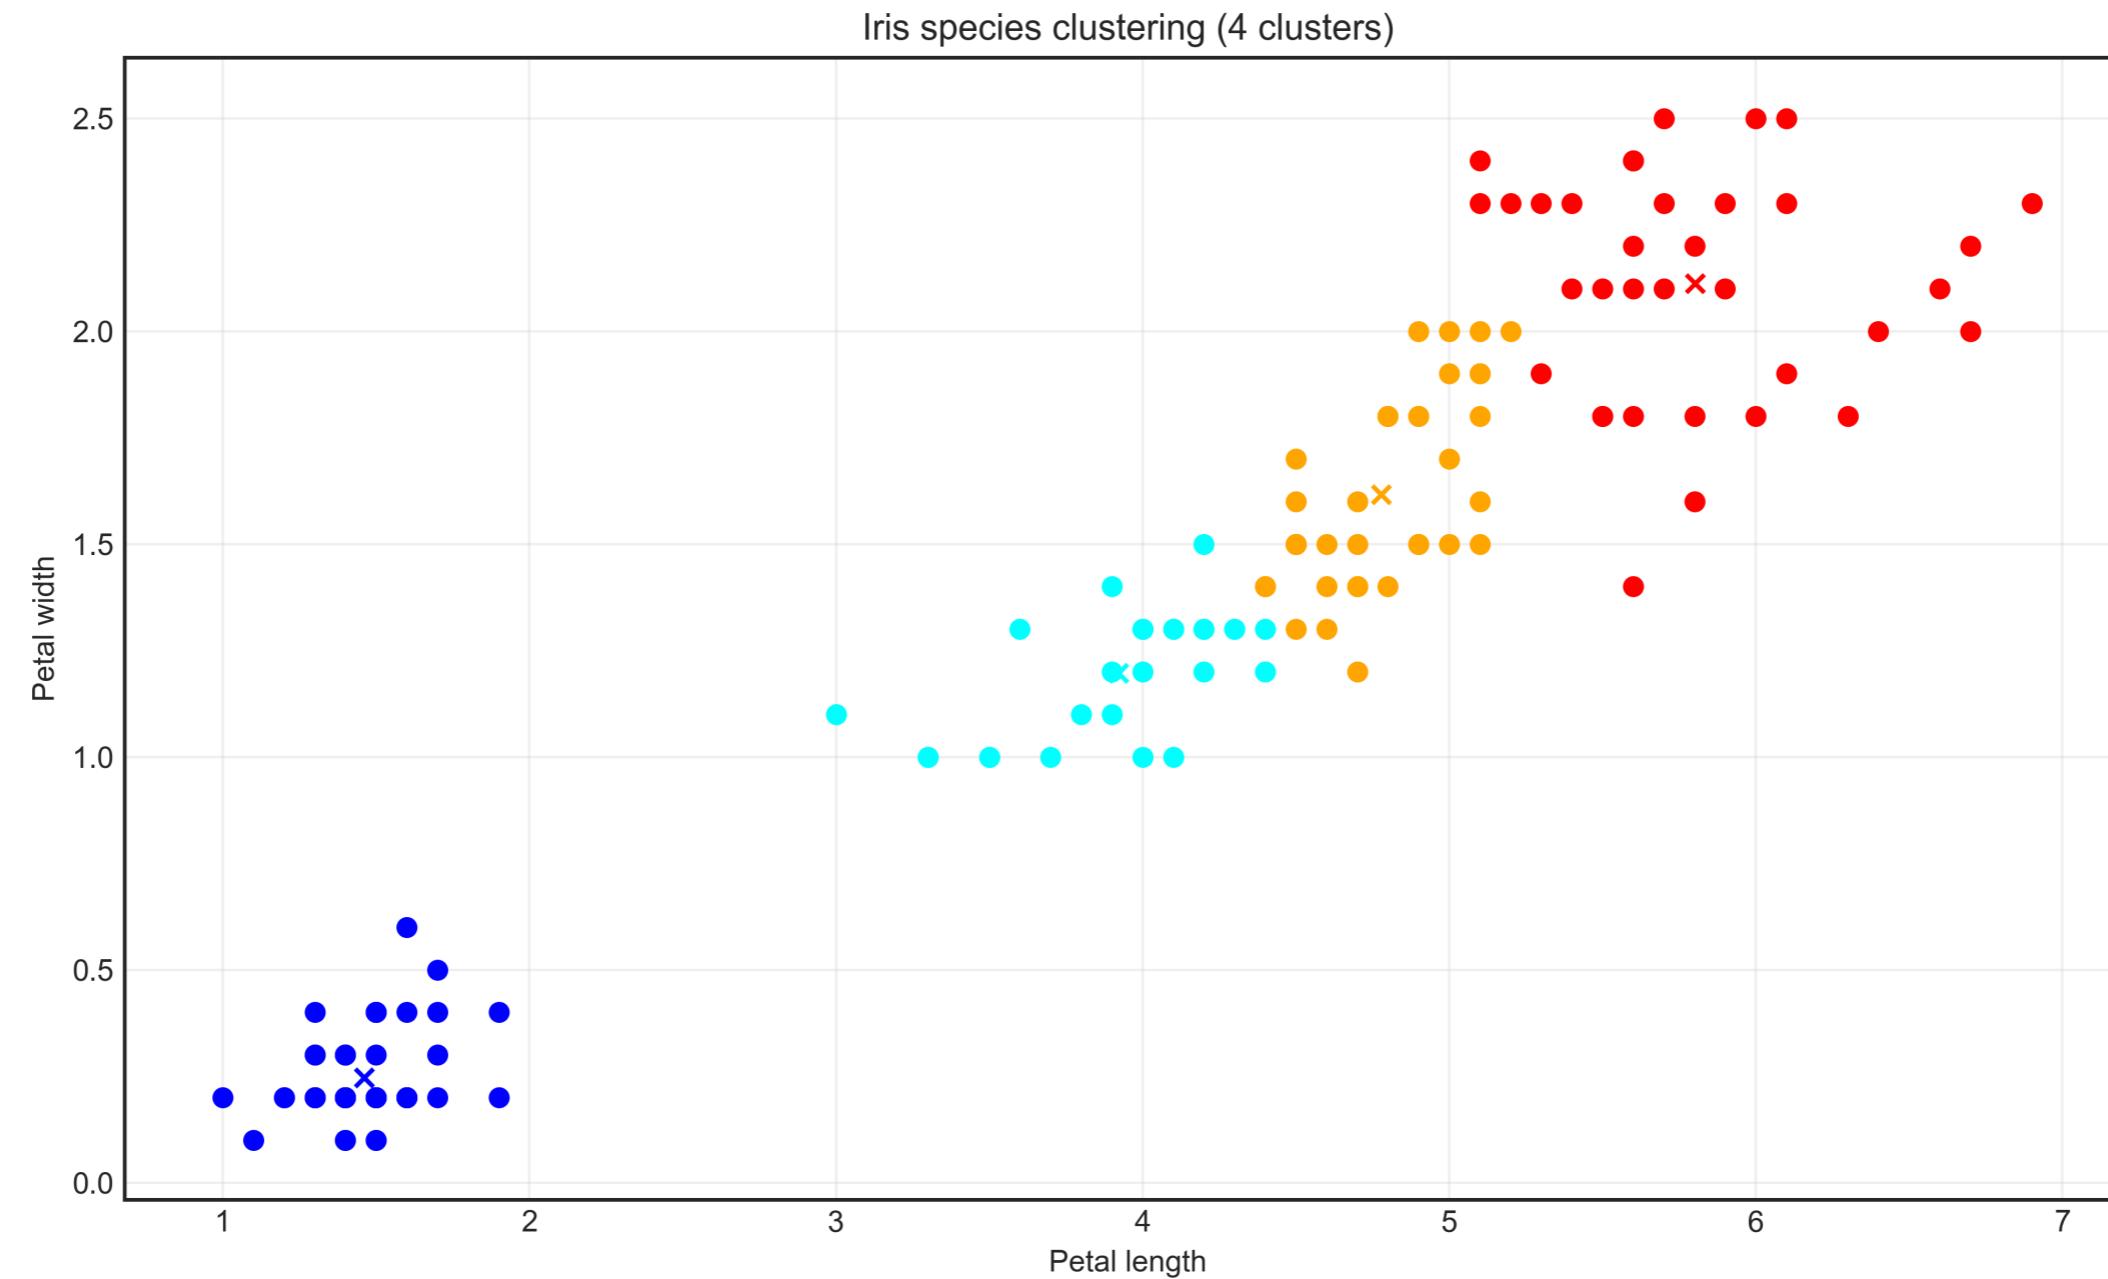
Clustering models

- **K Means:**
 - Specify the **number of clusters**
- **DBSCAN** (density-based spatial clustering of applications with noise):
 - Specify **what constitutes a cluster**

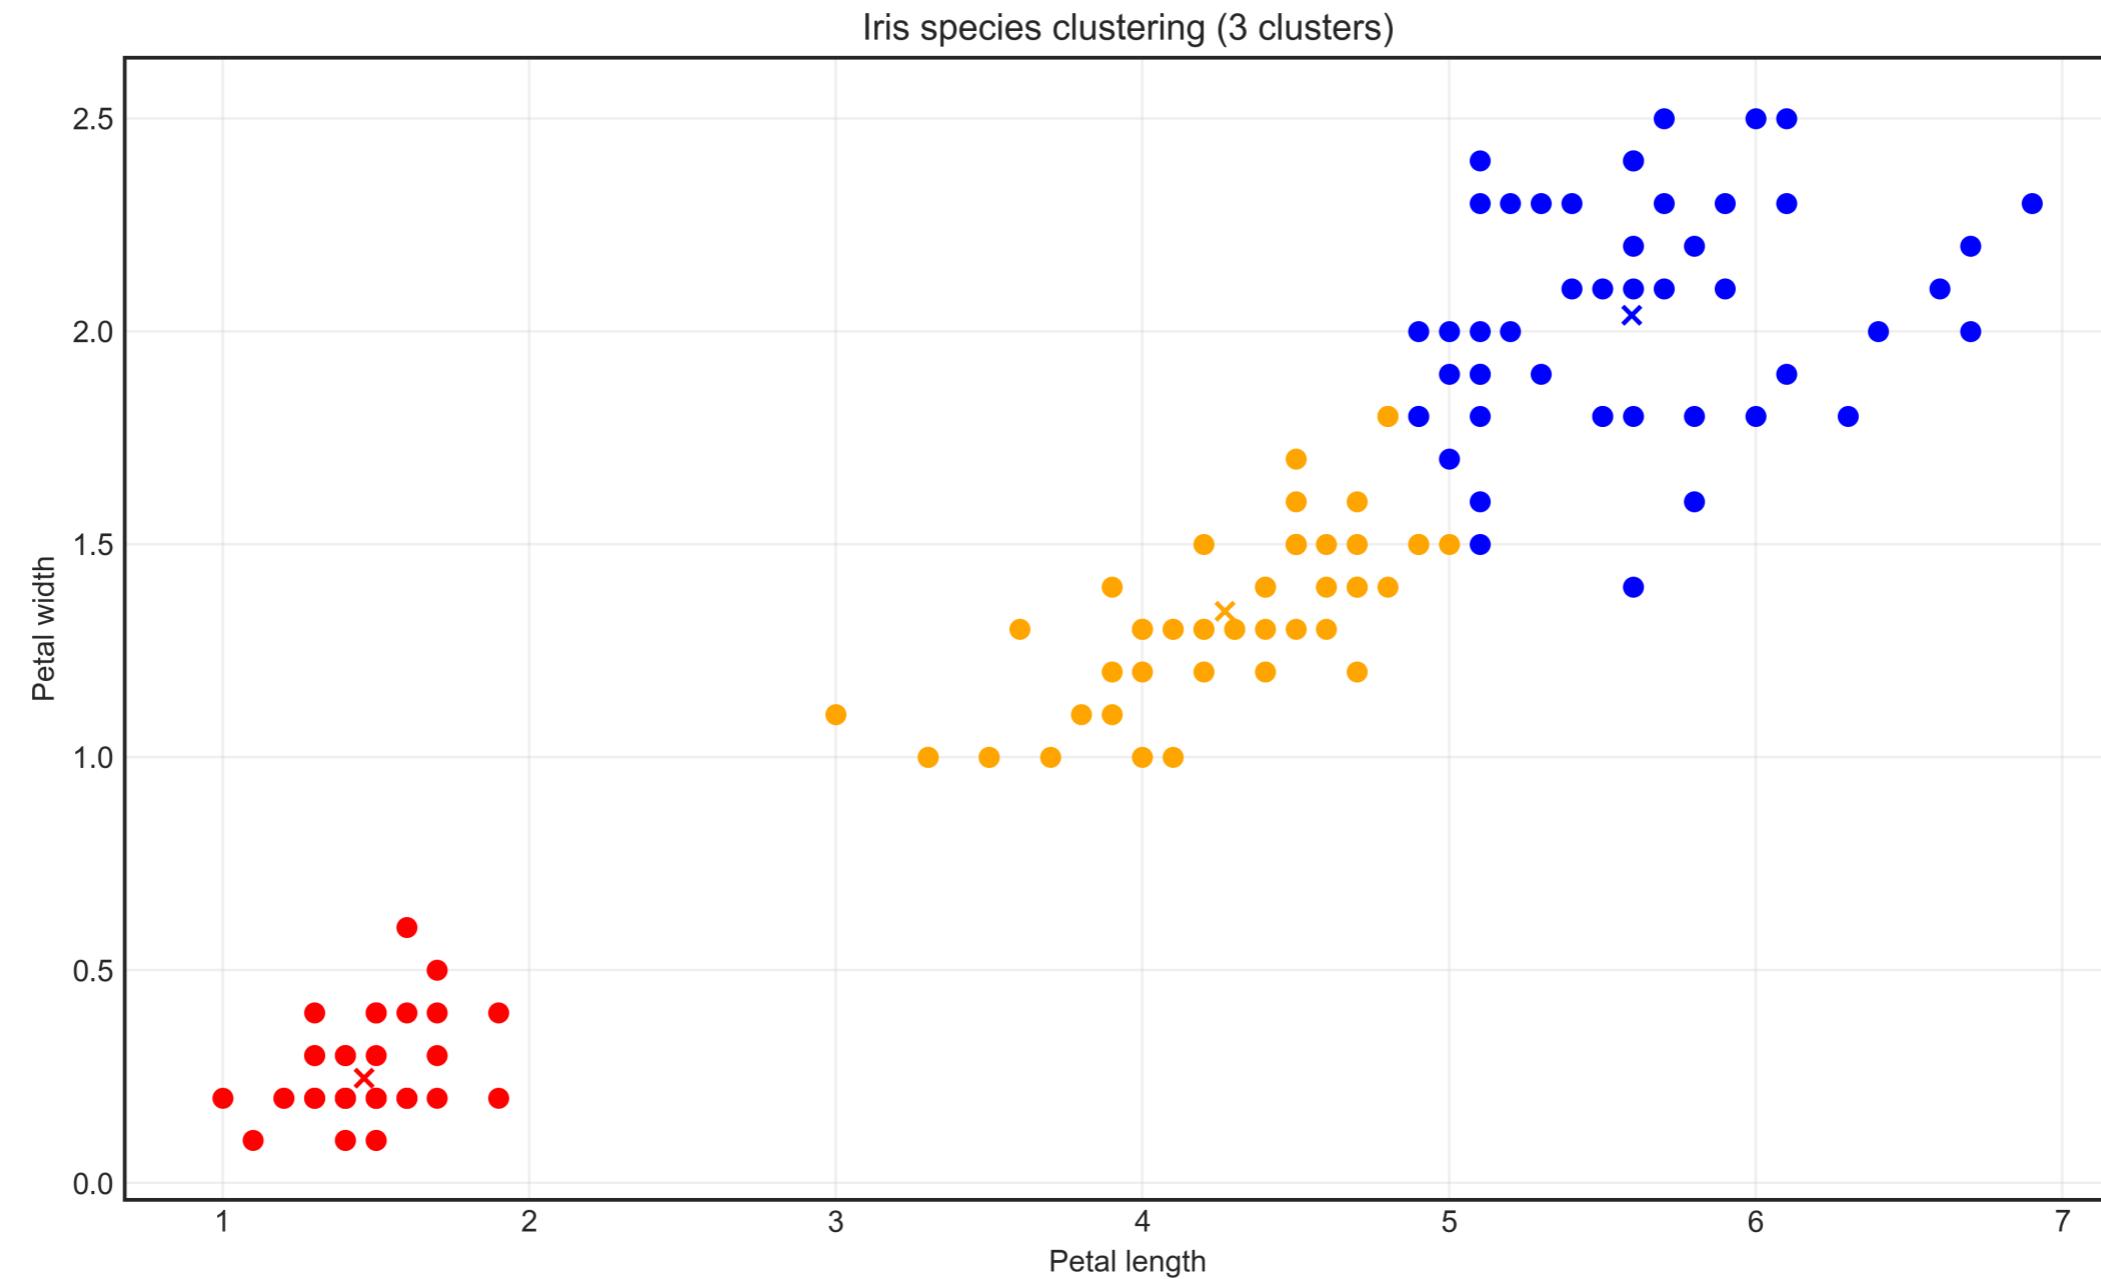
Iris table

	Petal length	Petal width
0	1.4	0.2
1	1.4	0.2
2	1.3	0.2
3	5.1	1.9
...

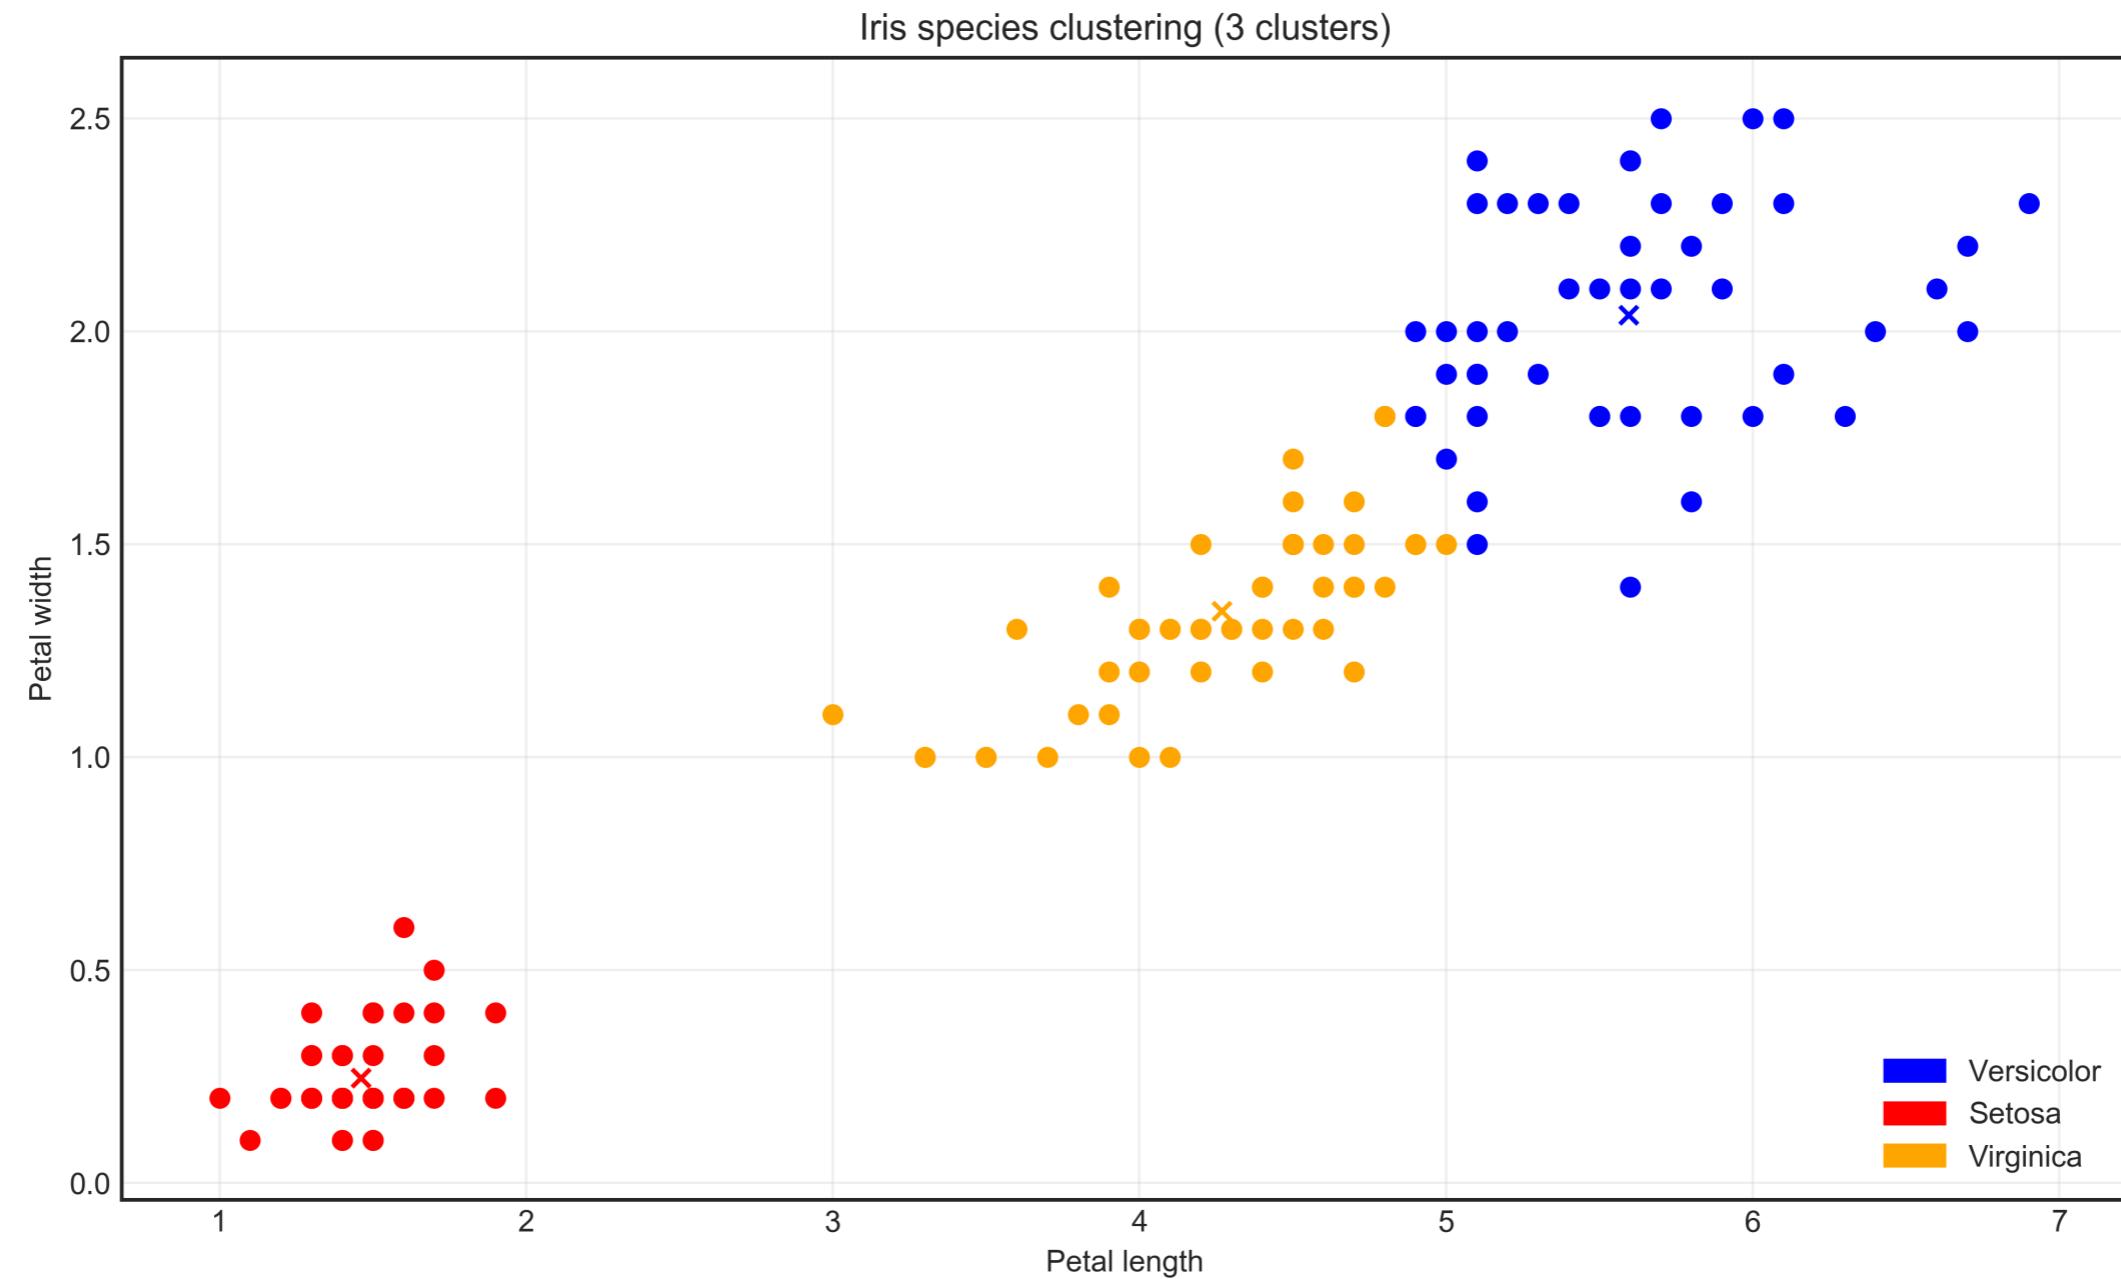
K-Means with 4 clusters



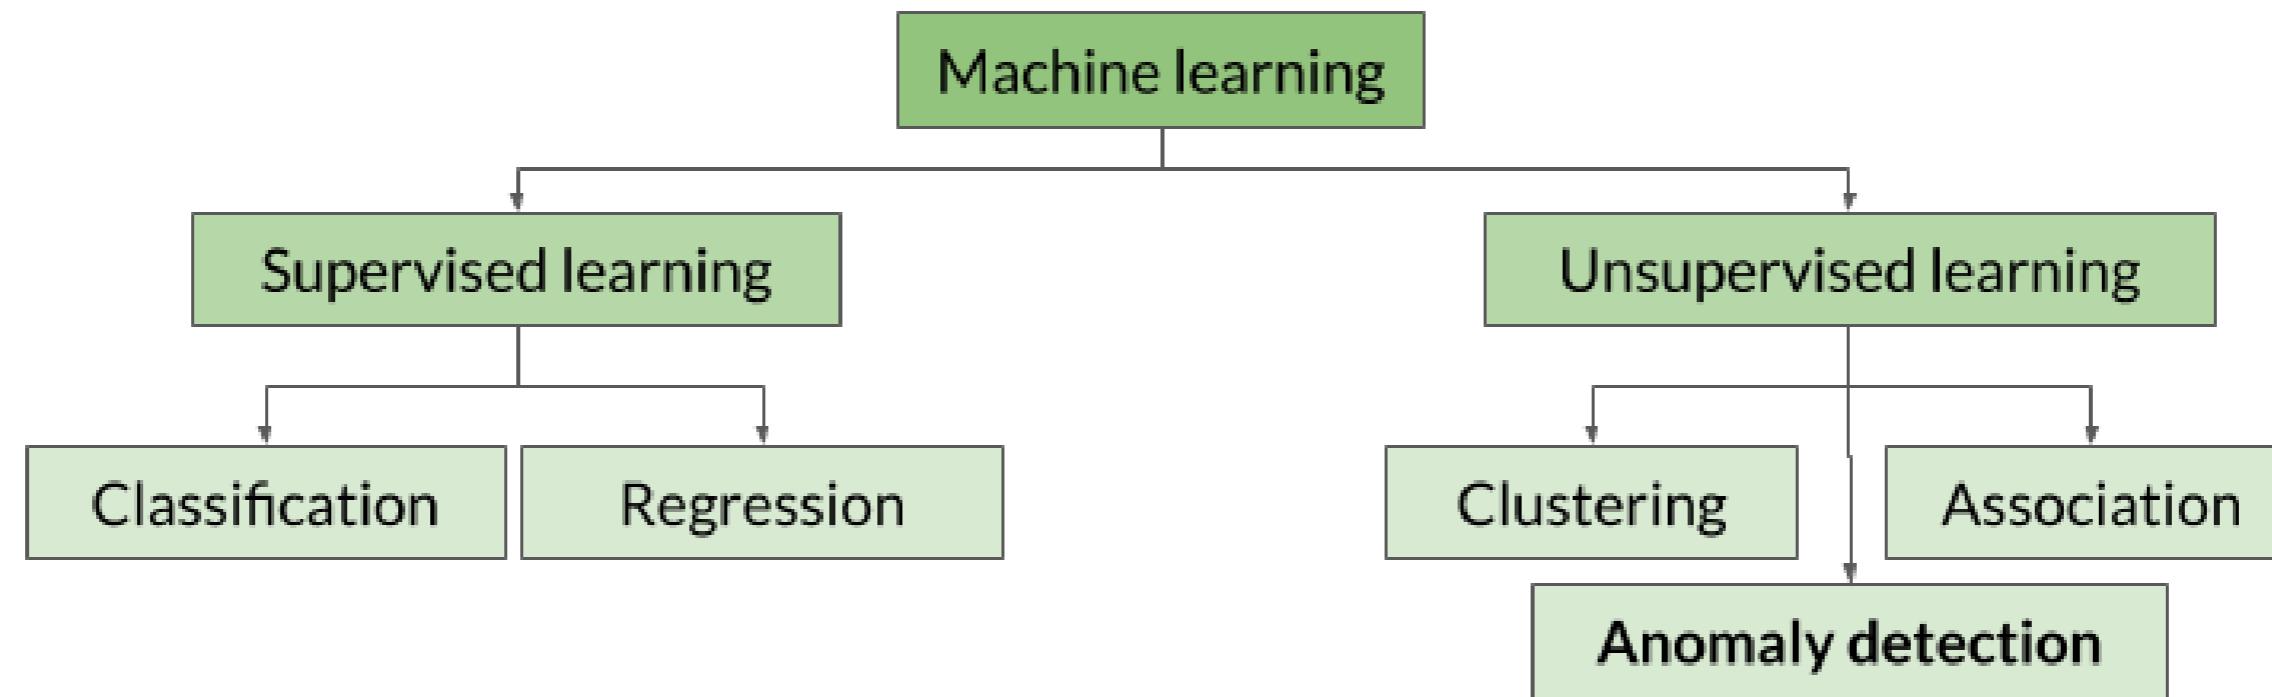
K-Means with 3 clusters



Ground truth



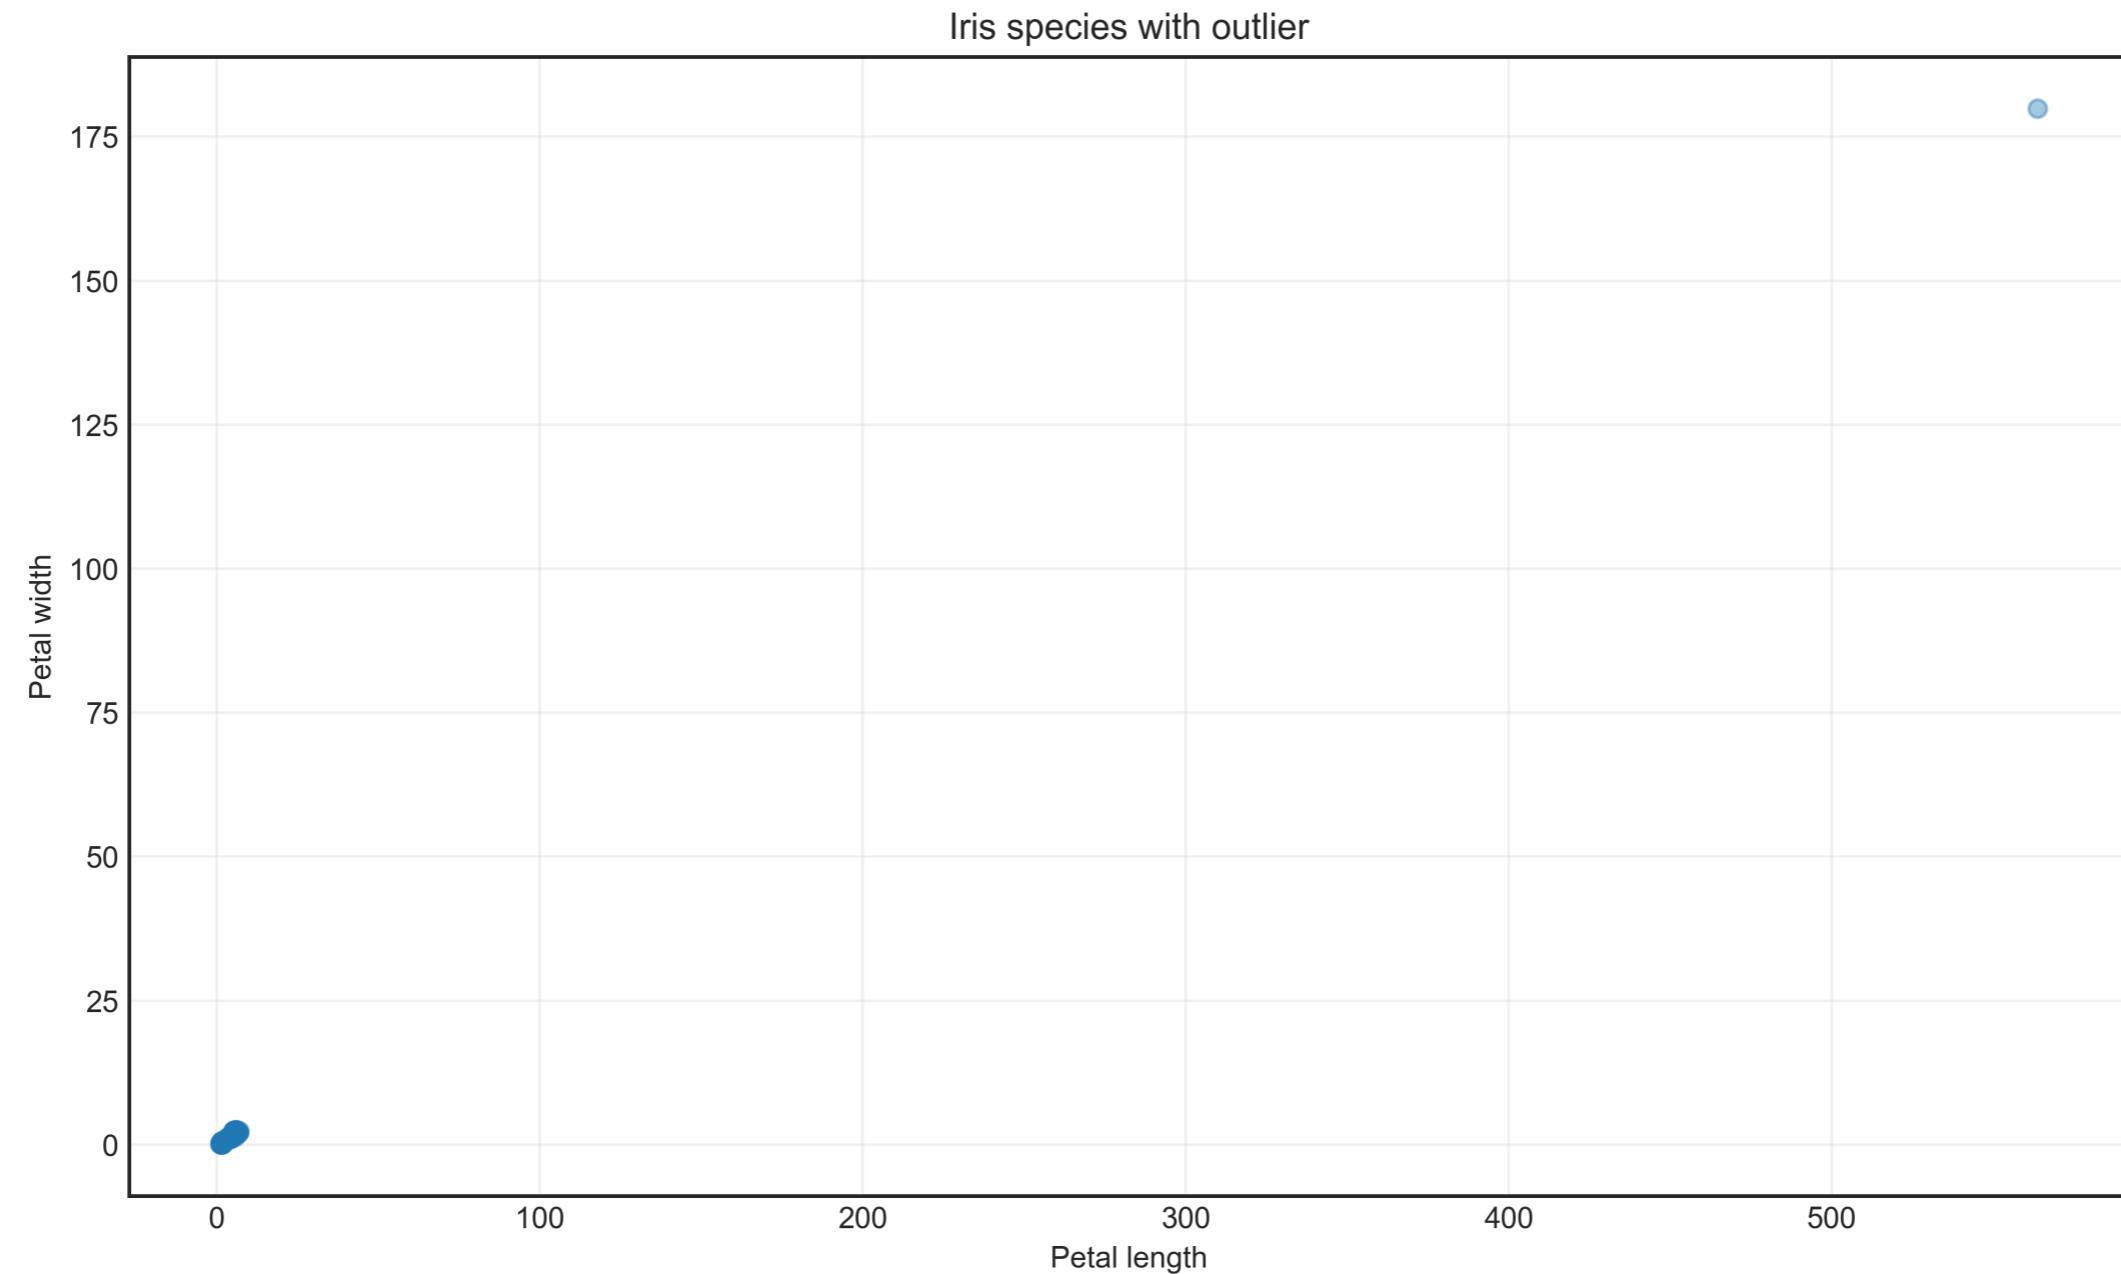
Anomaly detection



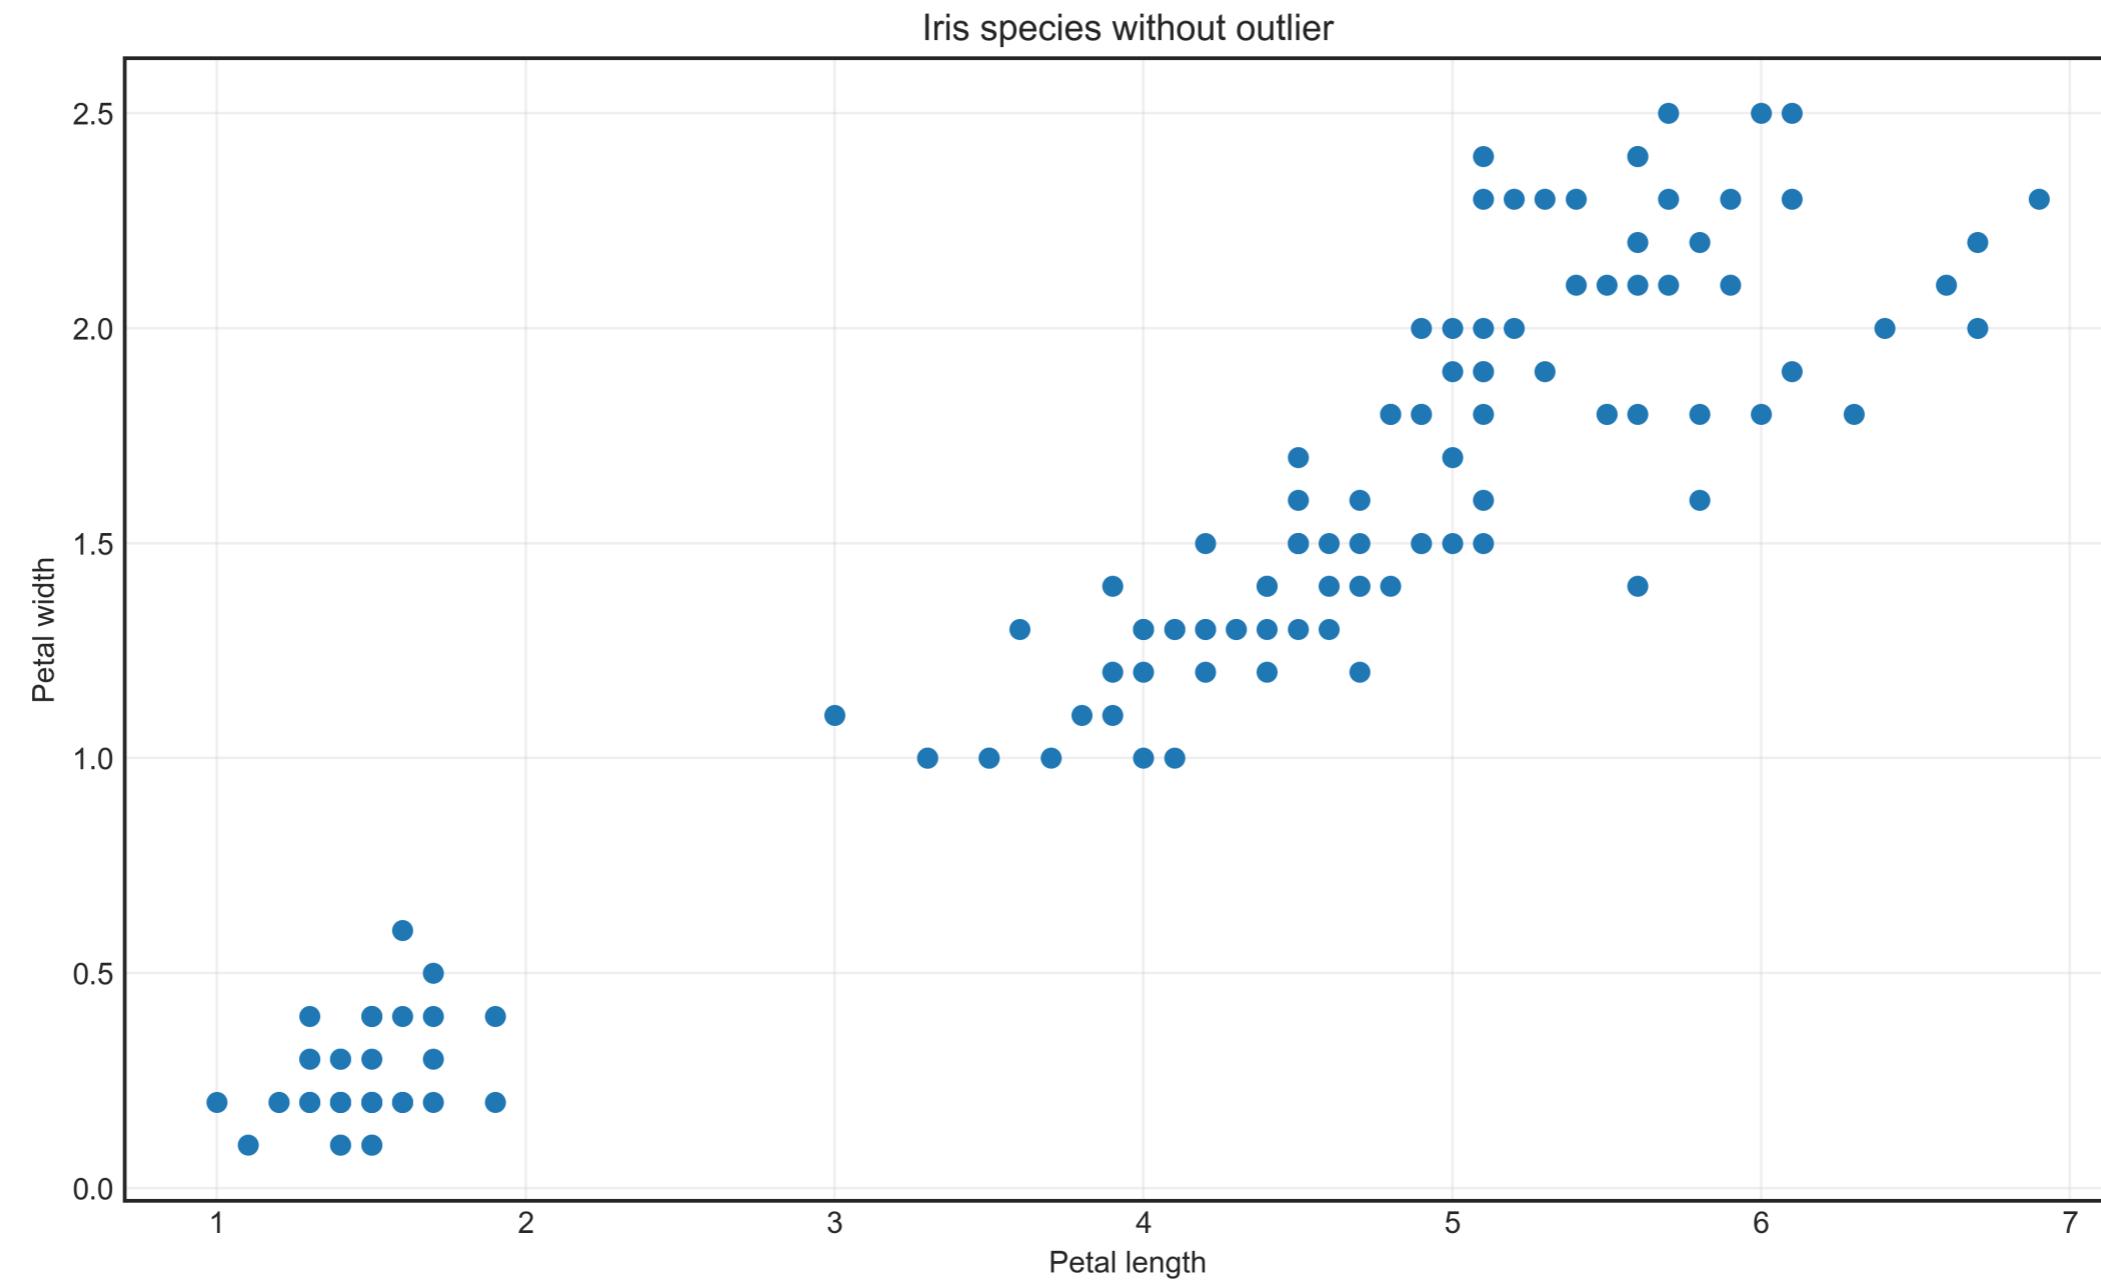
Detecting outliers

- Anomaly detection = **detecting outliers**
- Outliers = observations that **differ from the rest**

Outliers



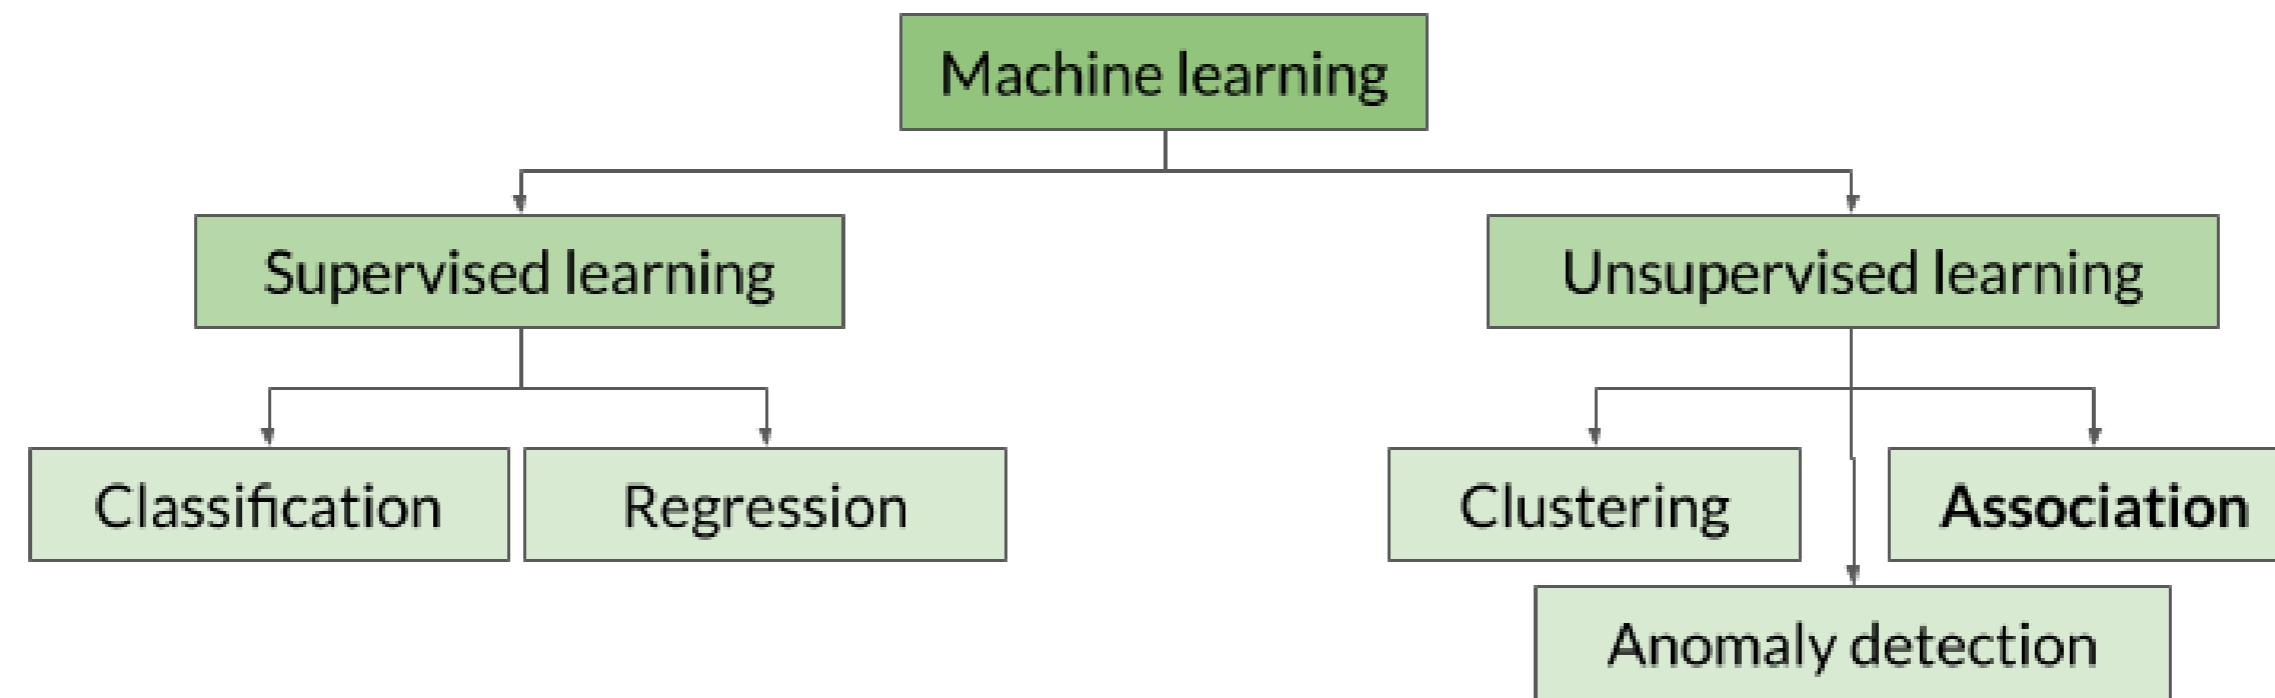
Removing outliers



Some anomaly detection use cases

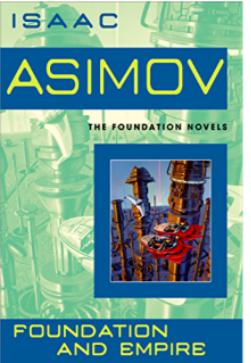
- Discover devices that fail faster or last longer
- Discover fraudsters that manage trick the system
- Discover patients that resist a fatal disease
- ...

Association



Association

Customers who bought this item also bought



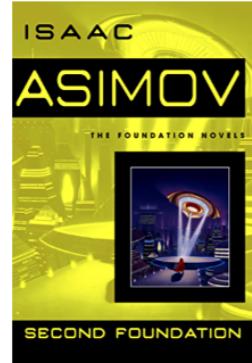
Foundation and Empire

› Isaac Asimov

★★★★★ 335

Kindle Edition

1 offer from \$5.99



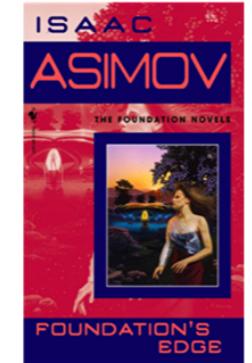
Second Foundation

› Isaac Asimov

★★★★★ 253

Kindle Edition

1 offer from \$6.99



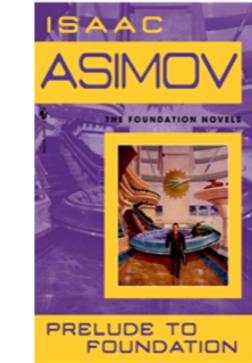
Foundation's Edge

› Isaac Asimov

★★★★★ 277

Kindle Edition

1 offer from \$6.99



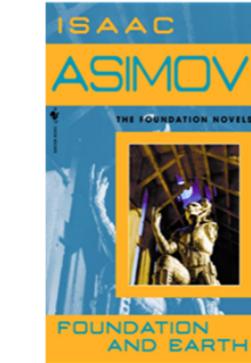
Prelude to Foundation

› Isaac Asimov

★★★★★ 371

Kindle Edition

1 offer from \$8.99



Foundation and Earth

› Isaac Asimov

★★★★★ 322

Kindle Edition

1 offer from \$6.99



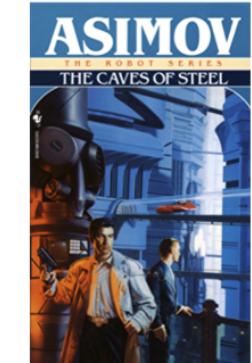
I, Robot (The Robot Series)

› Isaac Asimov

★★★★★ 714

Kindle Edition

1 offer from \$7.99



The Caves of Steel (The Robot Series Book 1)

› Isaac Asimov

★★★★★ 557

Kindle Edition

1 offer from \$7.99

Let's practice!

MACHINE LEARNING FOR EVERYONE

Evaluating performance

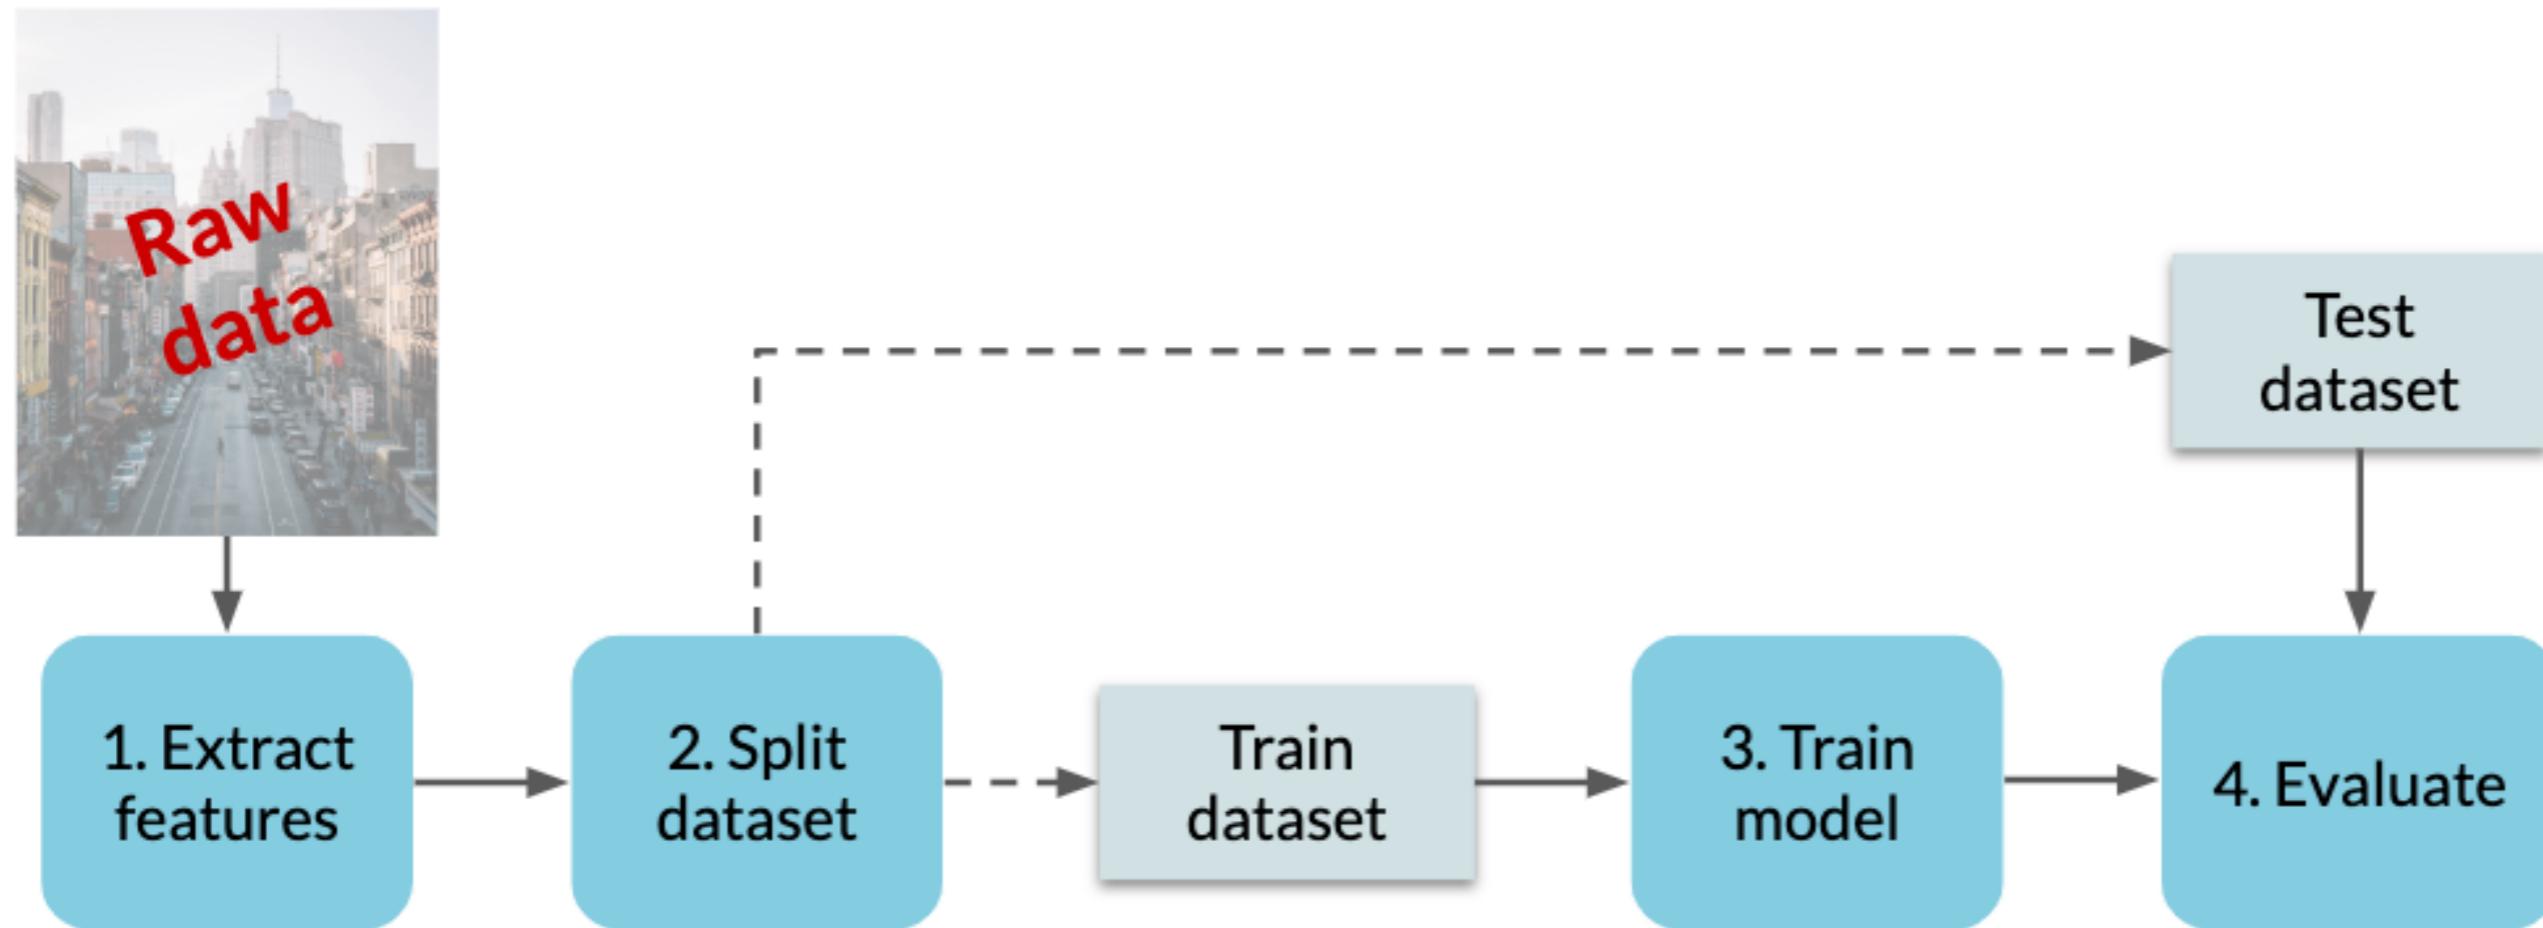
MACHINE LEARNING FOR EVERYONE



Hadrien Lacroix

Content Developer at DataCamp

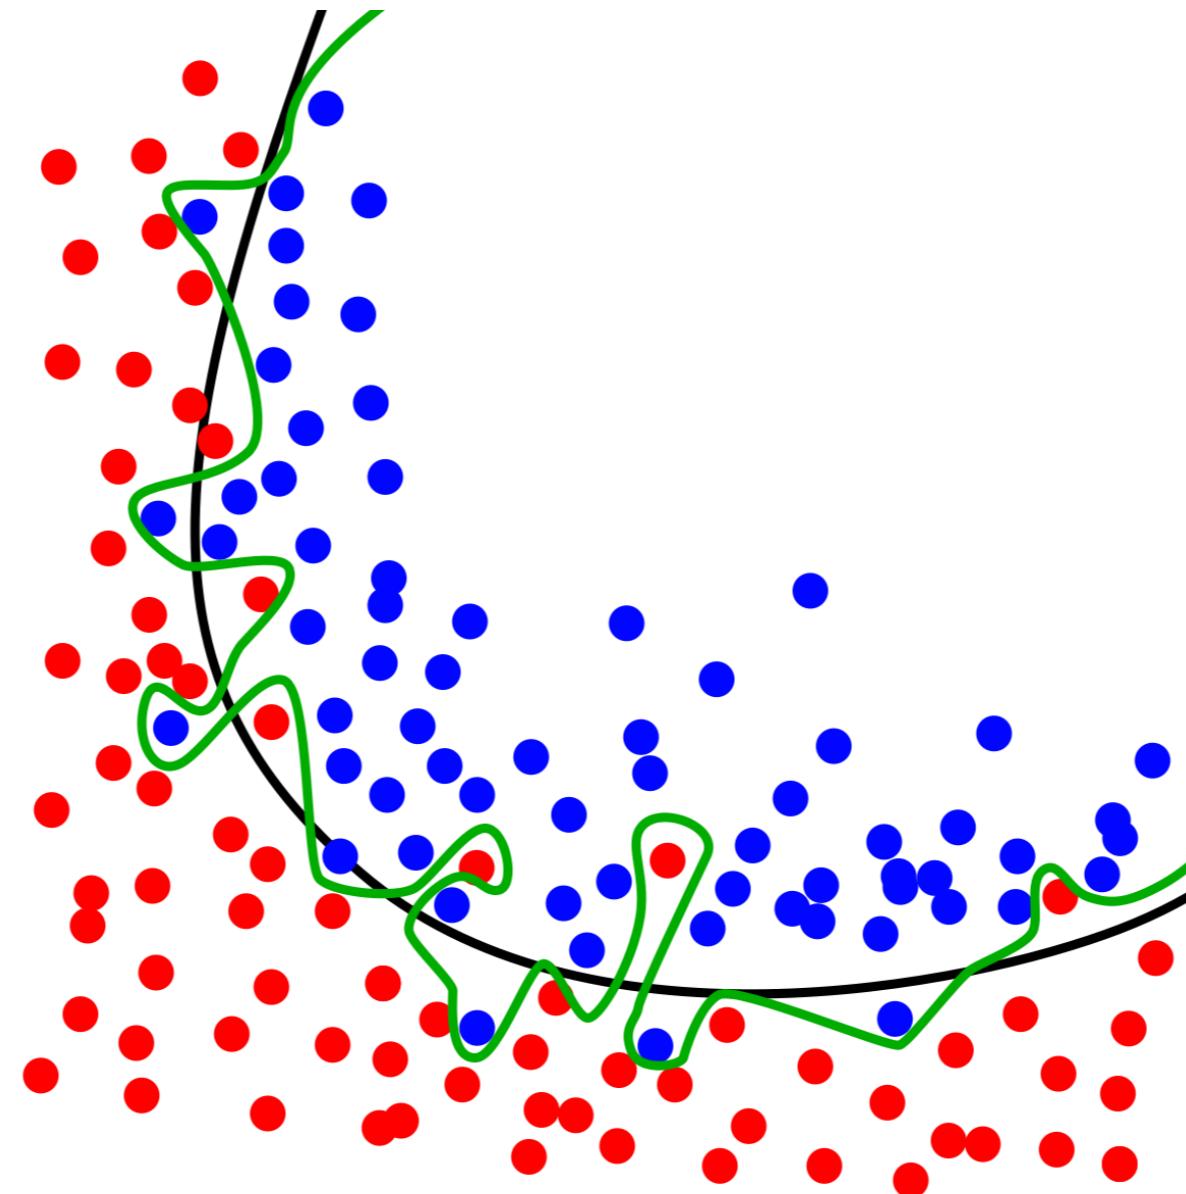
Evaluate step



Overfitting

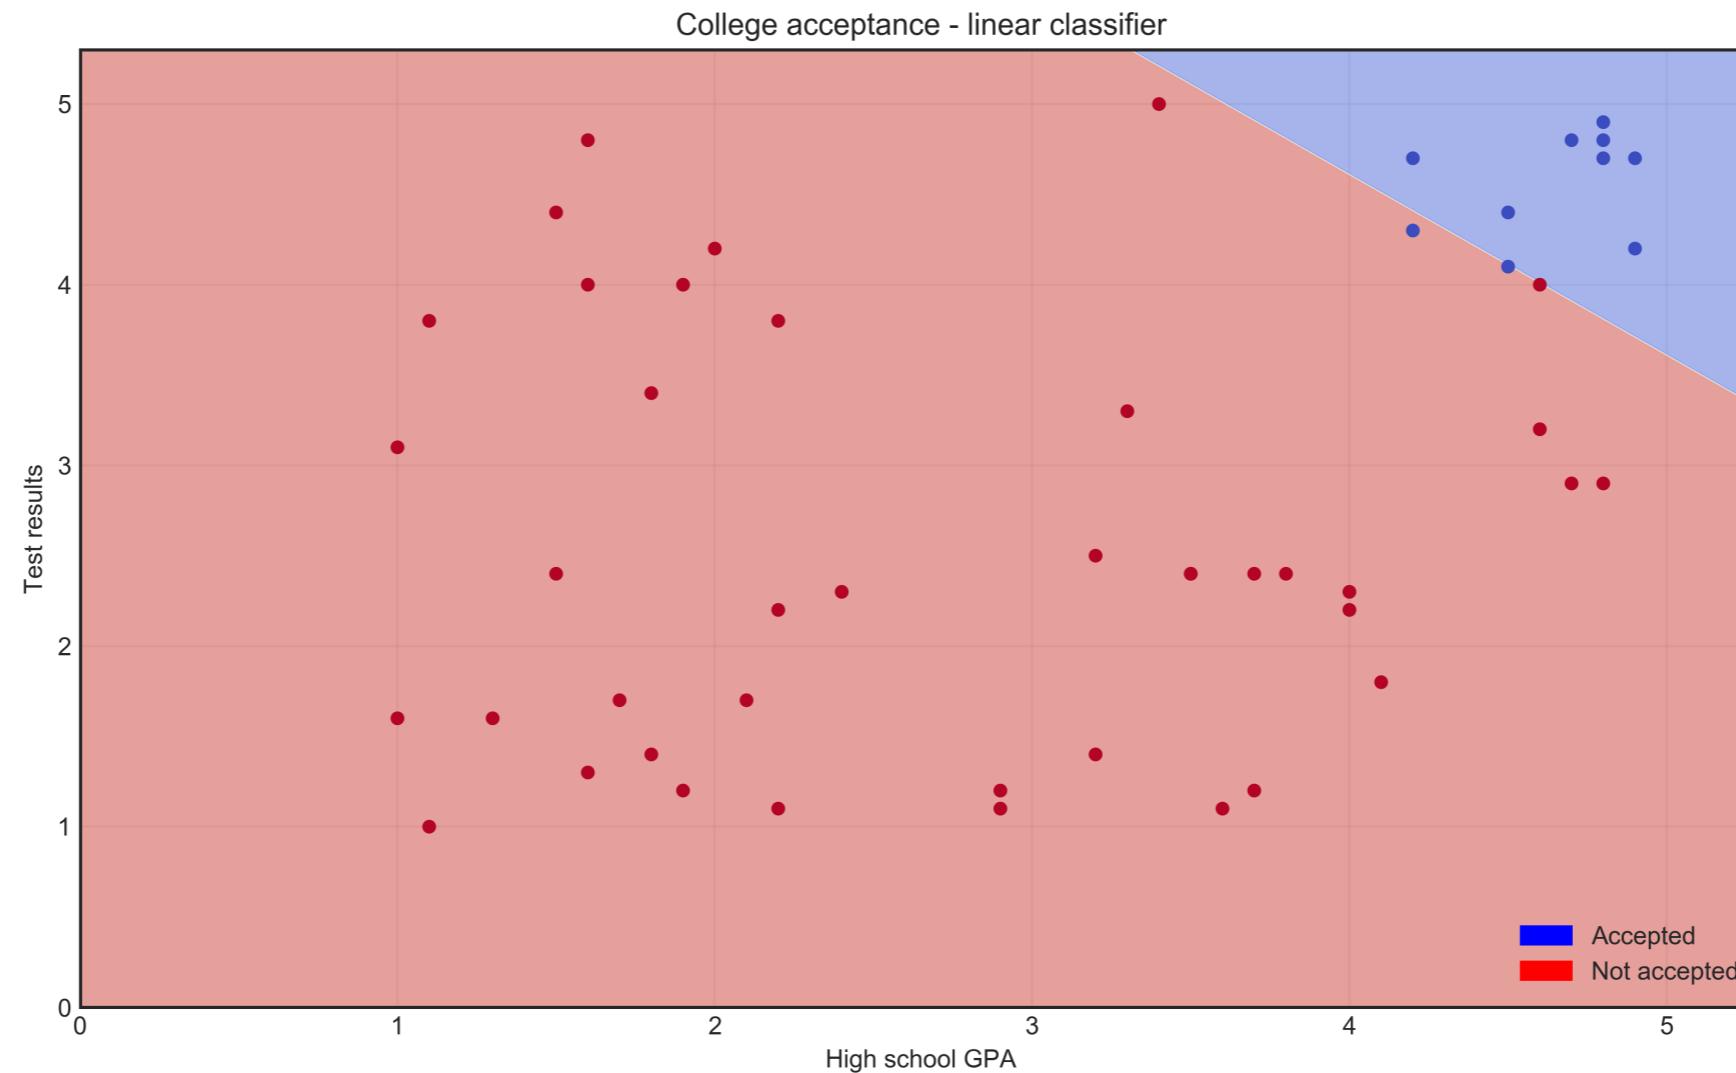
- Performs **great** on **training** data
- Performs **poorly** on **testing** data
- Model memorized **training** data and can't generalize learnings to new data
- Use **testing** set to check model **performance**

Illustrating overfitting

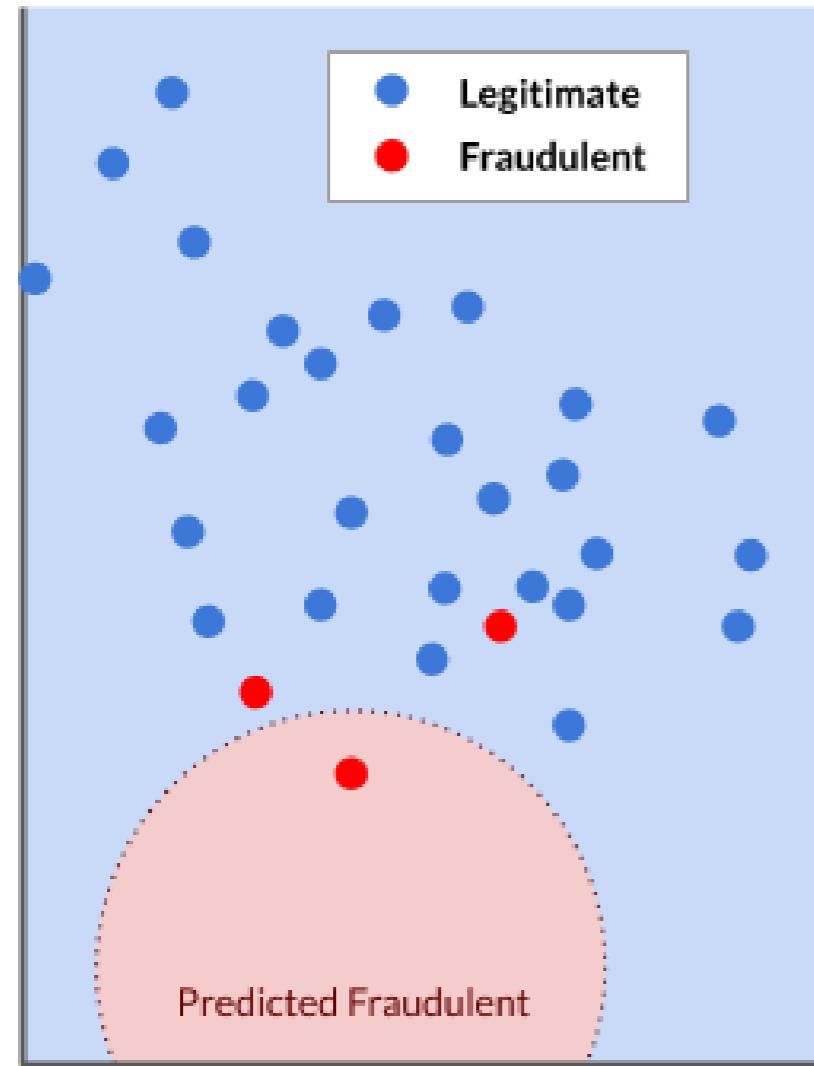


Accuracy

- Accuracy = **correctly classified observations / all observations**
- $48 / 50 = 96\%$



Limits of accuracy: fraud example

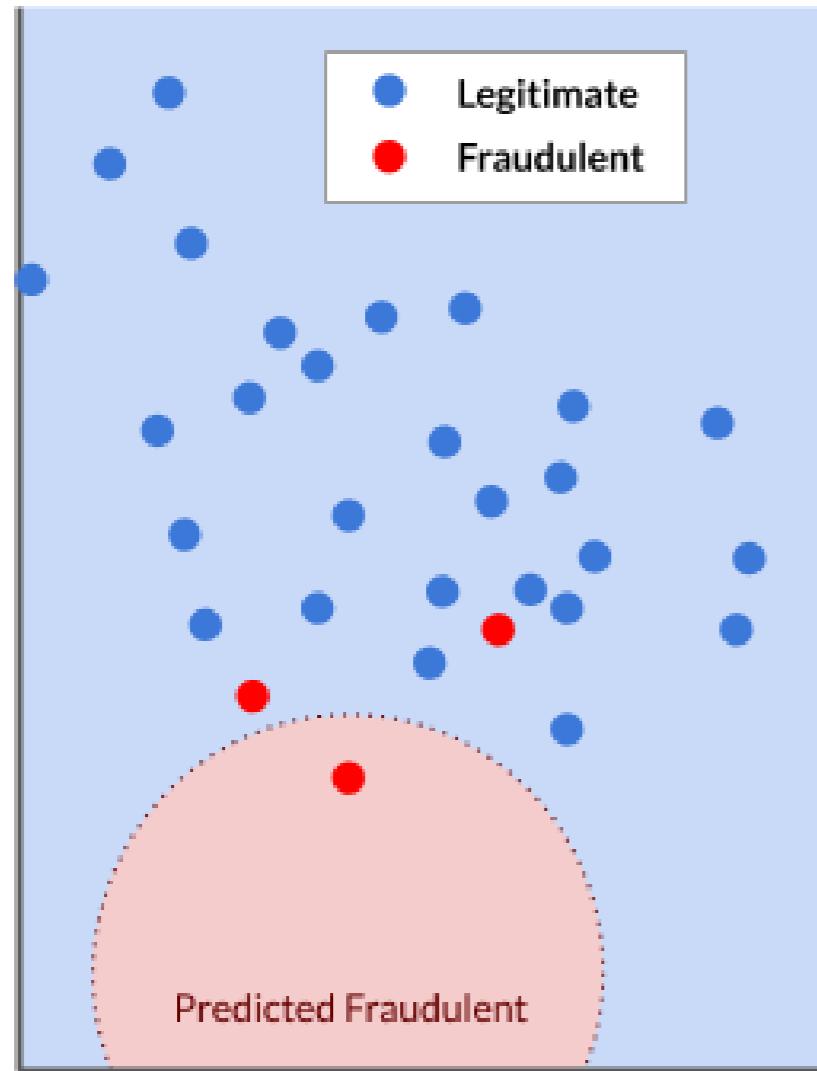


Accuracy of this model:

$$\frac{28 \text{ correctly classified}}{30 \text{ total points}} = 93.33\%$$

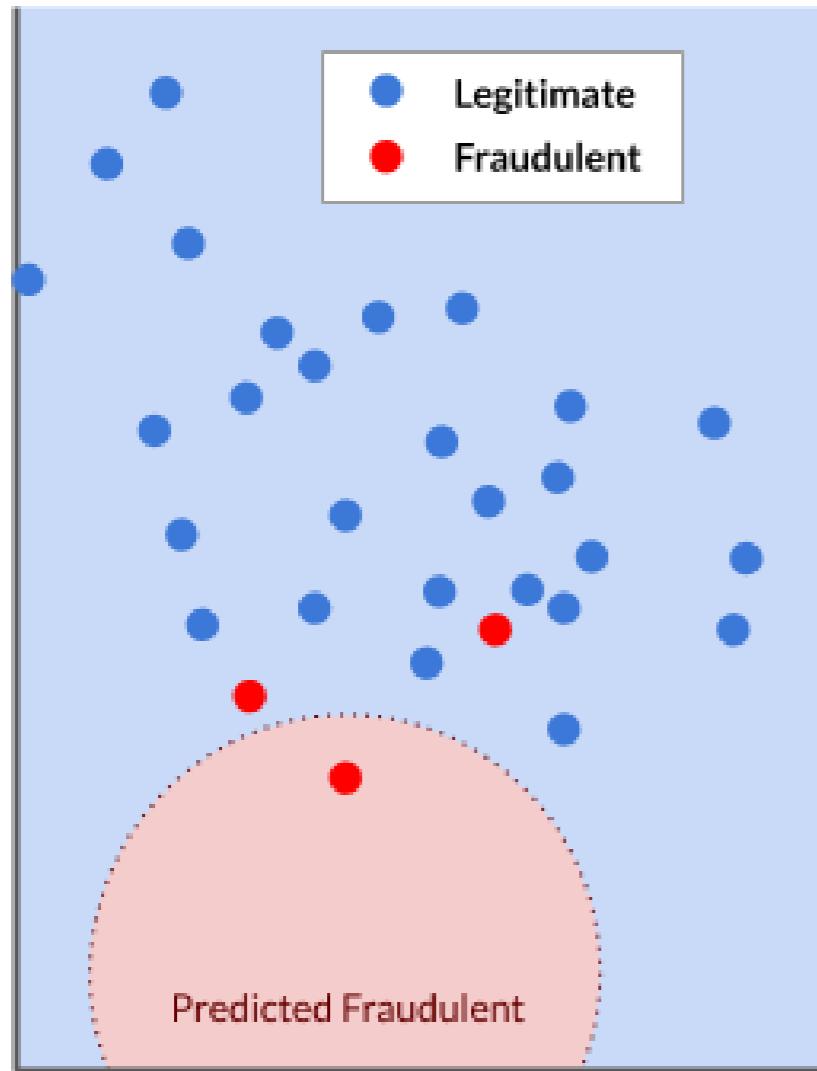
- Misses majority of fraudulent transactions
- Need a better metric

Confusion matrix



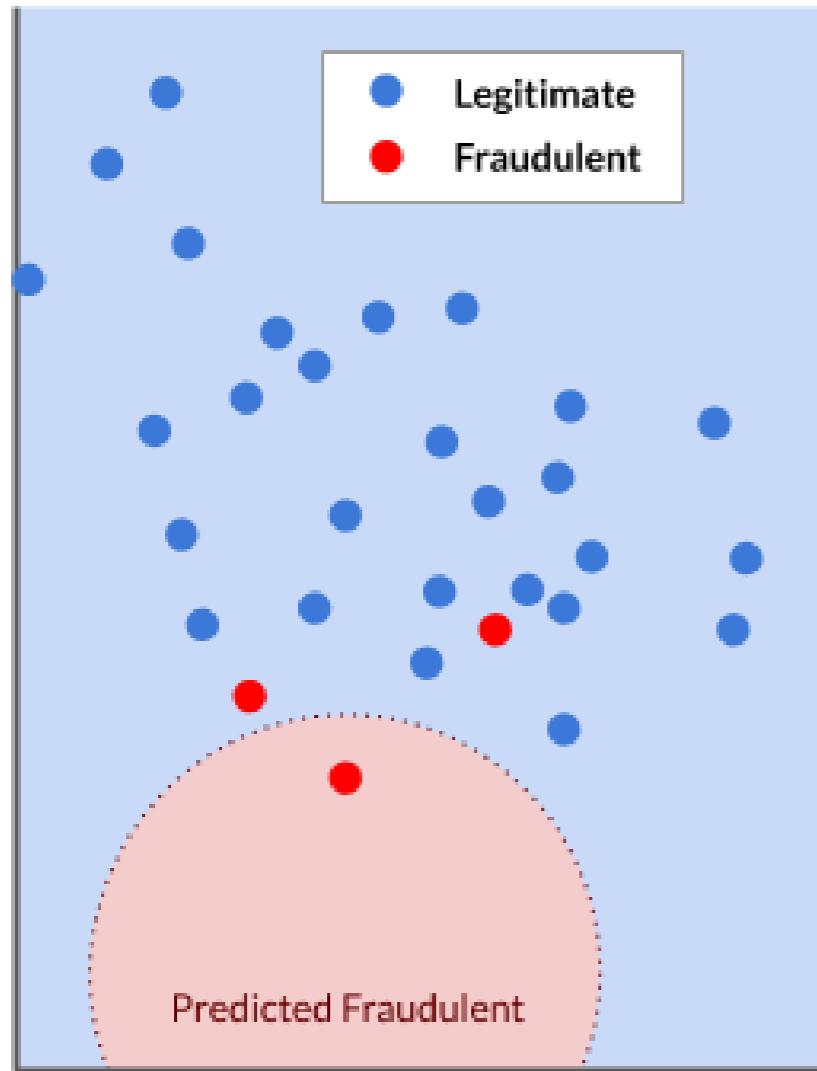
		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent		
	Not Fraudulent		

True positives



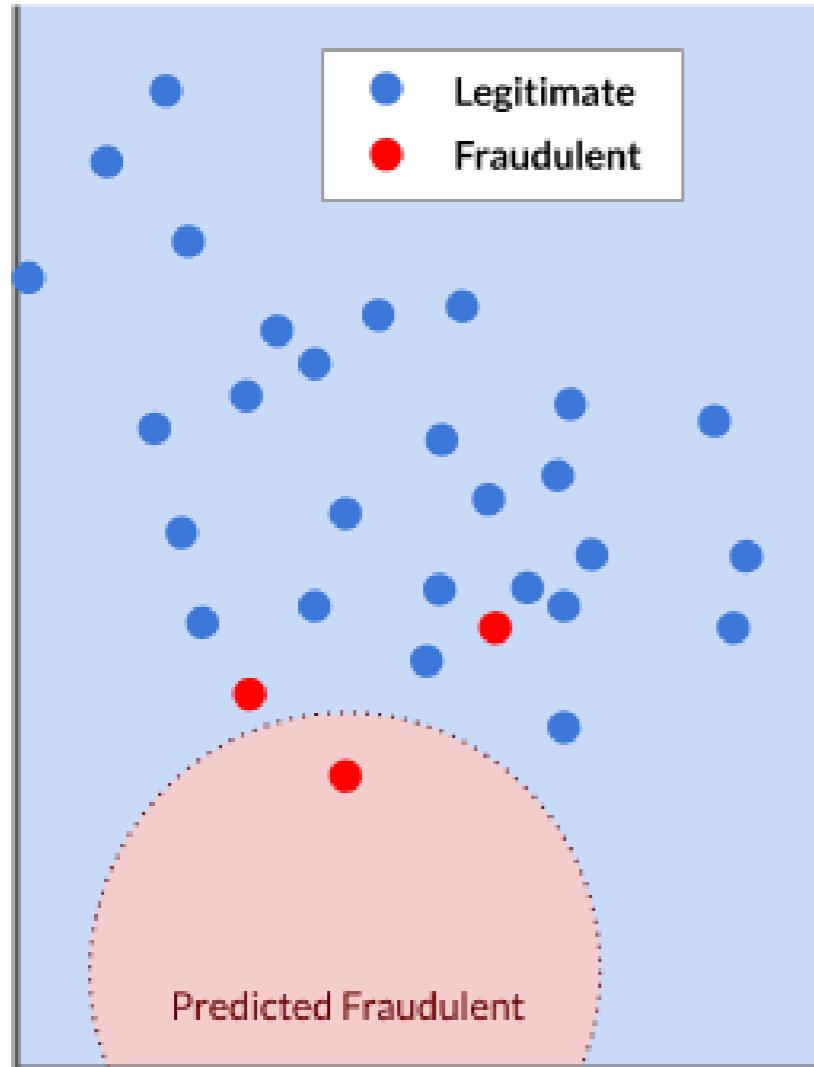
		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent		
	Not Fraudulent		

True positives



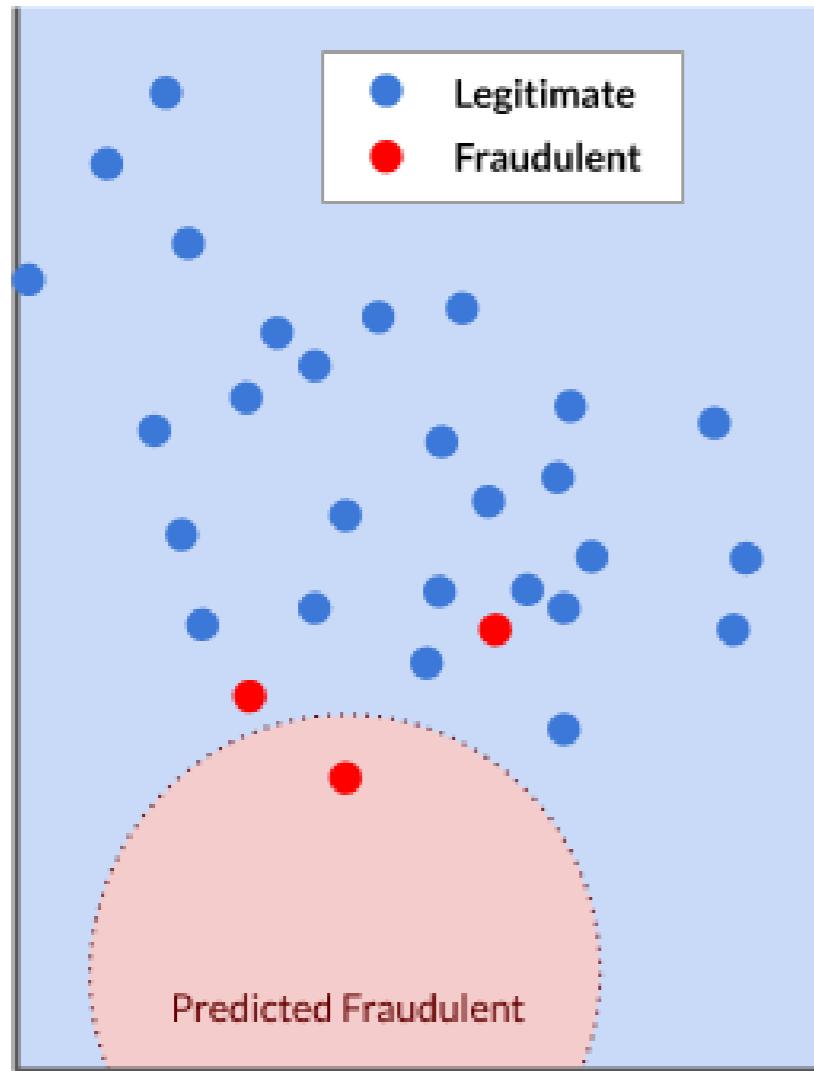
		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent	1 true positives	
	Not Fraudulent		

False negatives



		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent	1 true positives	
	Not Fraudulent		

False negatives

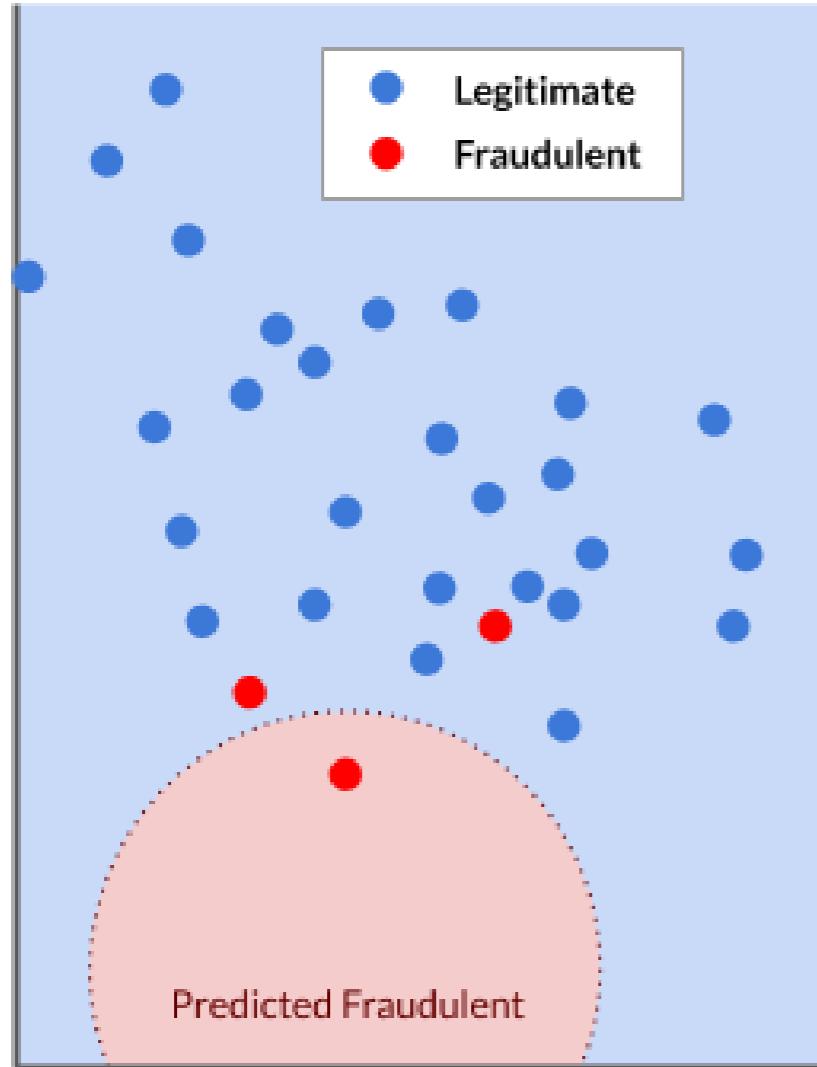


		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent	1 true positives	
	Not Fraudulent	2 false negatives	

Remembering False Negatives

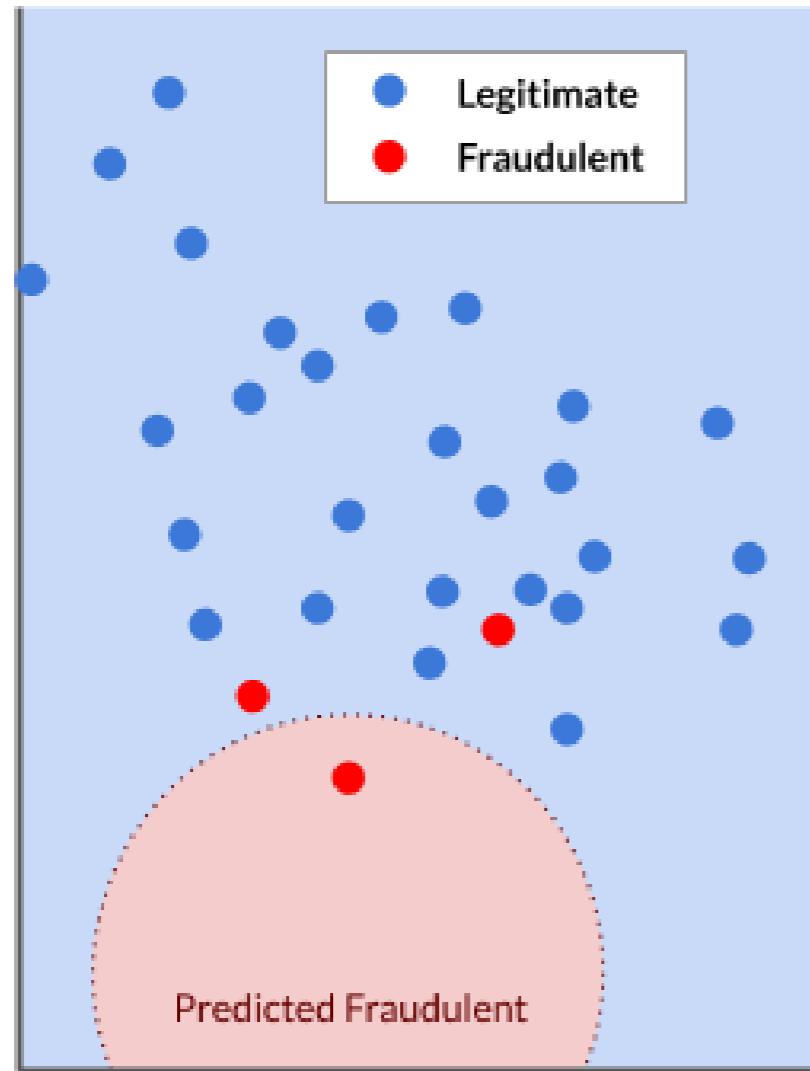


Fill out the rest...



		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent	1 true positives	
	Not Fraudulent	2 false negatives	

False positives, true negatives



		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent	1 true positives	0 false positives
	Not Fraudulent	2 false negatives	27 true negatives

Remembering False Positives



¹ <https://www.flickr.com/photos/59632563@N04/6104068209>

Sensitivity

		Actual values	
		Fraudulent	Not Fraudulent
Predicted	Fraudulent	1 true positives	0 false positives
	Not Fraudulent	2 false negatives	27 true negatives

How many fraudulent transactions did we classify correctly?

$$\text{Sensitivity} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}} = 1/3 = 33.33\%$$

- Rather mark legitimate transactions as suspicious than authorize fraudulent transactions

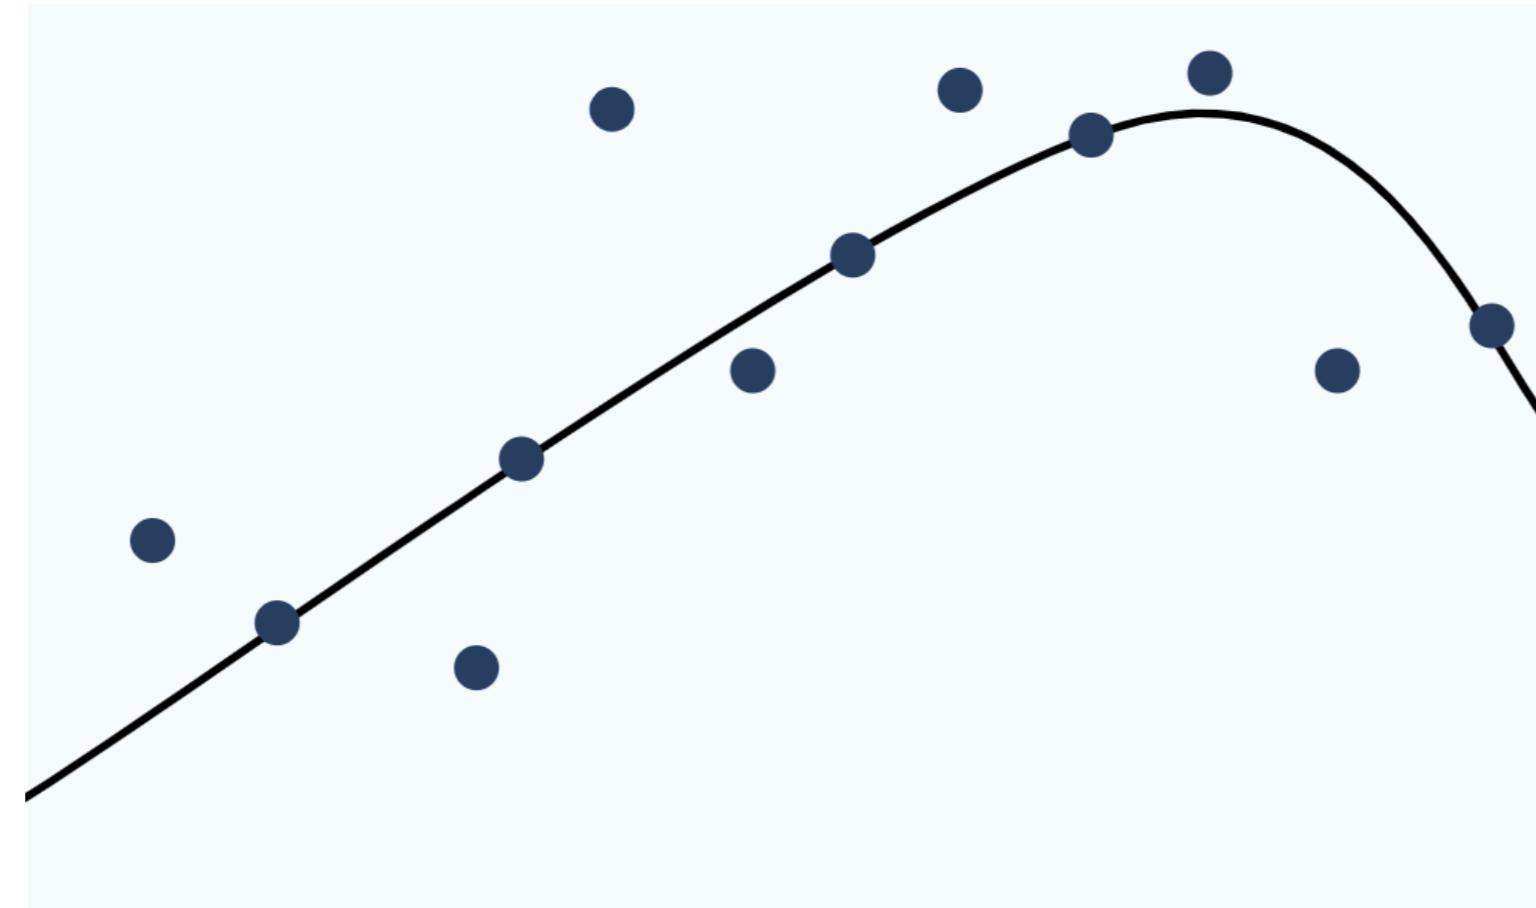
Specificity

$$Specificity = \frac{true\ negatives}{true\ negatives + false\ positives}$$

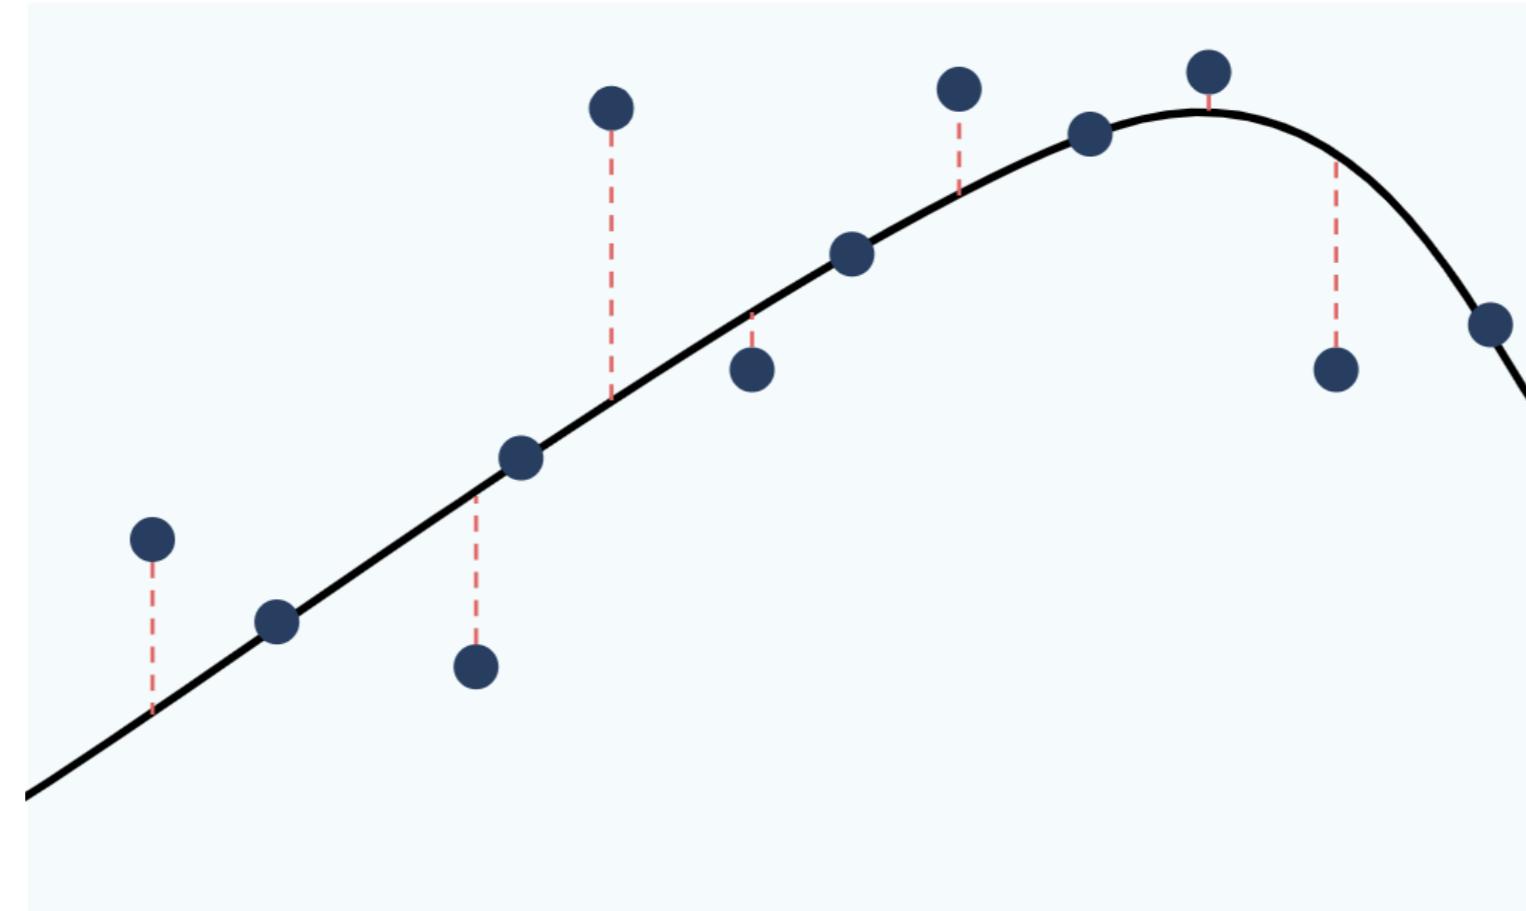
Spam filter:

- Rather send spam to inbox than send real emails to the spam folder

Evaluating regression



Evaluating regression



- Error = distance between point (actual value) and line (predicted value)
- Many ways calculate this. e.g, root mean square error

Unsupervised learning



¹ <https://www.flickr.com/photos/micahdowty/8540188997>

Let's practice!

MACHINE LEARNING FOR EVERYONE

Improving performance

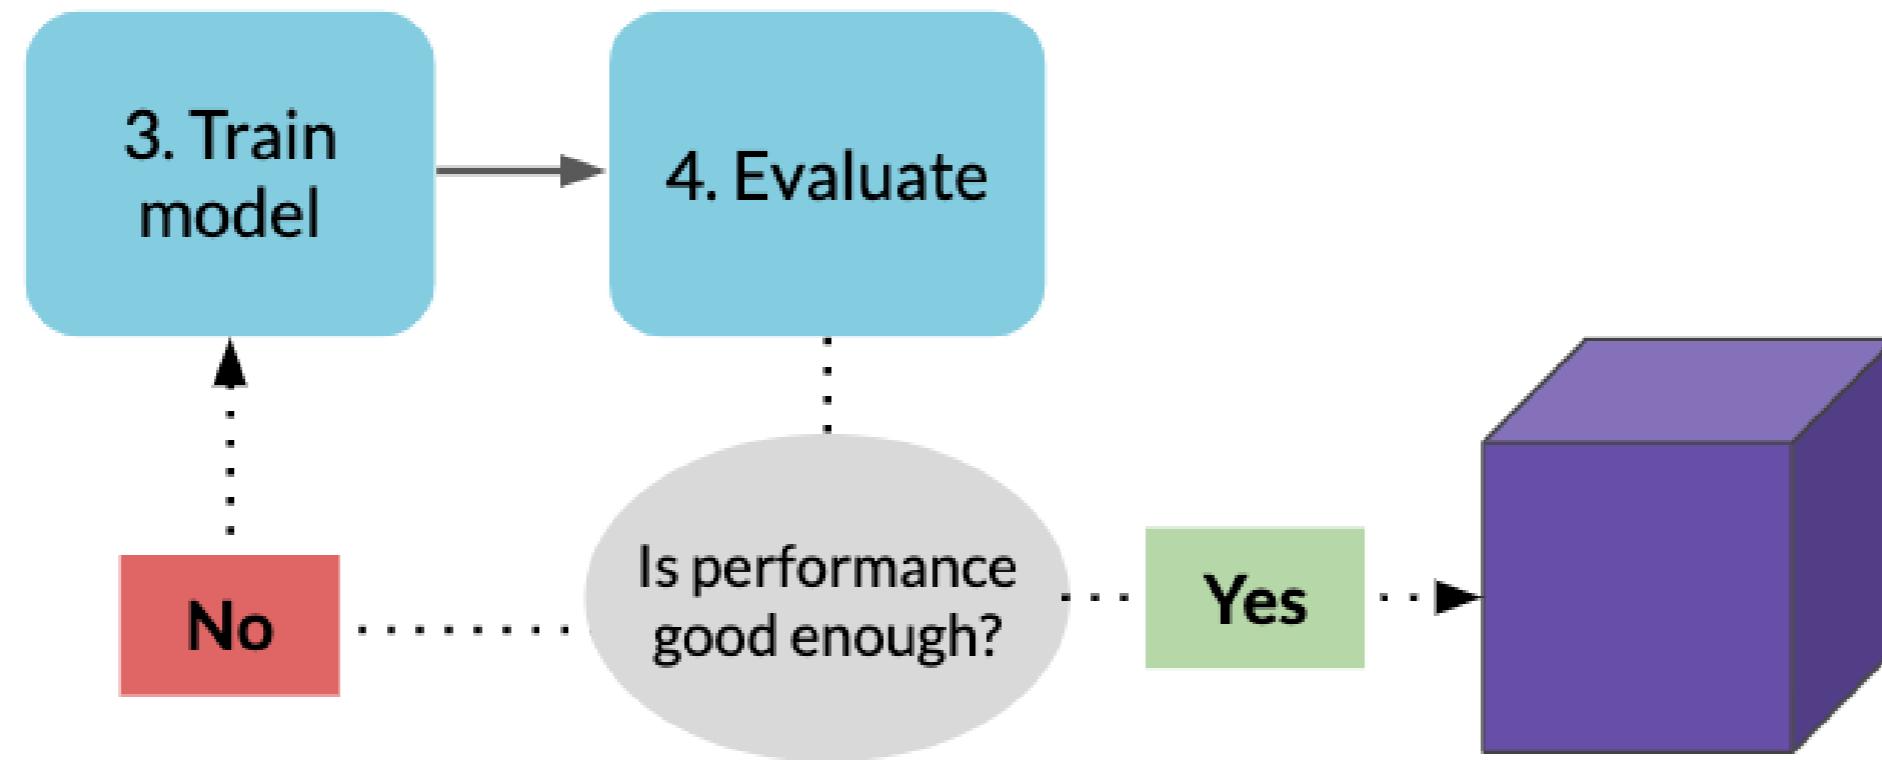
MACHINE LEARNING FOR EVERYONE



Hadrien Lacroix

Content Developer at DataCamp

Machine learning workflow



Several options

- Dimensionality reduction
- Hyperparameter tuning
- Ensemble methods

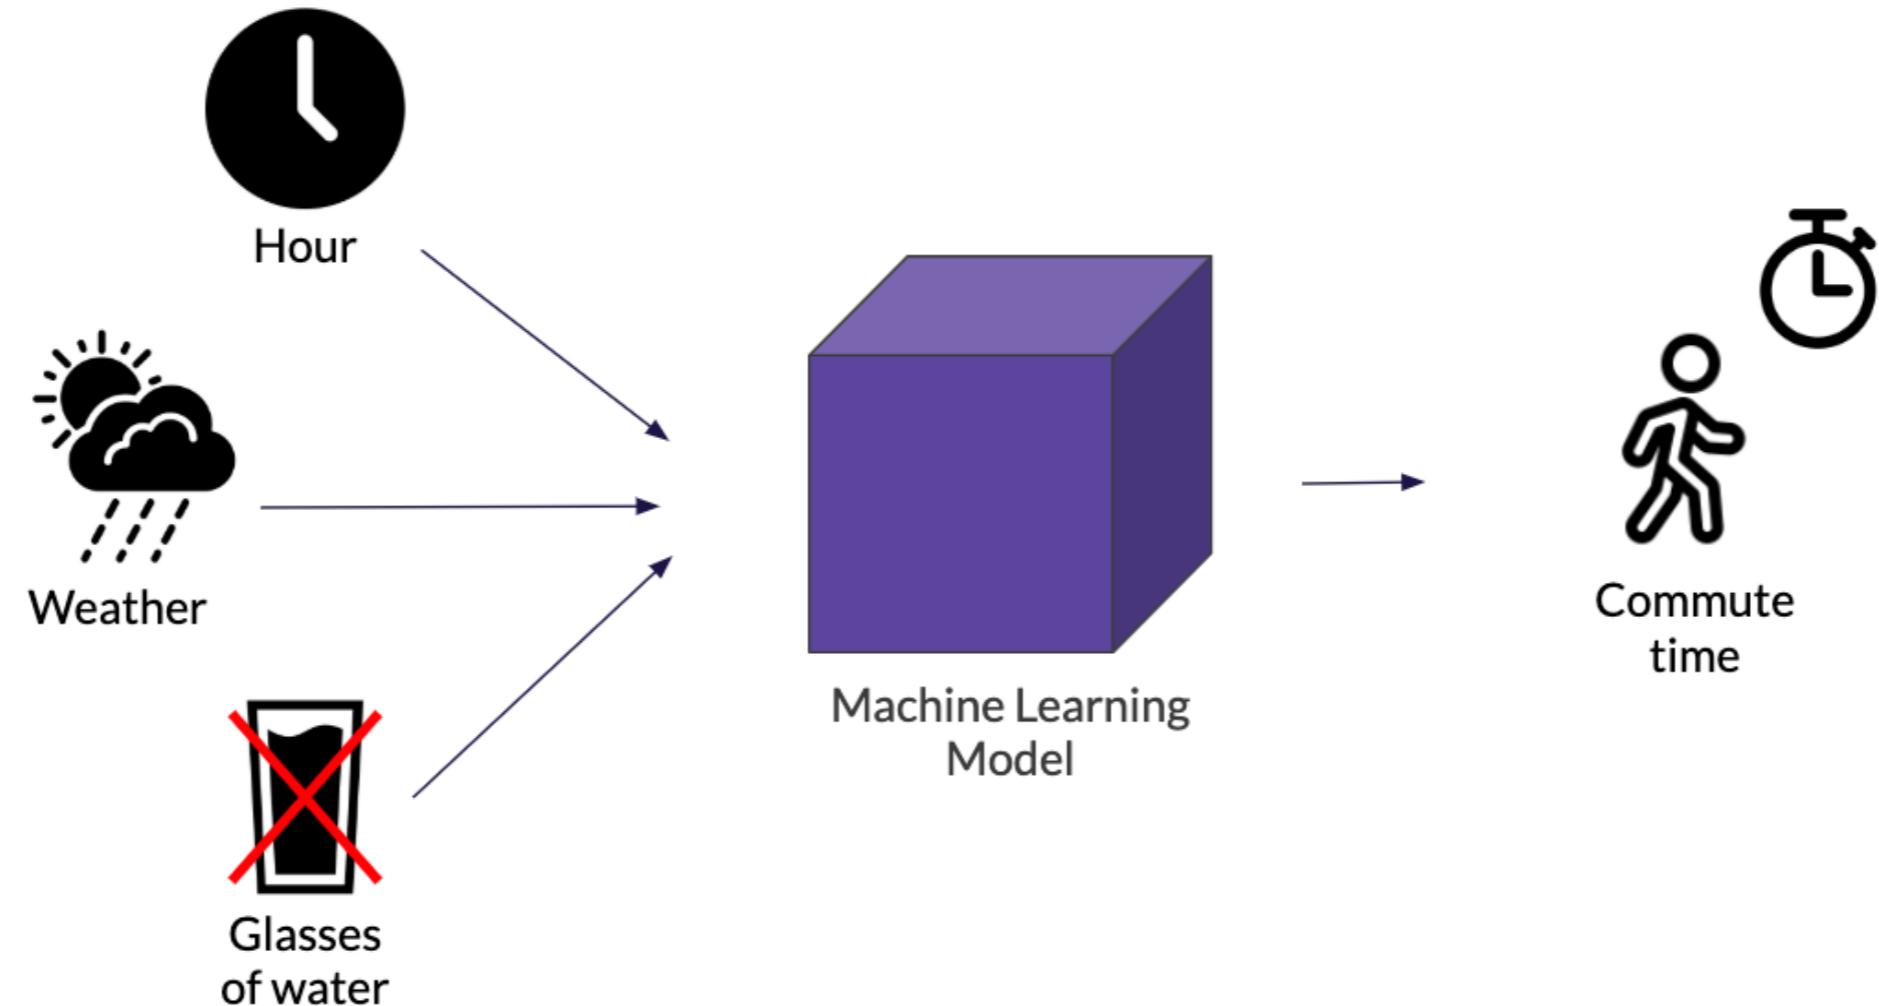
Dimensionality reduction

Reducing the number of features



Dimensionality reduction: example

Irrelevance: some features don't carry useful information



Dimensionality reduction: example

Correlation: some features carry similar information

- Keep only one feature
 - e.g. *height* and *shoe size* --> *height*
- Collapse multiple features into one underlying feature
 - e.g. *height* and *weight* --> *Body Mass Index*

Hyperparameter tuning



Hyperparameter tuning



Hyperparameter tuning



Hyperparameter tuning



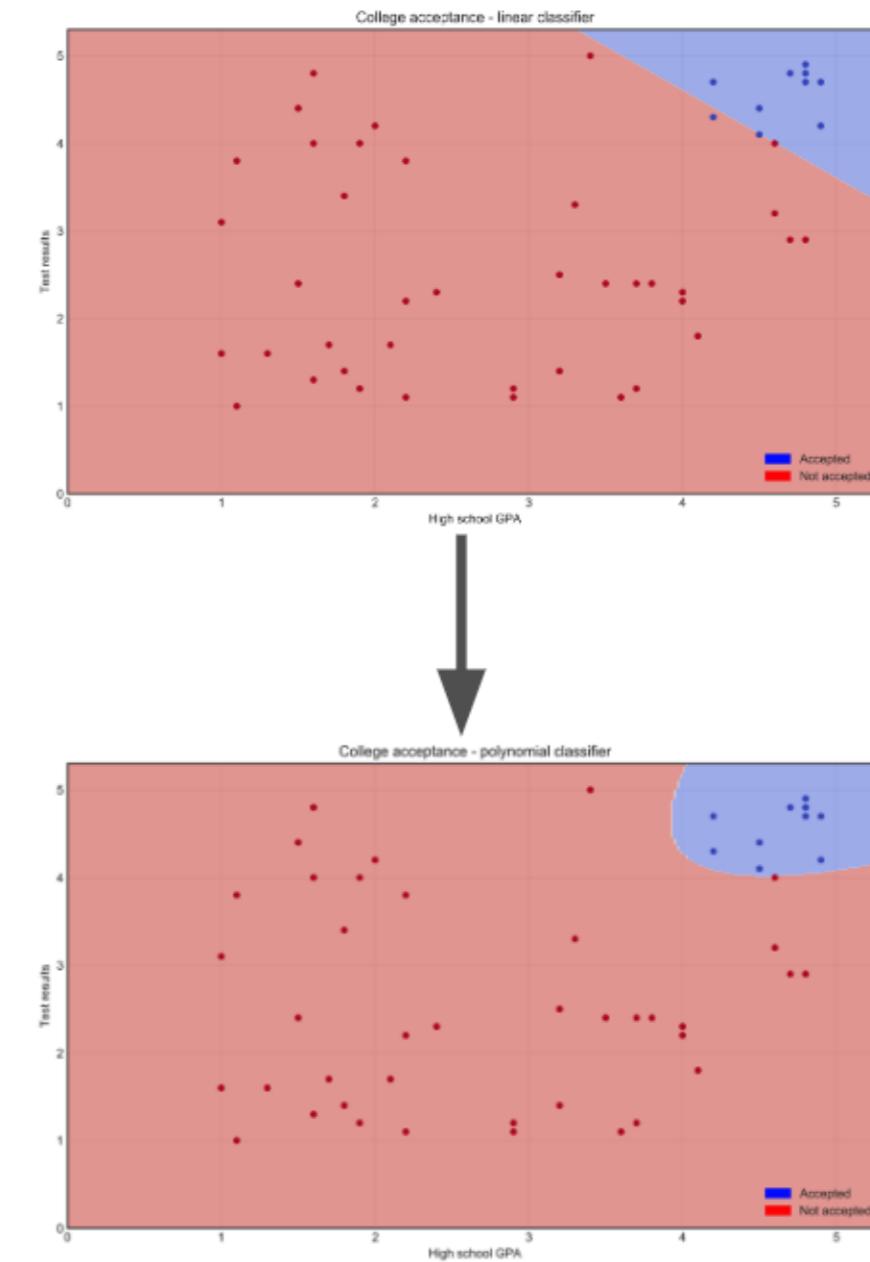
Hyperparameter tuning



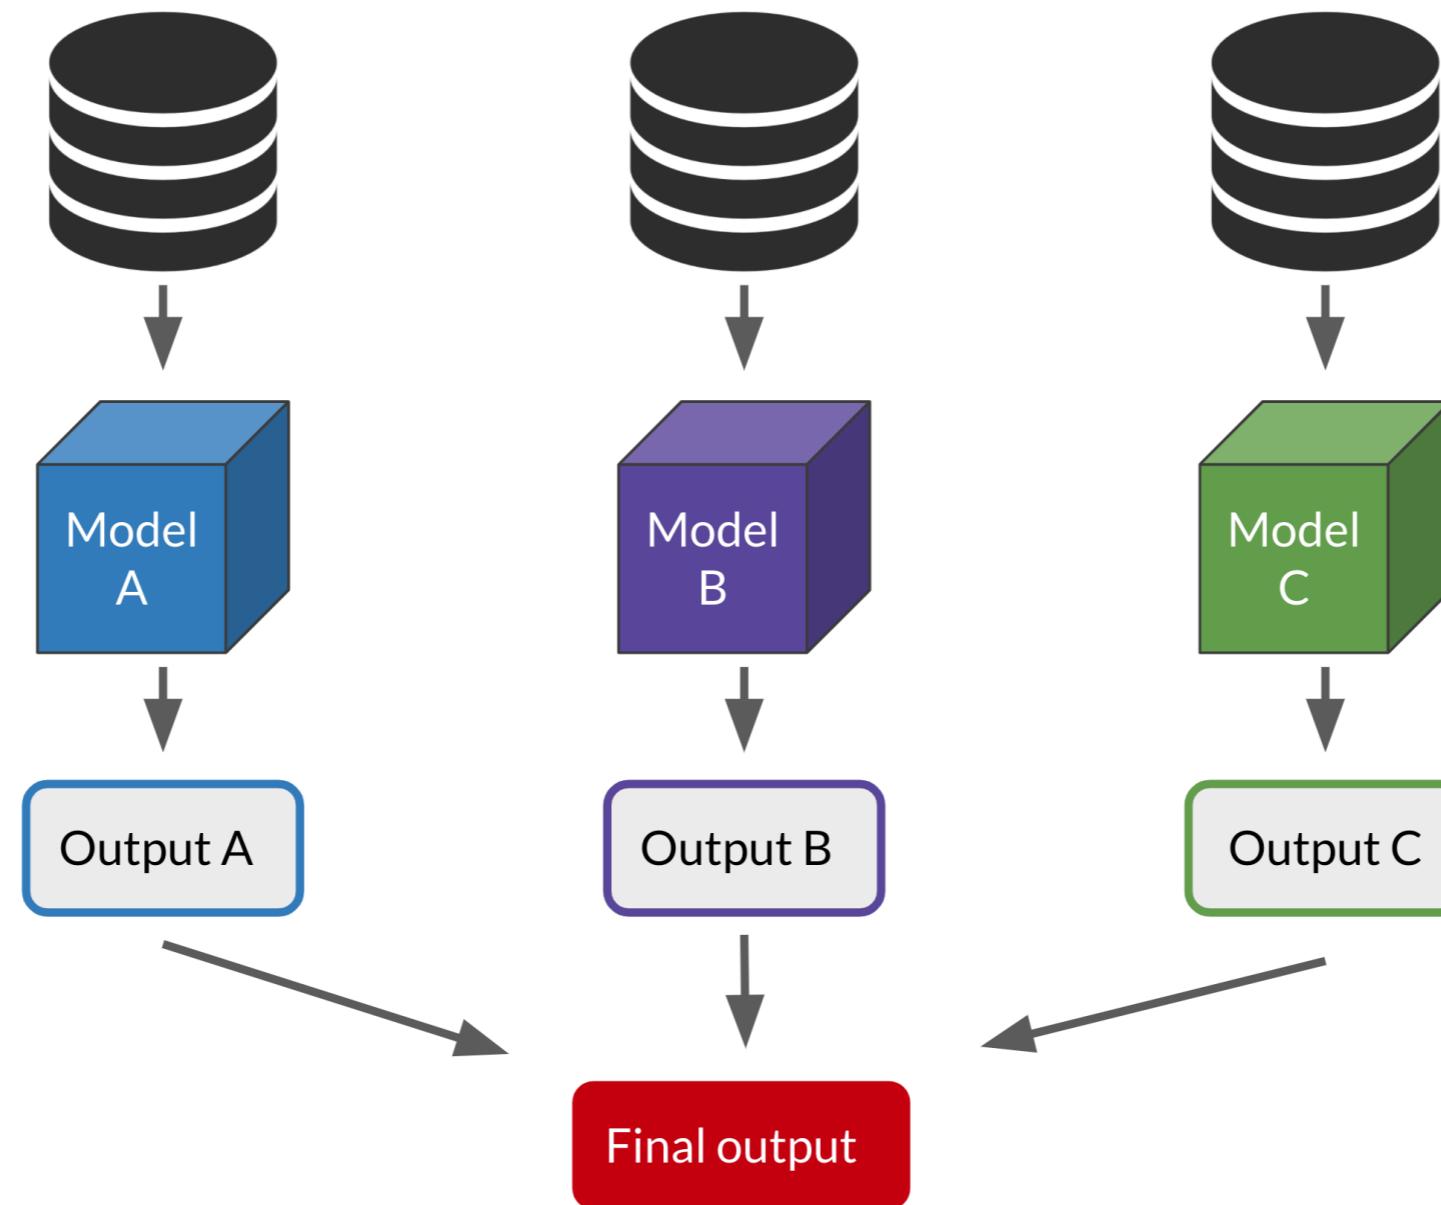
Hyperparameter tuning: example

SVM algorithm hyperparameters:

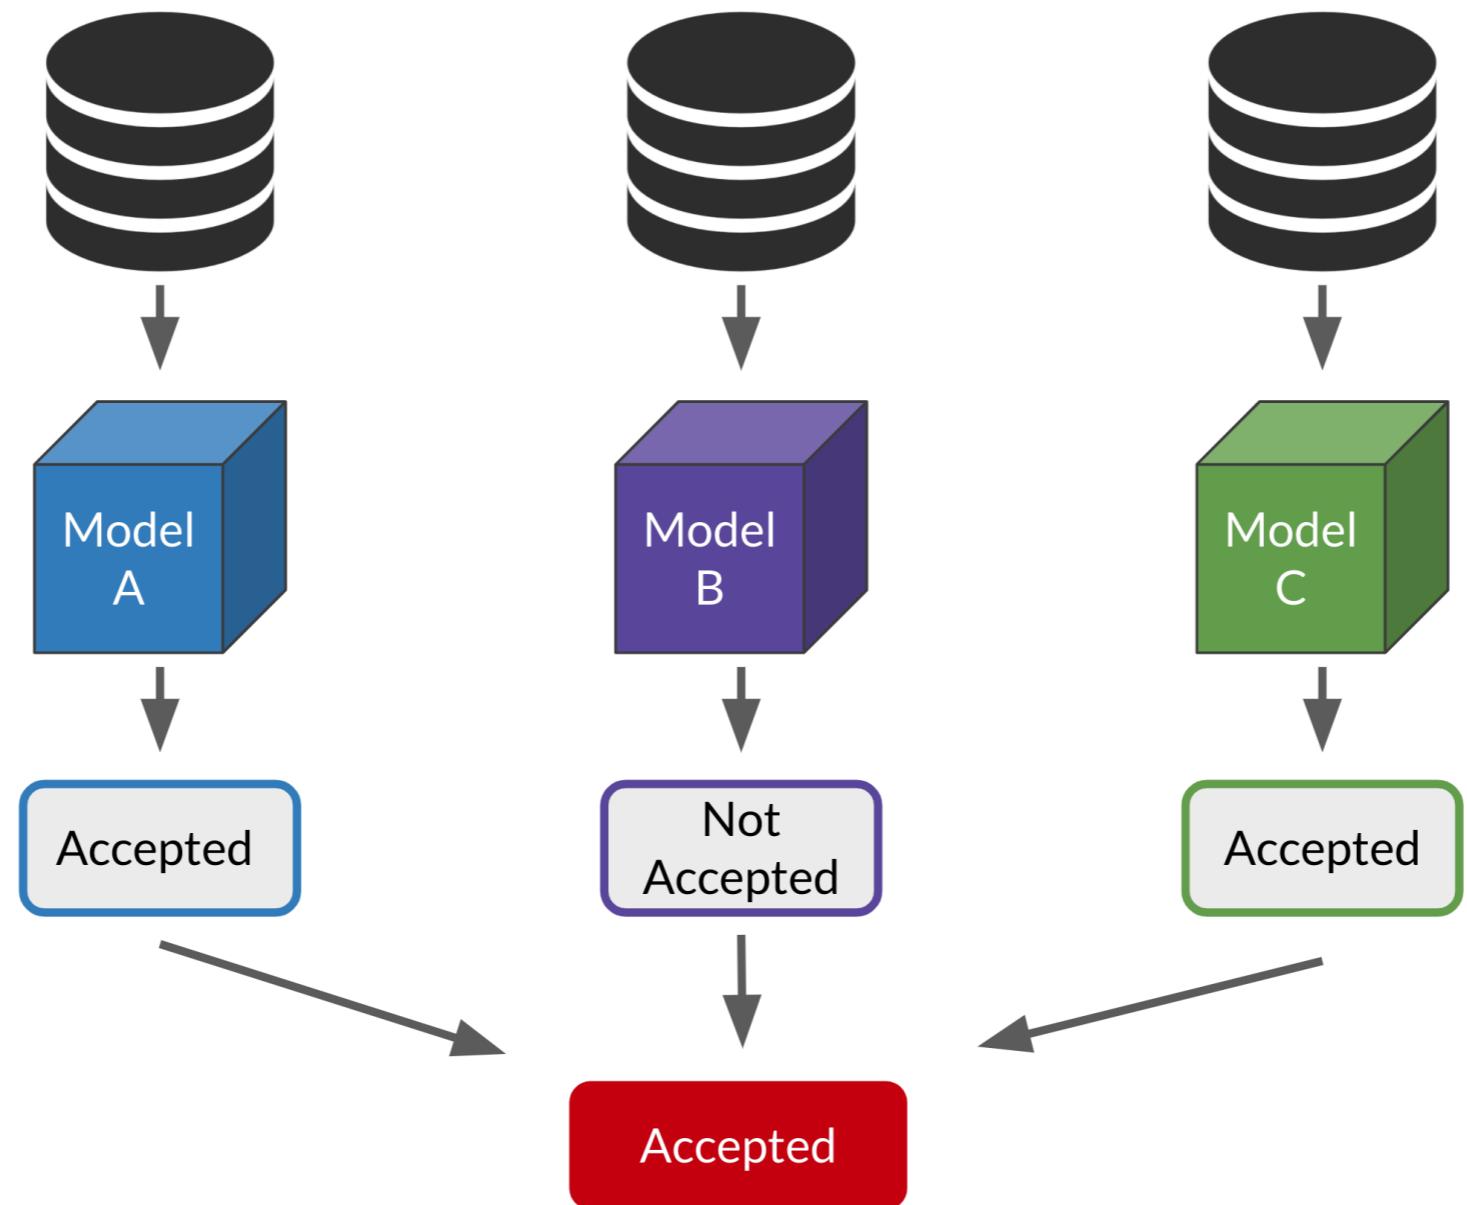
- kernel : "linear" --> "poly"
- C
- degree
- gamma
- shrinking
- coef0
- tol
- ...



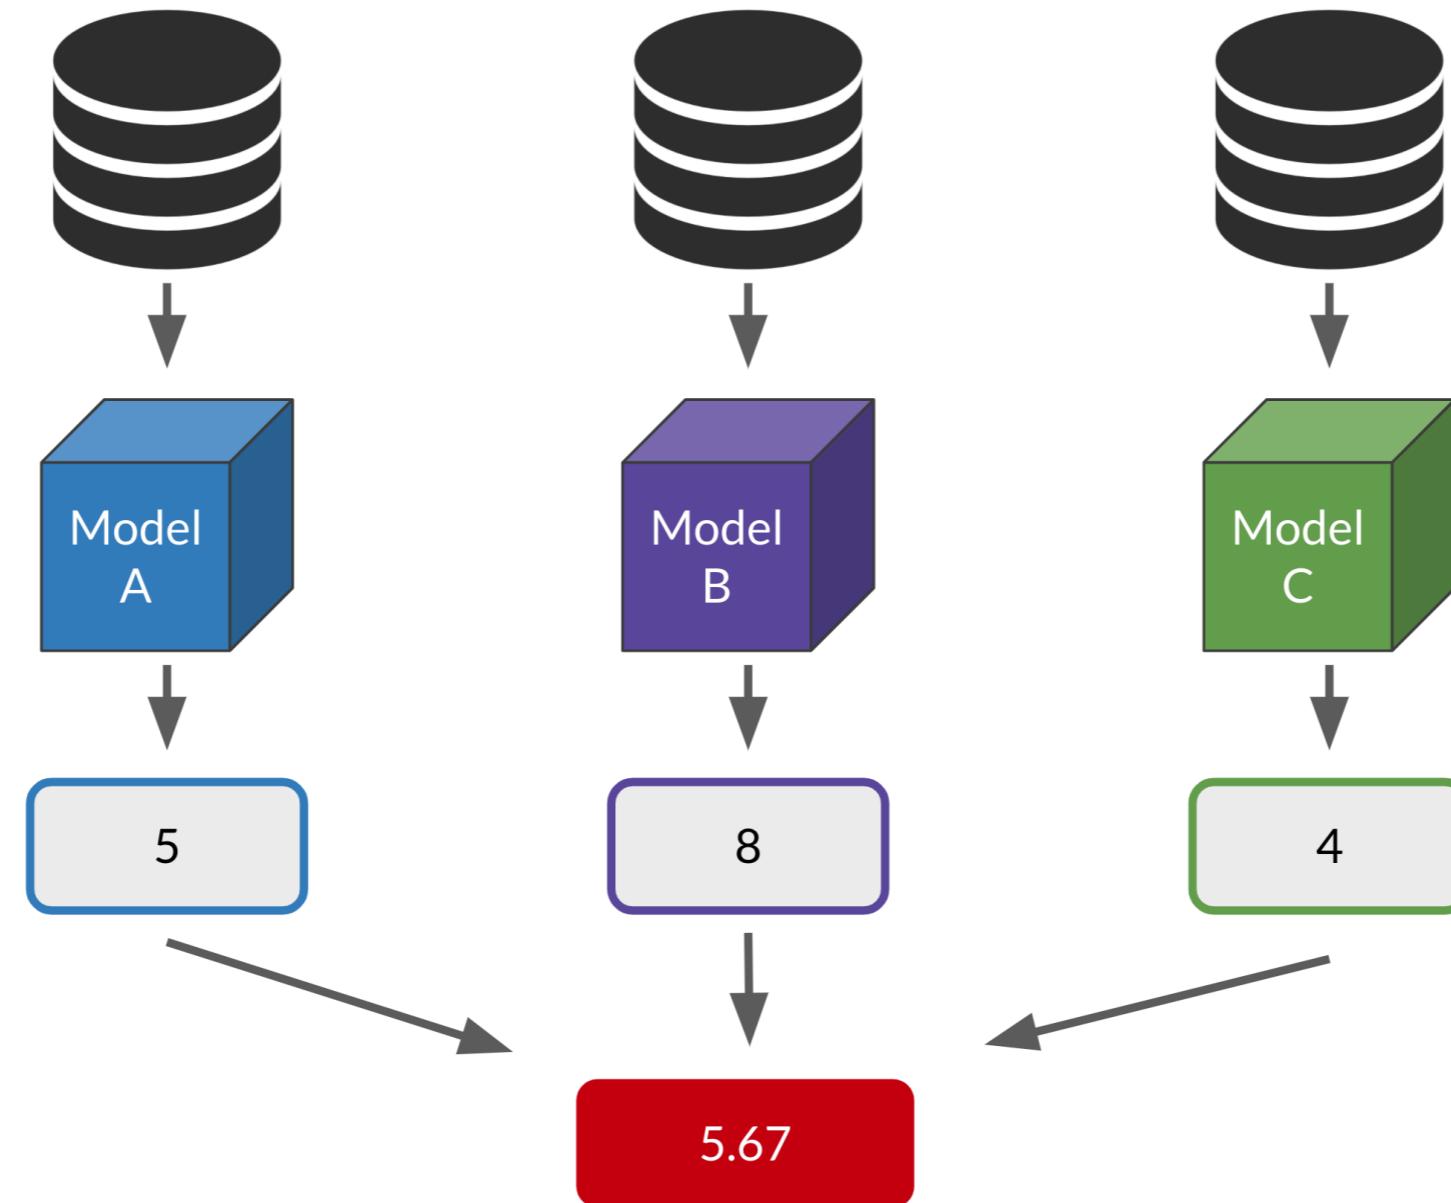
Ensemble methods



Ensemble methods: classification



Ensemble methods: regression



Let's practice!

MACHINE LEARNING FOR EVERYONE